

Relational Database Design

Charles Severance
www.wa4e.com



<http://www.wa4e.com/lectures/SQL-02-MySQL-Design-Handout.txt>

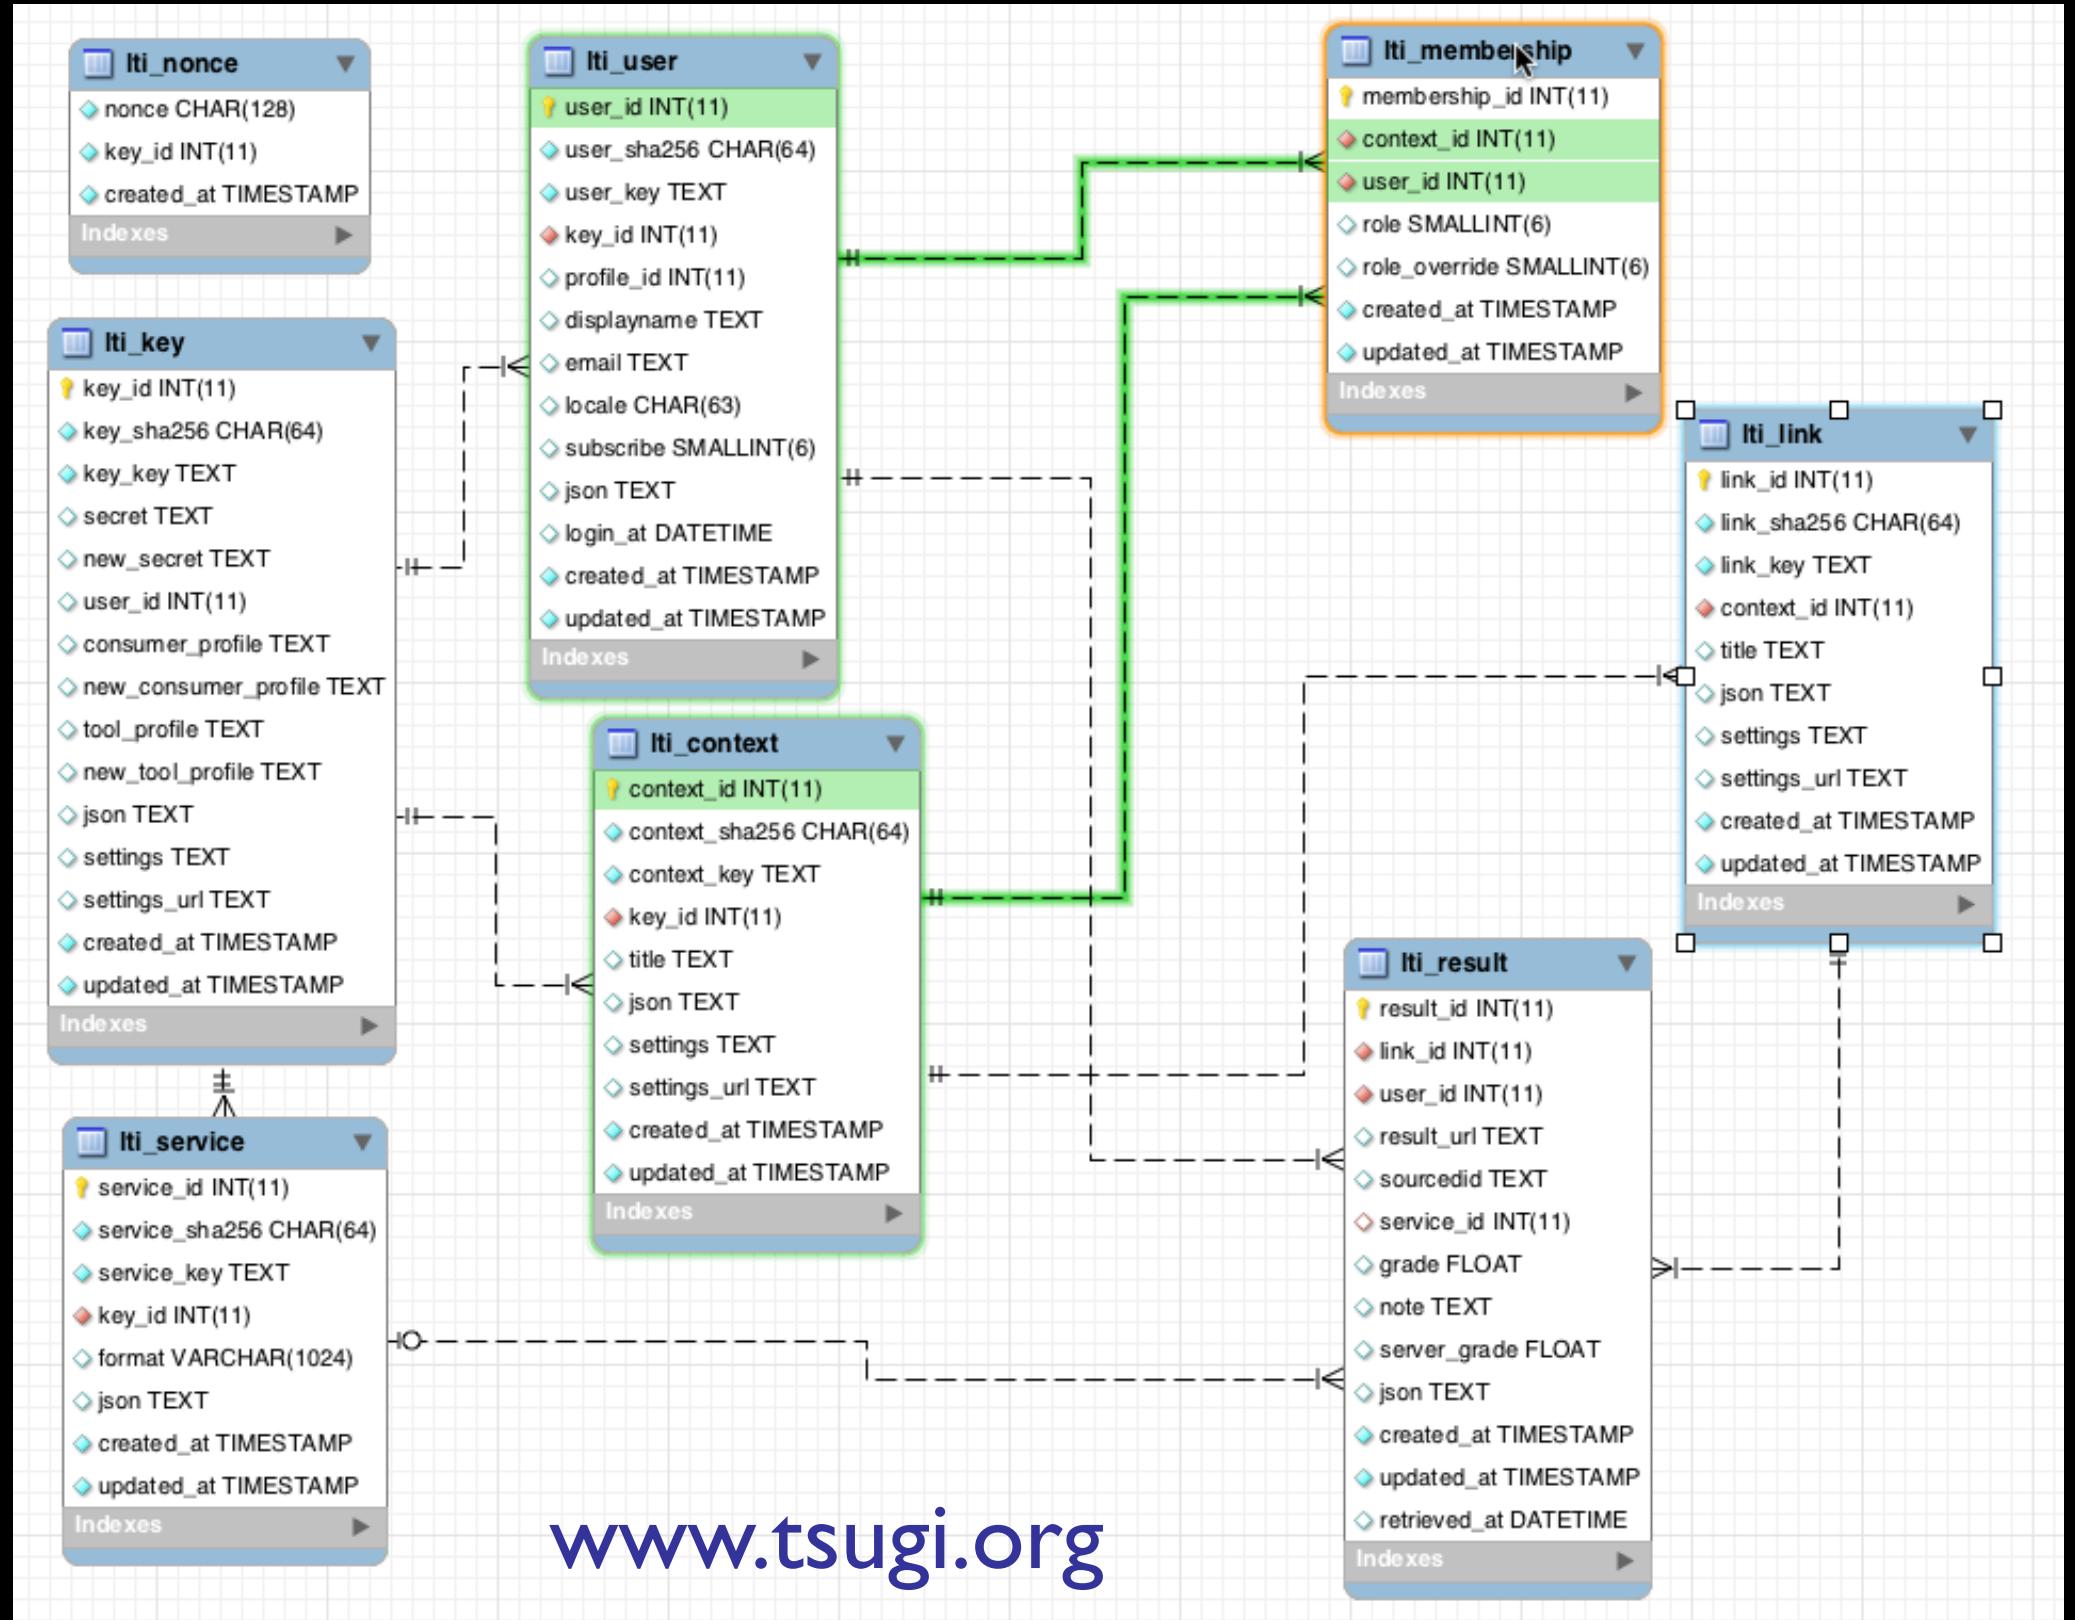


Relational Database Design

http://en.wikipedia.org/wiki/Relational_model

Database Design

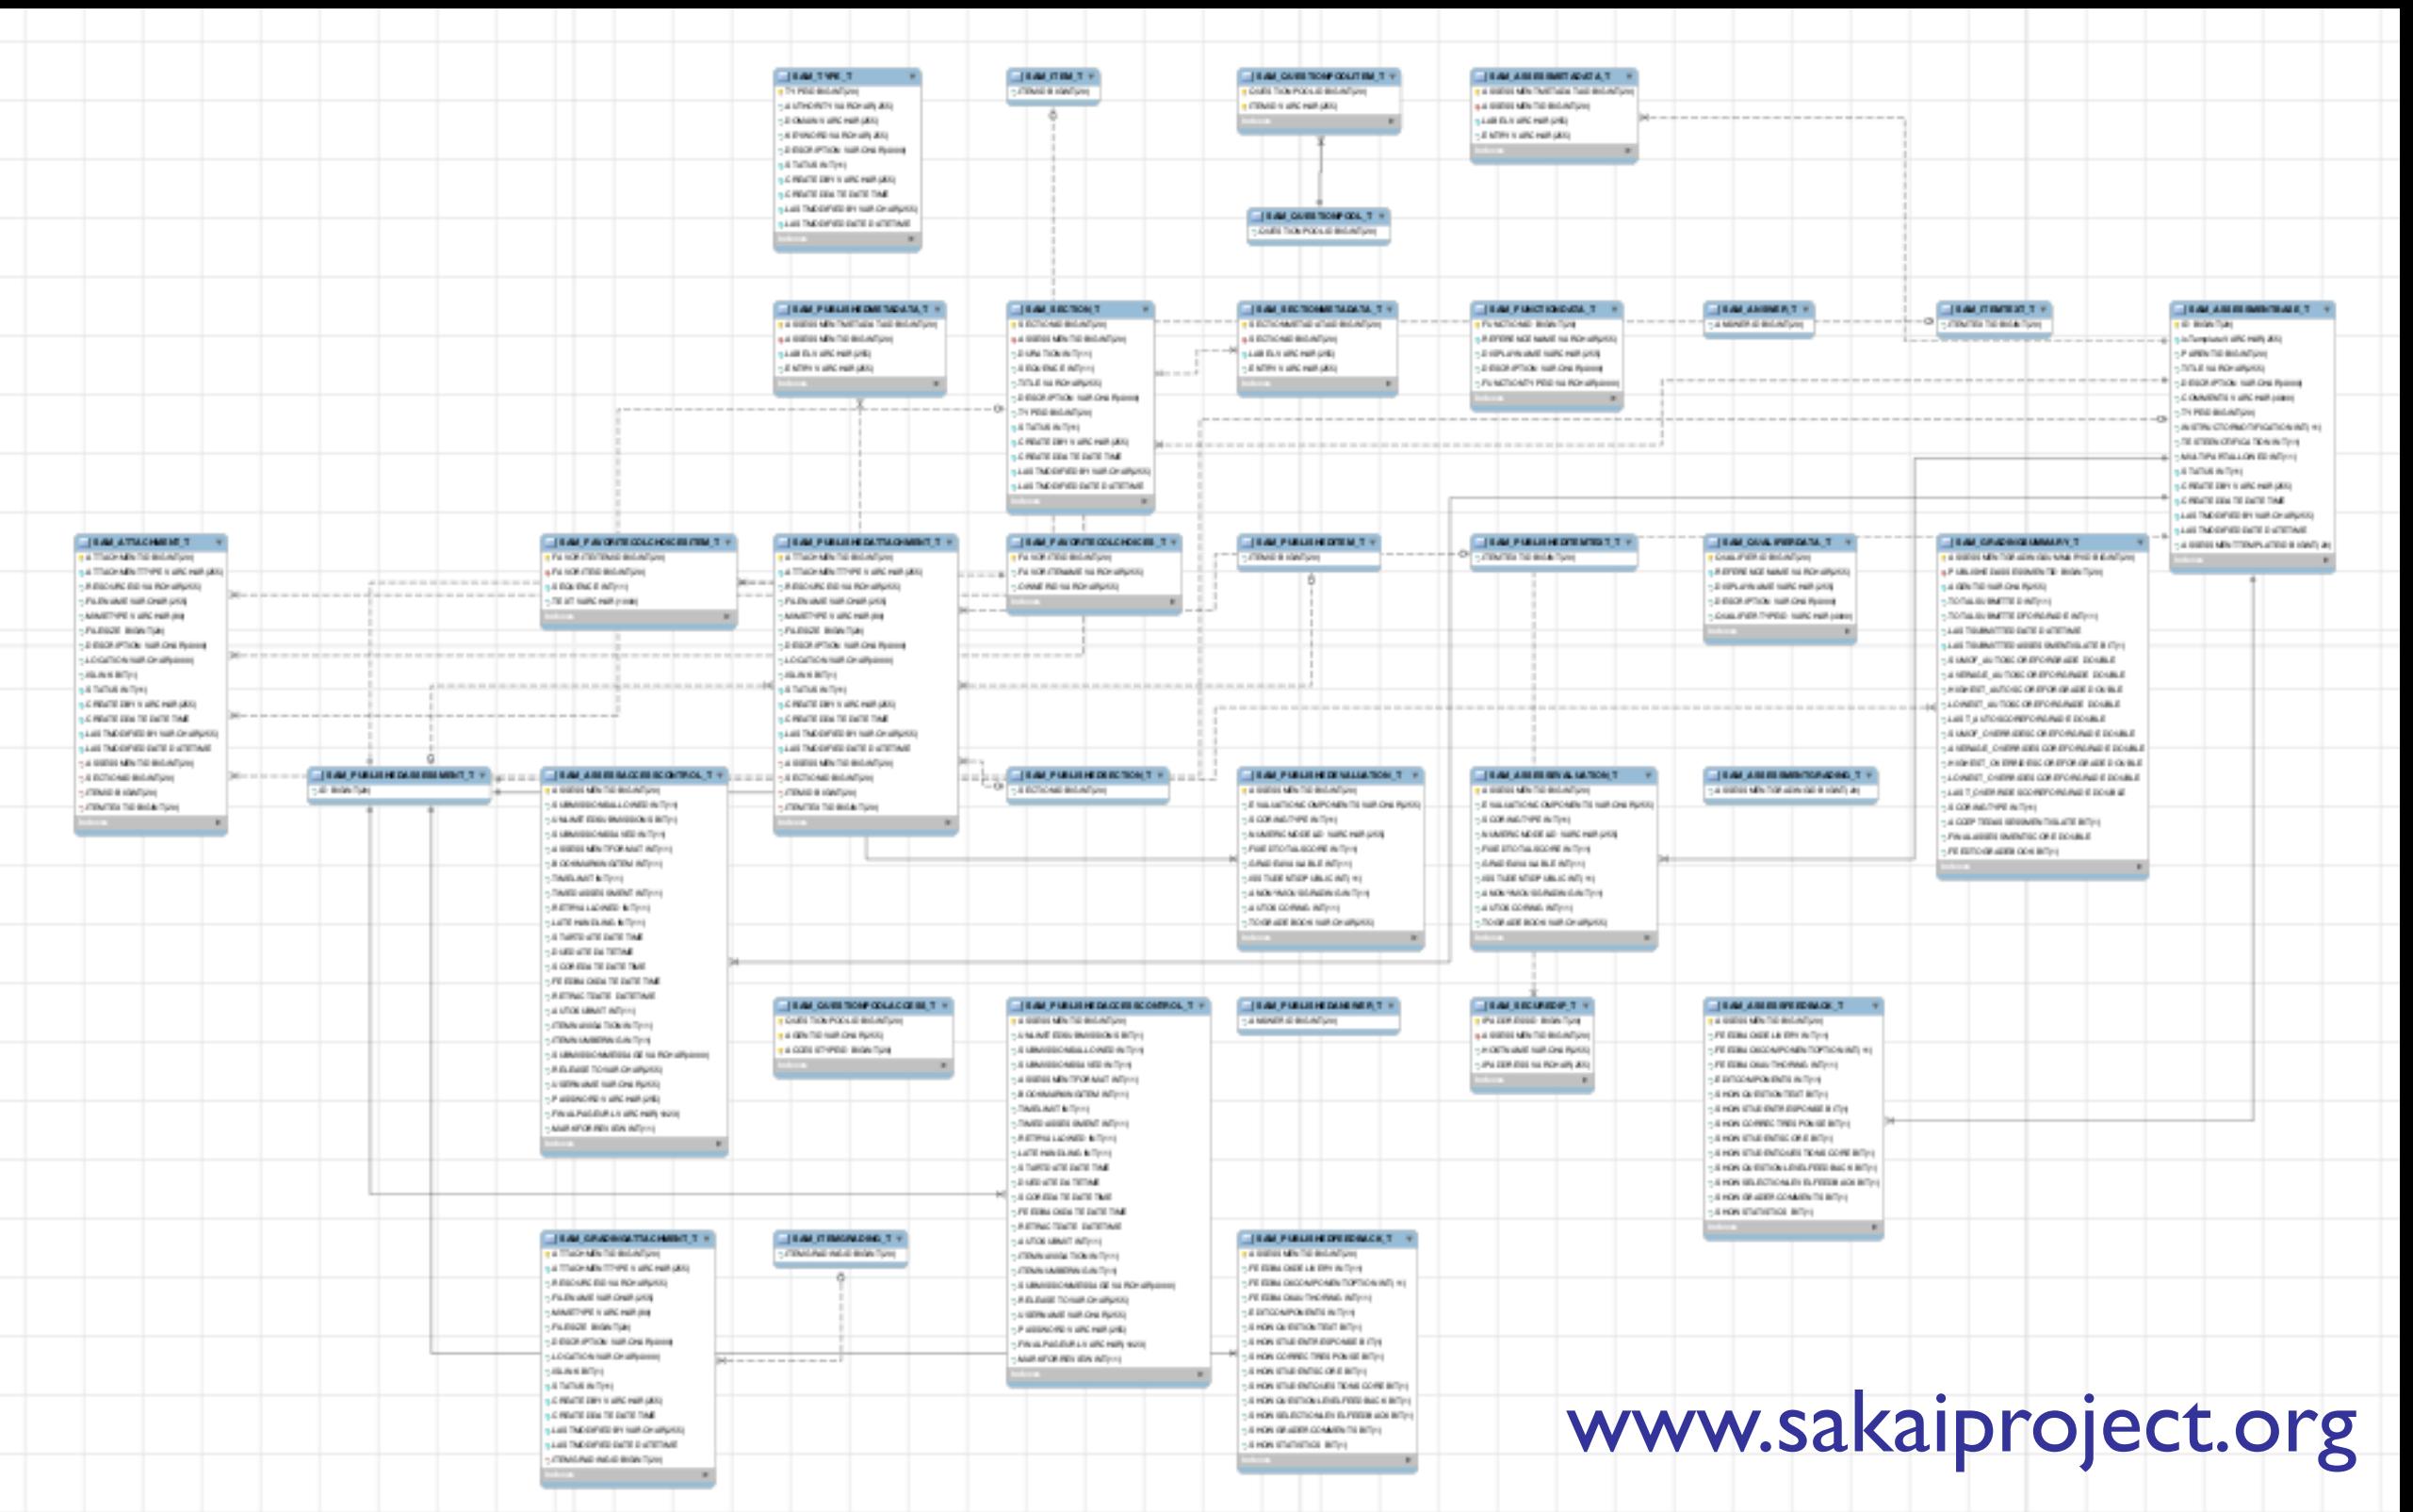
- Database design is an **art form** of its own with particular skills and experience.
- Our goal is to avoid the really bad mistakes and design clean and easily understood databases.
- Others may performance tune things later.
- Database design starts with a picture...



Relational Database Design

WEB APPLICATIONS FOR EVERYBODY

The Michigan Wolverines logo, consisting of a large, bold, yellow block letter 'M'.



Building a Data Model

- Drawing a picture of the data objects for our application and then figuring out how to represent the objects and their relationships
- Basic Rule: Don't put the same string data in twice - use a relationship instead
- When there is one thing in the “real world” there should only be one copy of that thing in the database



Track	Len	Artist	Album	Genre	Rating	Count
✓ Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
✓ Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
✓ For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
✓ Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
✓ Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
✓ Tin Man	3:30	America	Greatest Hits	Easy Listen...	★★★★★	23
✓ Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...	★★★★★	24
✓ Track 01	4:22	Billy Price	Danger Zone	Blues/R&B	★★★★★	26
✓ Track 02	2:45	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
✓ Track 03	3:26	Billy Price	Danger Zone	Blues/R&B	★★★★★	22
✓ Track 04	4:17	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
✓ Track 05	3:50	Billy Price	Danger Zone	Blues/R&B	★★★★★	21
✓ War Pigs/Luke's Wall	7:58	Black Sabbath	Paranoid	Metal	★★★★★	25
✓ Paranoid	2:53	Black Sabbath	Paranoid	Metal	★★★★★	22
✓ Planet Caravan	4:35	Black Sabbath	Paranoid	Metal	★★★★★	25
✓ Iron Man	5:59	Black Sabbath	Paranoid	Metal	★★★★★	26
✓ Electric Funeral	4:53	Black Sabbath	Paranoid	Metal	★★★★★	22
✓ Hand of Doom	7:10	Black Sabbath	Paranoid	Metal	★★★★★	23
✓ Rat Salad	2:30	Black Sabbath	Paranoid	Metal	★★★★★	31
✓ Jack the Stripper/Fairies Wear ...	6:14	Black Sabbath	Paranoid	Metal	★★★★★	24
✓ Bomb Squad (TECH)	3:28	Brent	Brent's Album			1
✓ clay techno	4:36	Brent	Brent's Album			2
✓ Heavy	3:08	Brent	Brent's Album			1
✓ Hi metal man	4:20	Brent	Brent's Album			1
✓ Mistro	2:58	Brent	Brent's Album			1

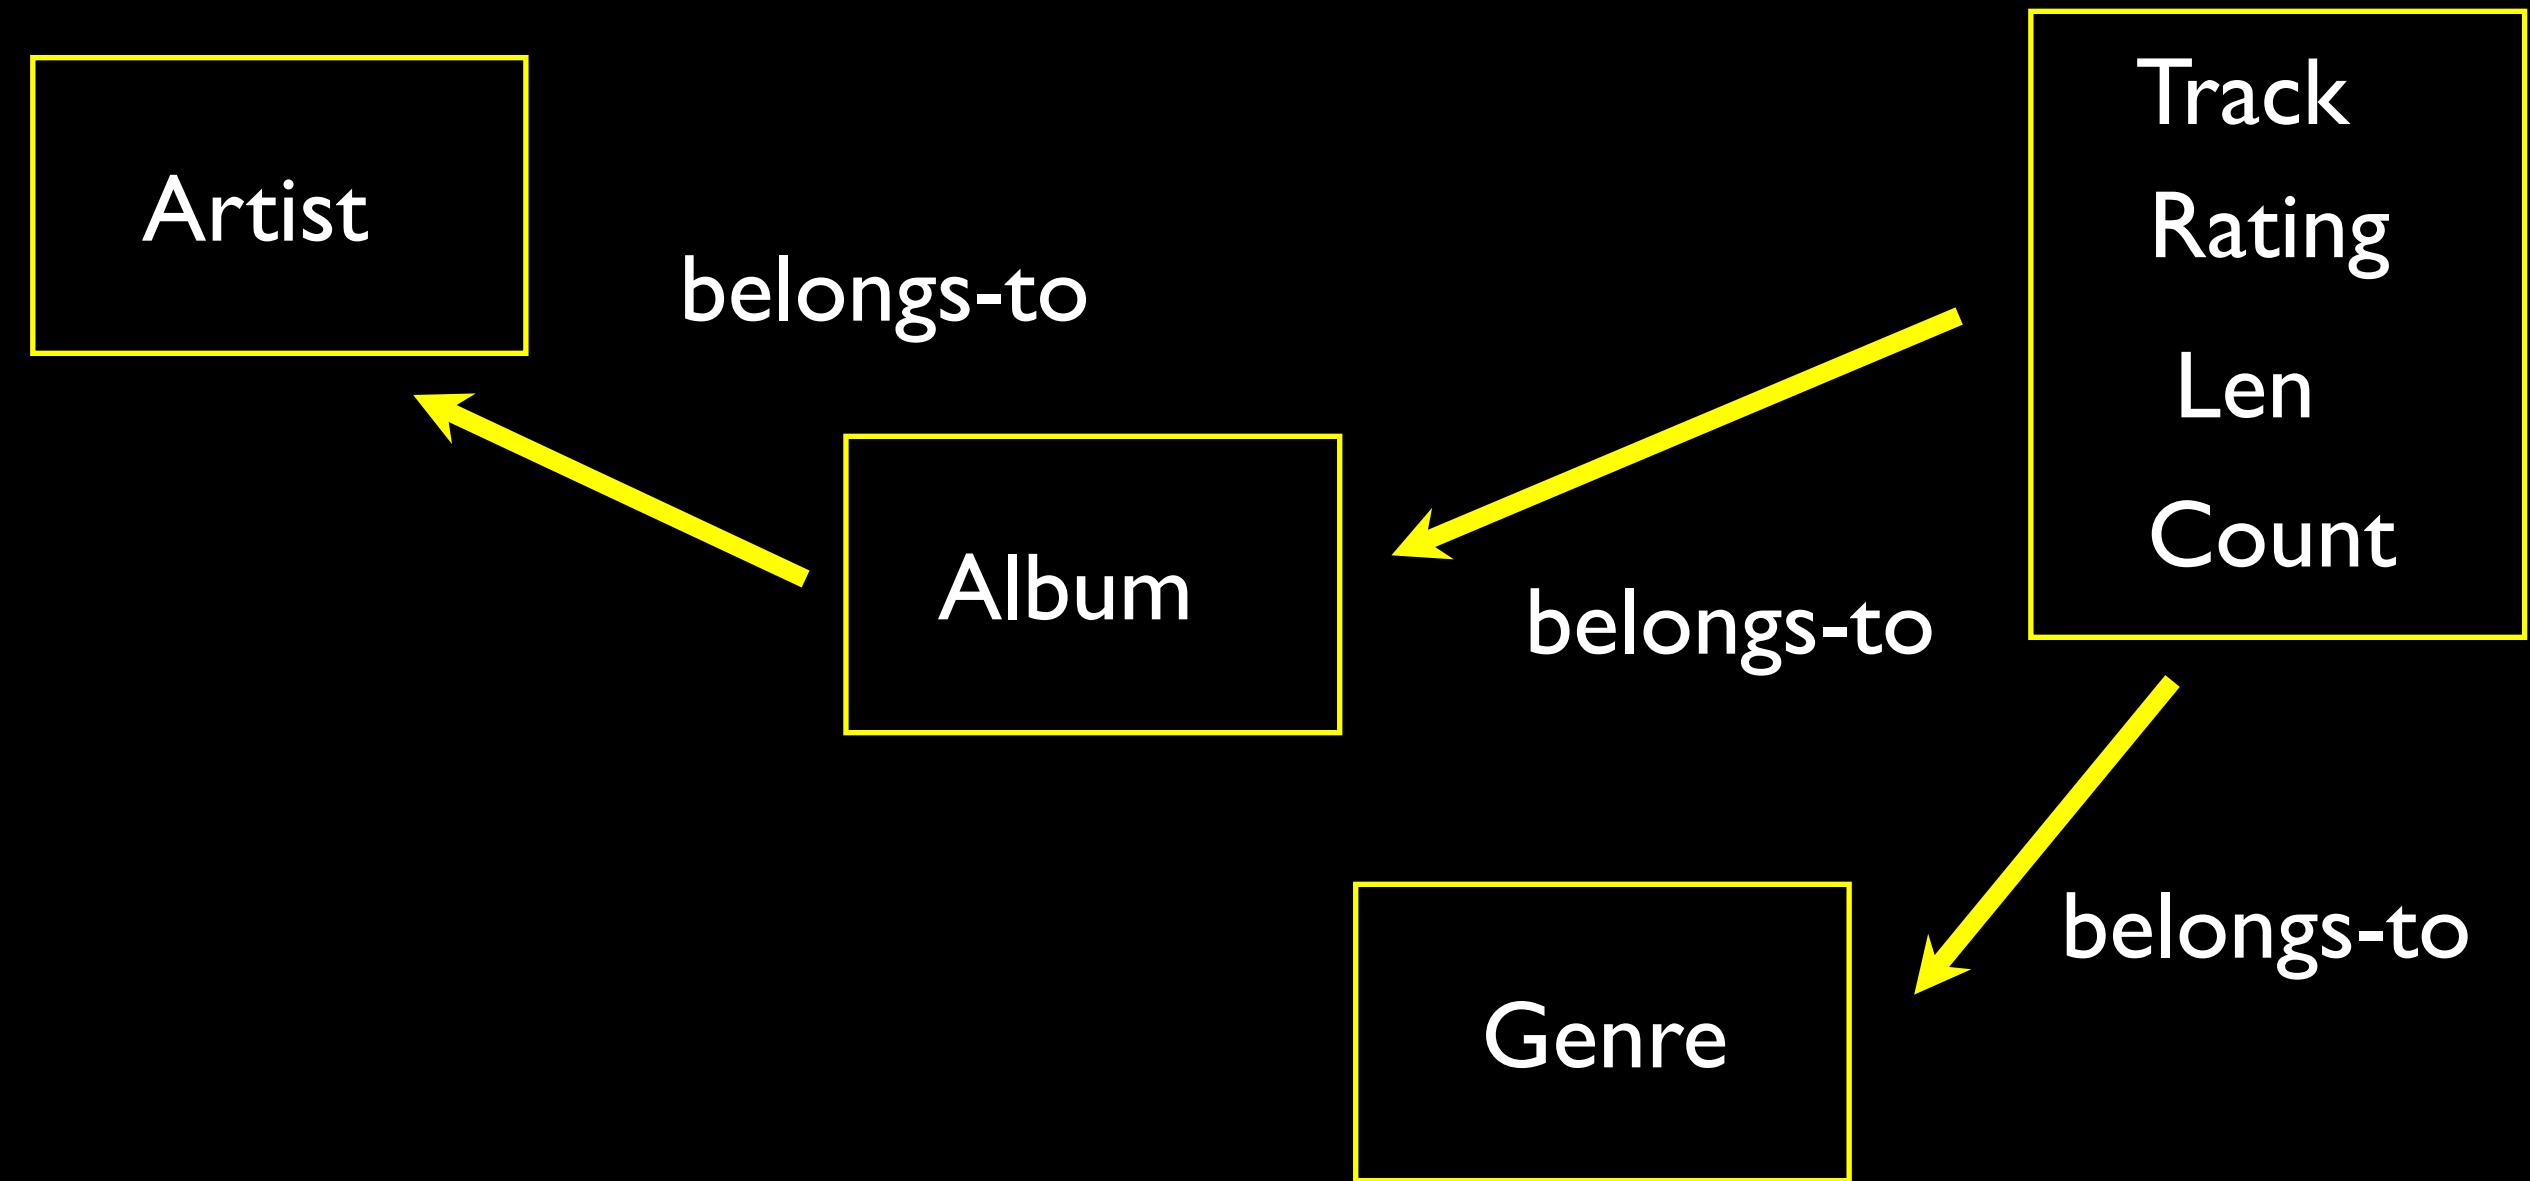
For each “piece of info”...

- Is the column an object or an attribute of another object?
- Once we define objects, we need to define the relationships between objects.

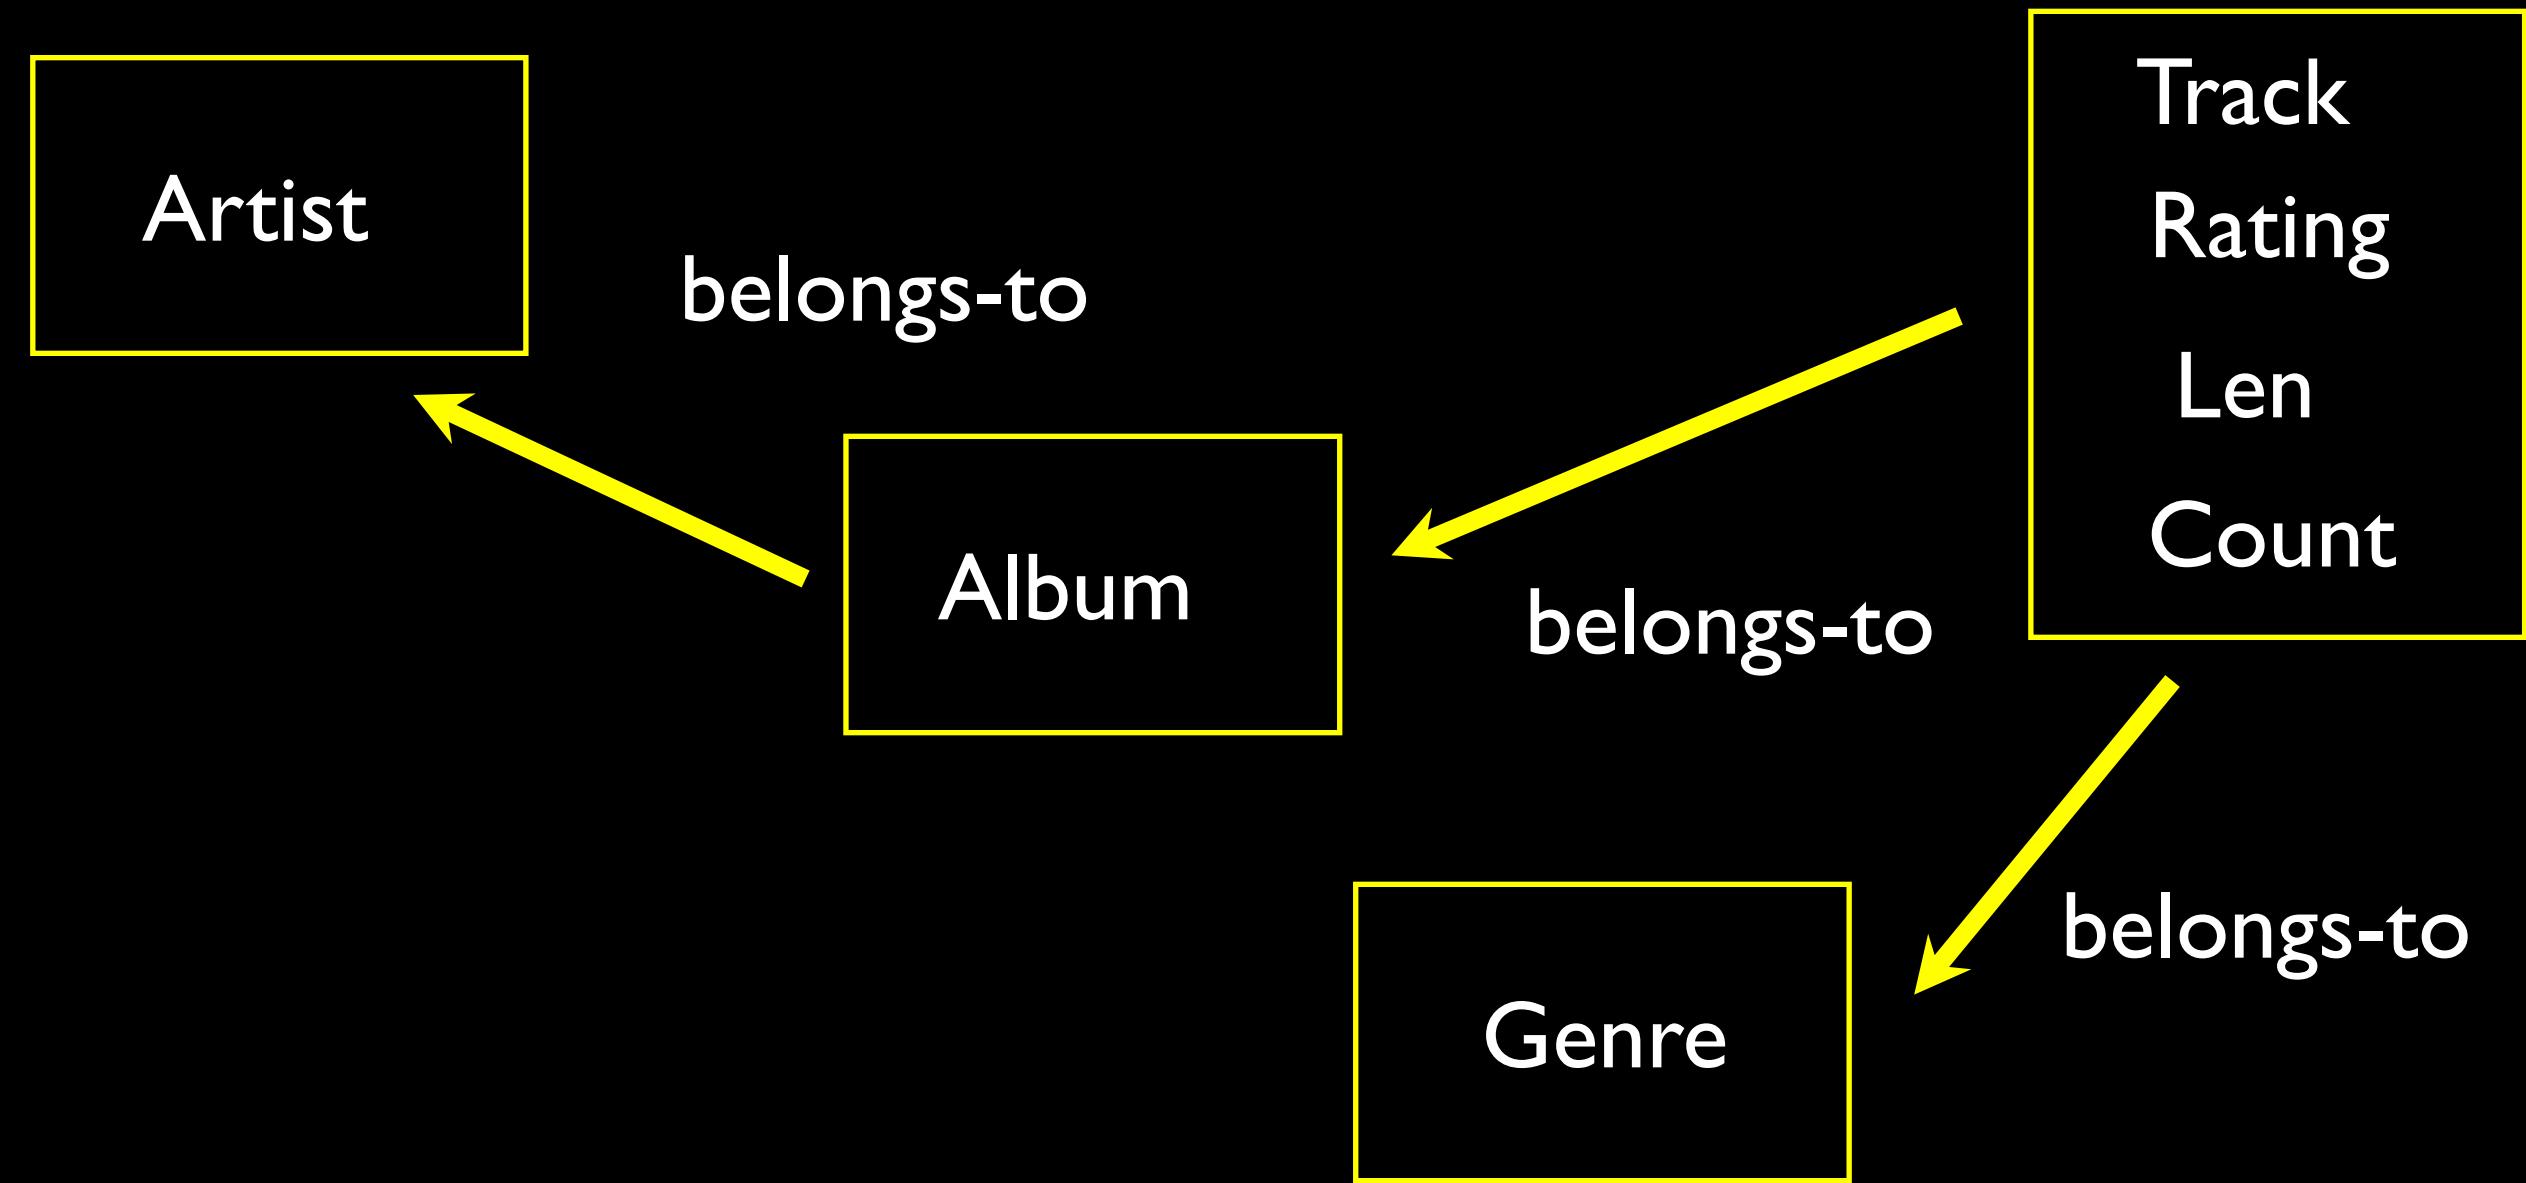
Len	Album	Genre	Artist	Rating	Track	Count
-----	-------	-------	--------	--------	-------	-------

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tie Me	2:20	America	Greatest Hits	Easy Listen...	★★★★★	22

Track
Album
Artist
Genre
Rating
Len
Count



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tie Man	7:20	America	Greatest Hits	Easy Listen...	★★★★★	22



✓ Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
✓ Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
✓ For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
✓ Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
✓ Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
✓ Tie Man	7:20	America	Greatest Hits	Easy Listen...	★★★★★	22



Normalization and Foreign Keys



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input checked="" type="checkbox"/> Tip Man	2:30	America	Greatest Hits	Easy Listen...		22

We want to keep track of which band is the “**creator**” of each music track...
What album does this song “belong to”?

Which album is this song related to?

Database Normalization (3NF)

There is *tons* of database theory - way too much to understand without excessive predicate calculus

- Do not replicate data. Instead, reference data. Point at data.
- Use integers for keys and for references.
- Add a special “key” column to each table, which you will make references to.

http://en.wikipedia.org/wiki/Database_normalization

Integer Reference Pattern

We use integer columns in one table to reference (or look up) rows in another table.

Artist

	artist_id	name
	1	Led Zeppelin
	2	AC/DC

Album

	album_id	title	artist_id
	1	Who Made Who	2
	2	IV	1



Key Terminology

Finding our way around....

Three Kinds of Keys

- **Primary key** - generally an integer auto-increment field
- **Logical key** - what the outside world uses for lookup
- **Foreign key** - generally an integer key pointing to a row in another table



Primary Key Rules

Best practices:

- Never use your **logical key** as the **primary key**.
- **Logical keys** can and do change, albeit slowly.
- **Relationships** that are based on matching string fields are less efficient than integers.

User

user_id

email

password

name

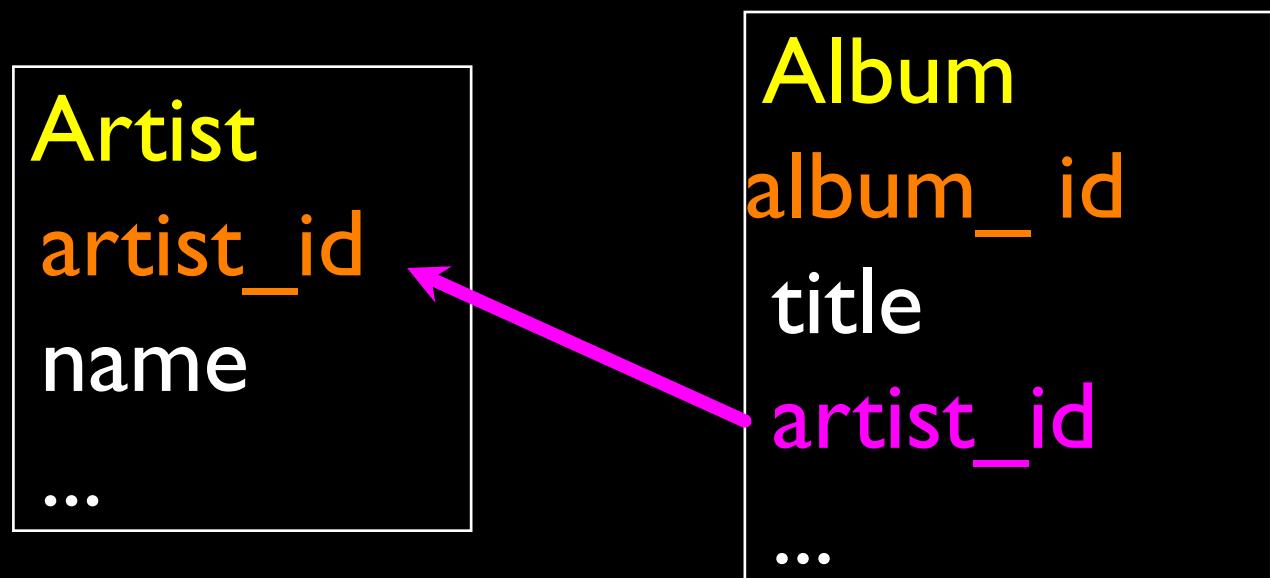
created_at

modified_at

login_at

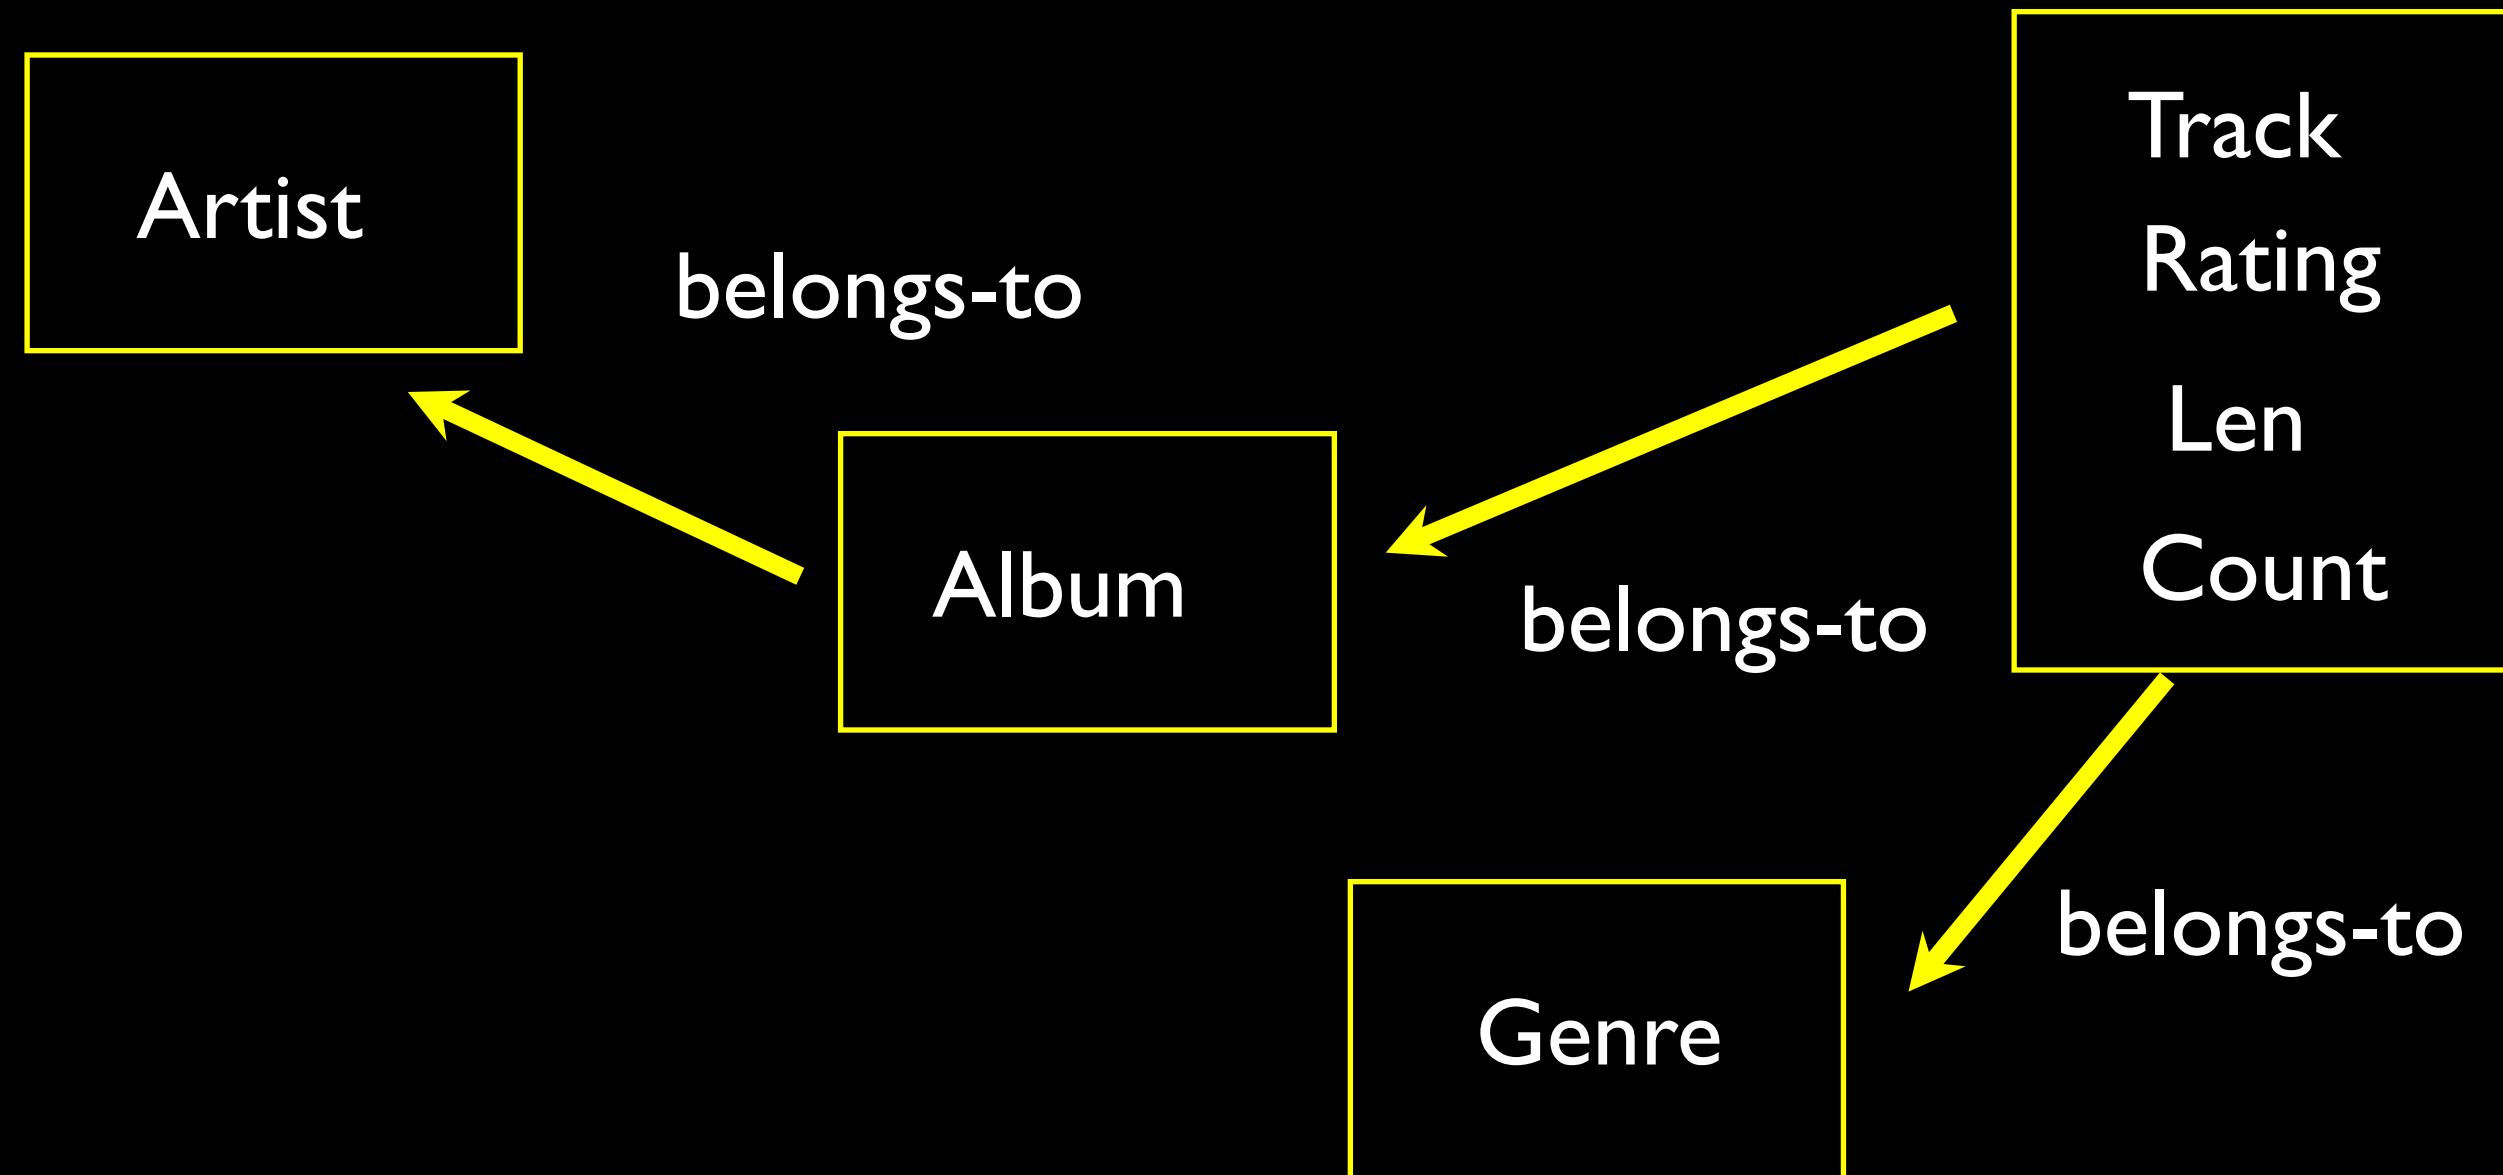
Foreign Keys

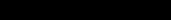
- A **foreign key** is when a table has a column containing a key that points to the **primary key** of another table.
- When all primary keys are integers, then all foreign keys are integers. This is good - very good.

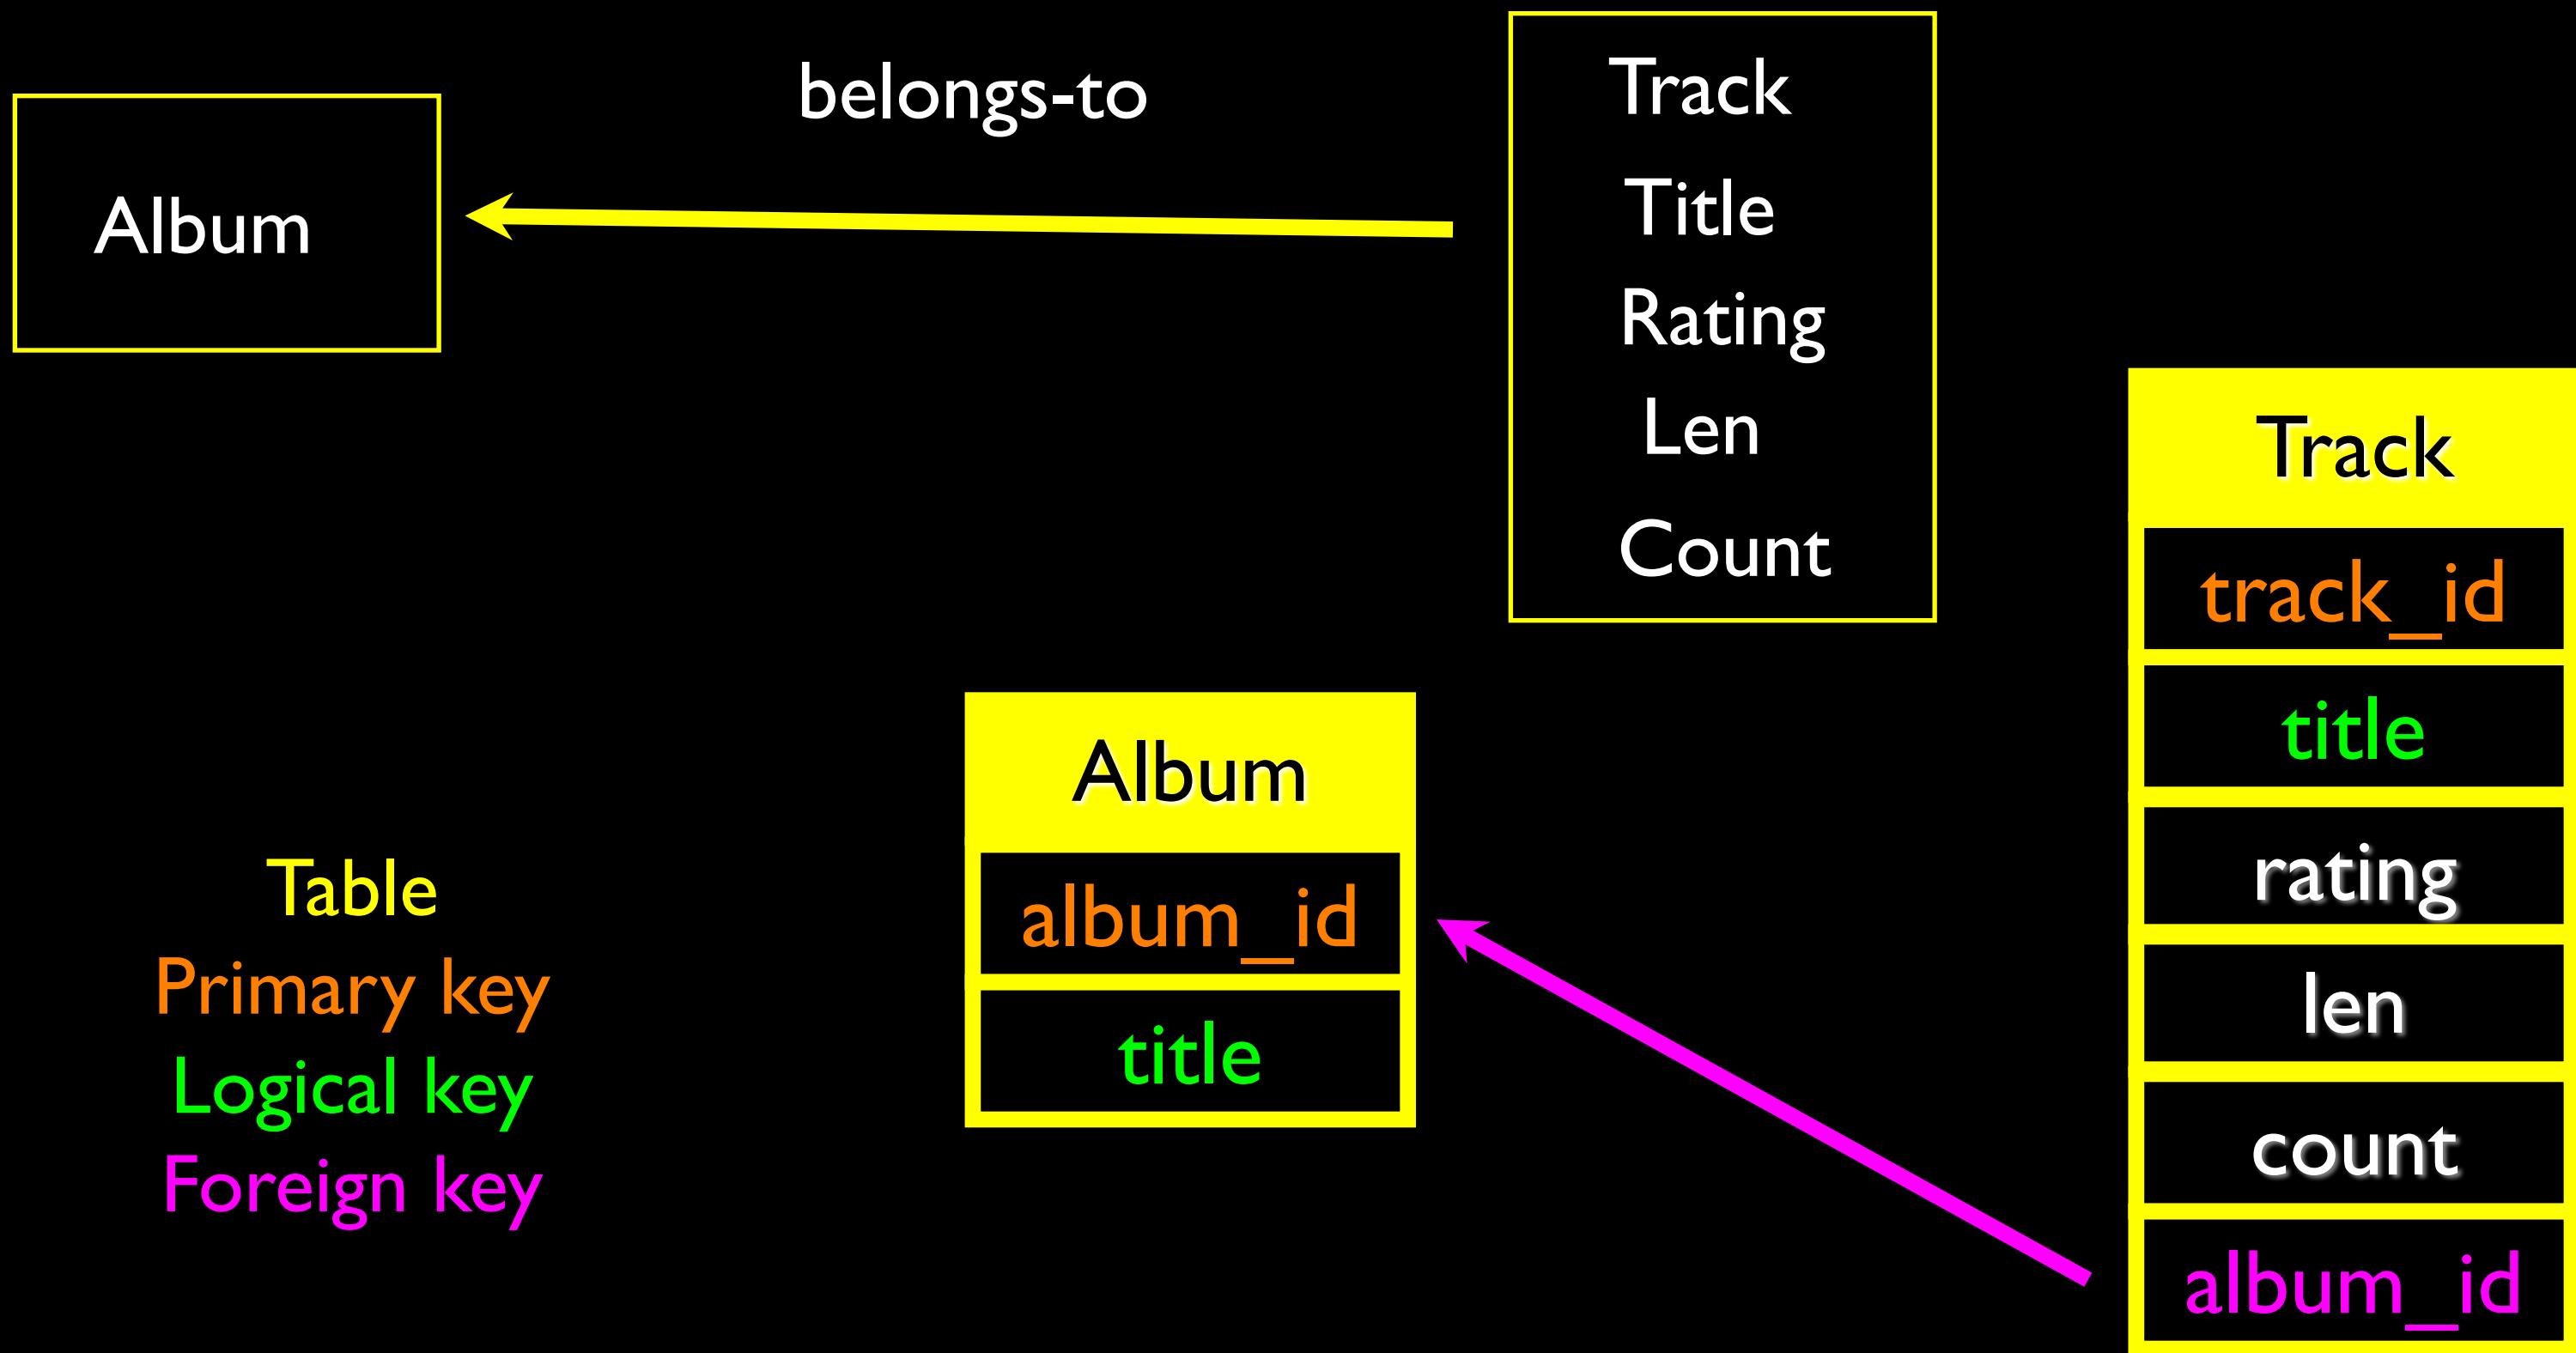


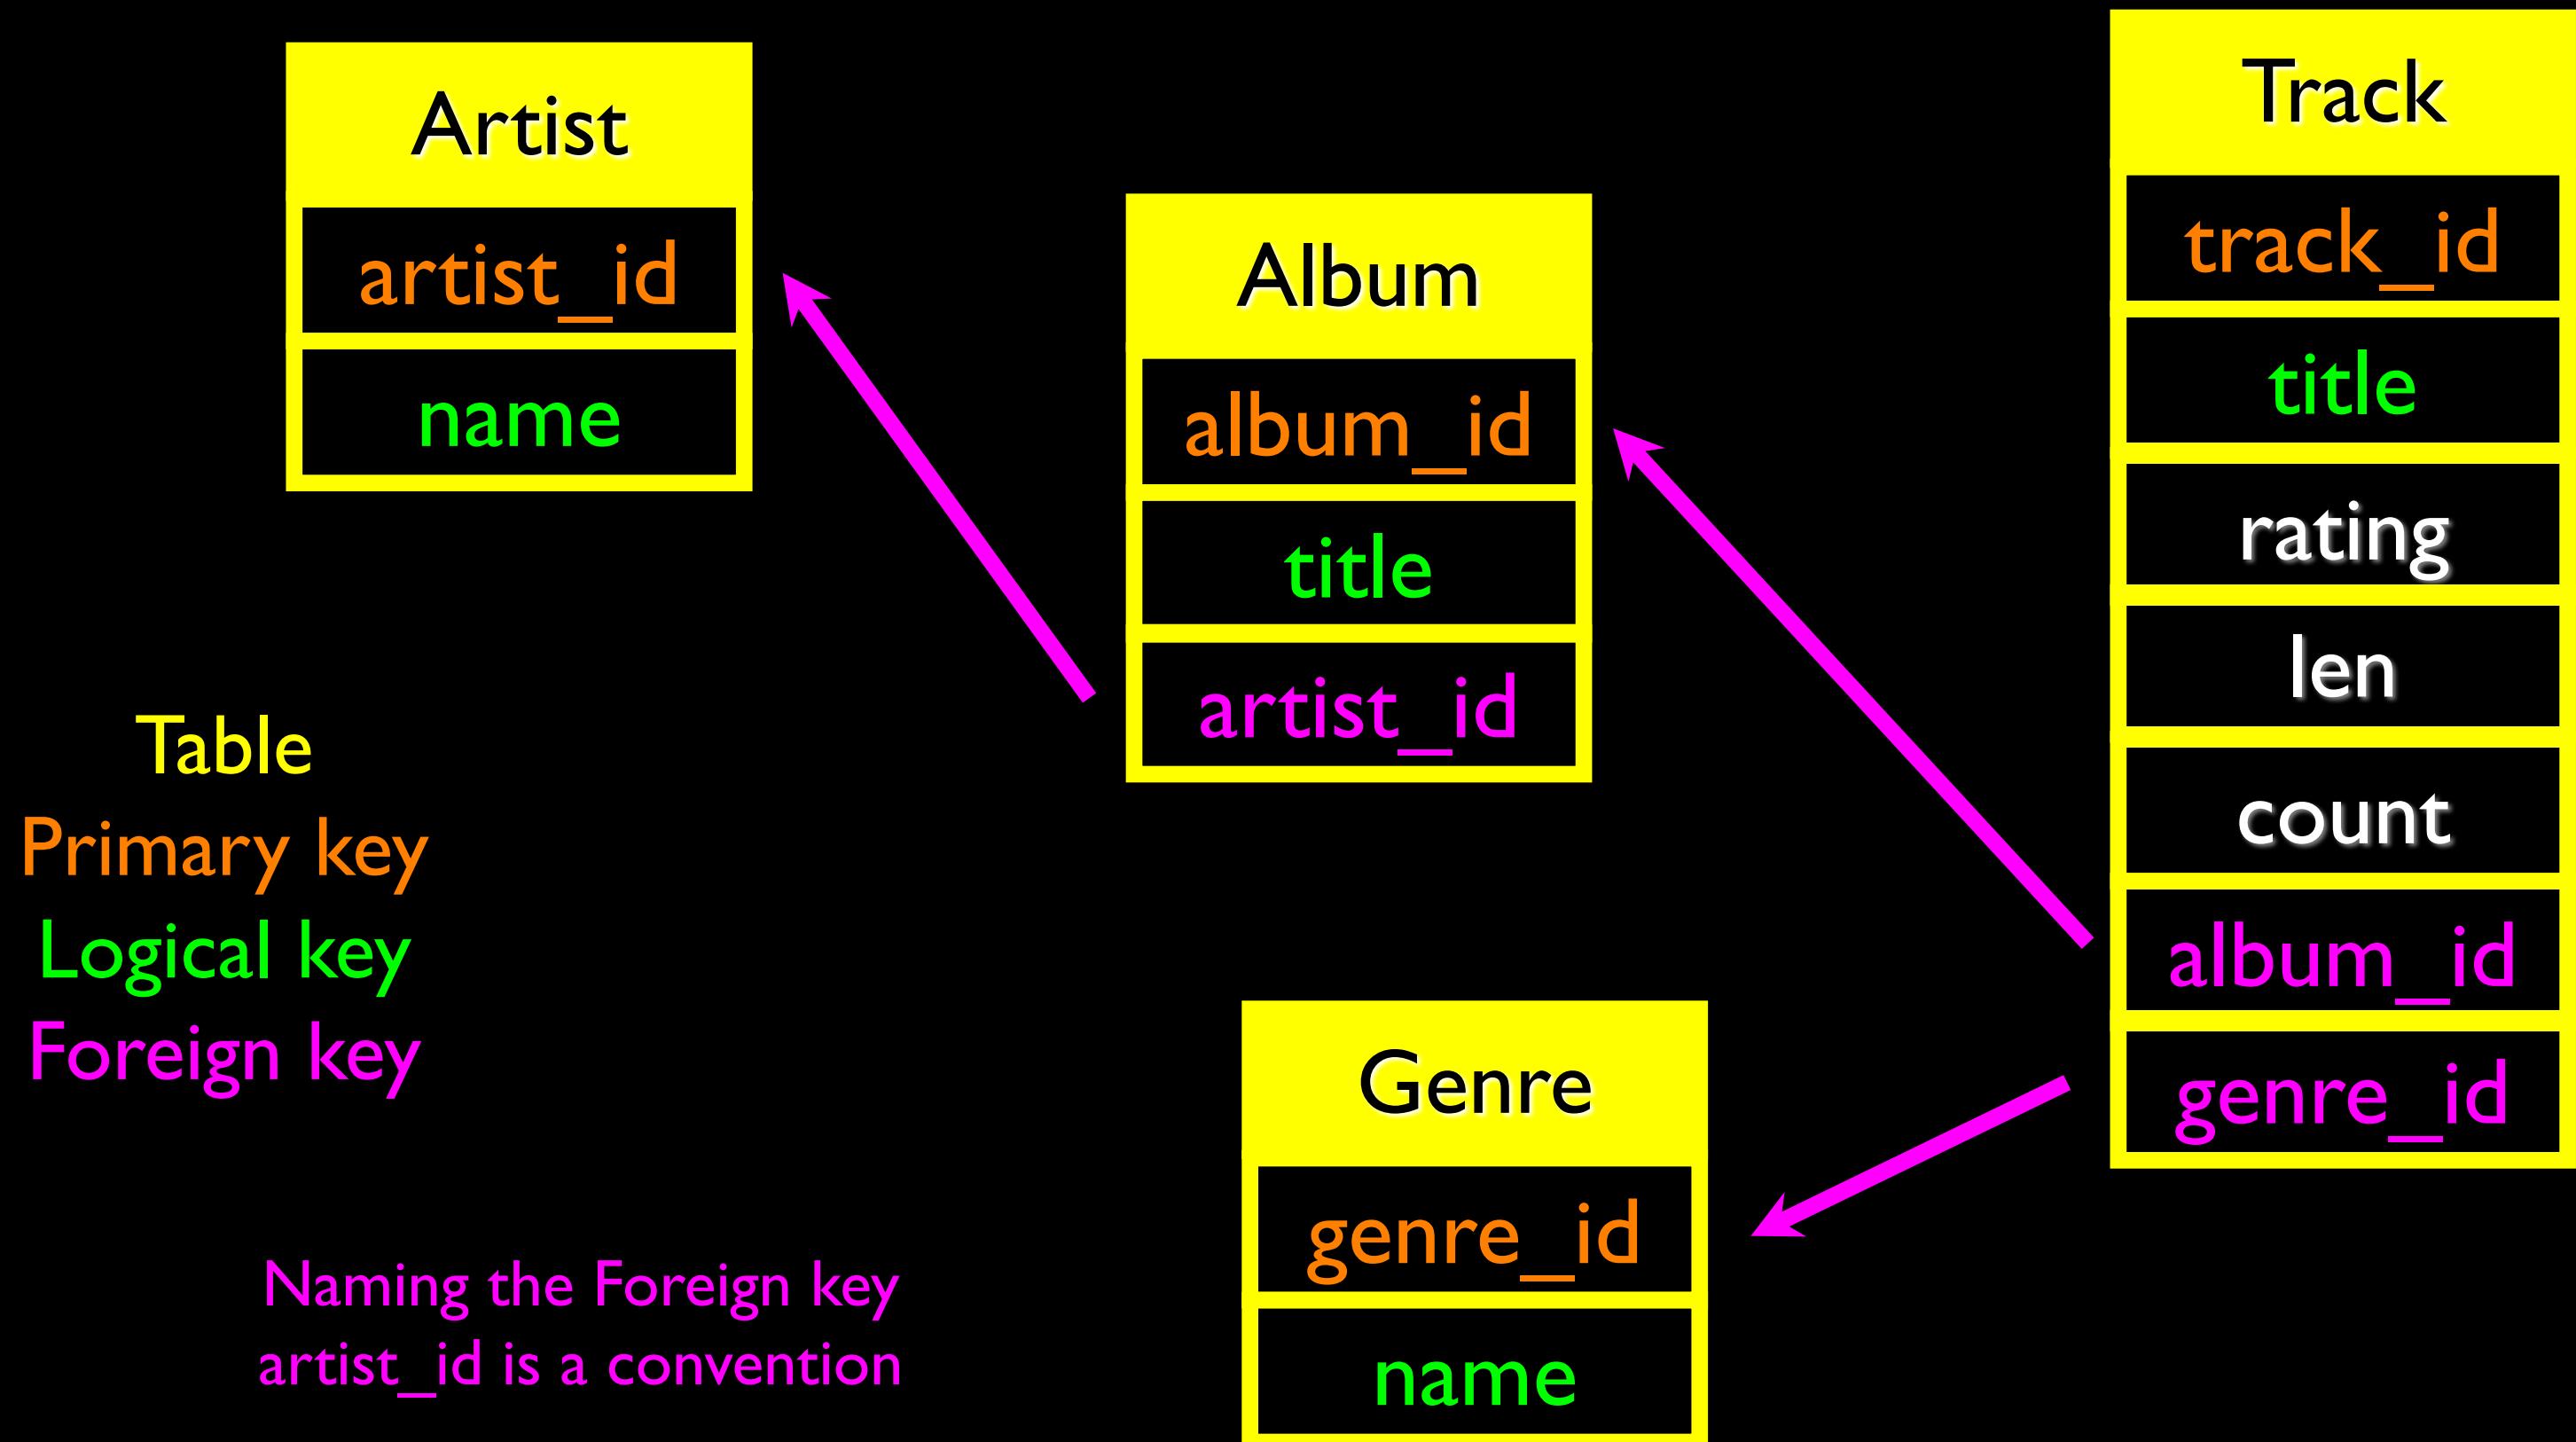


Building a Physical Data Schema



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input checked="" type="checkbox"/> Tie Me	7:20	America	Greatest Hits	Easy Listen...		22





Creating our Music Database

```
CREATE DATABASE Music  
DEFAULT CHARACTER SET utf8;
```

```
USE Music;
```

```
CREATE TABLE Artist (
    artist_id INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(255)
) ENGINE = InnoDB;
```

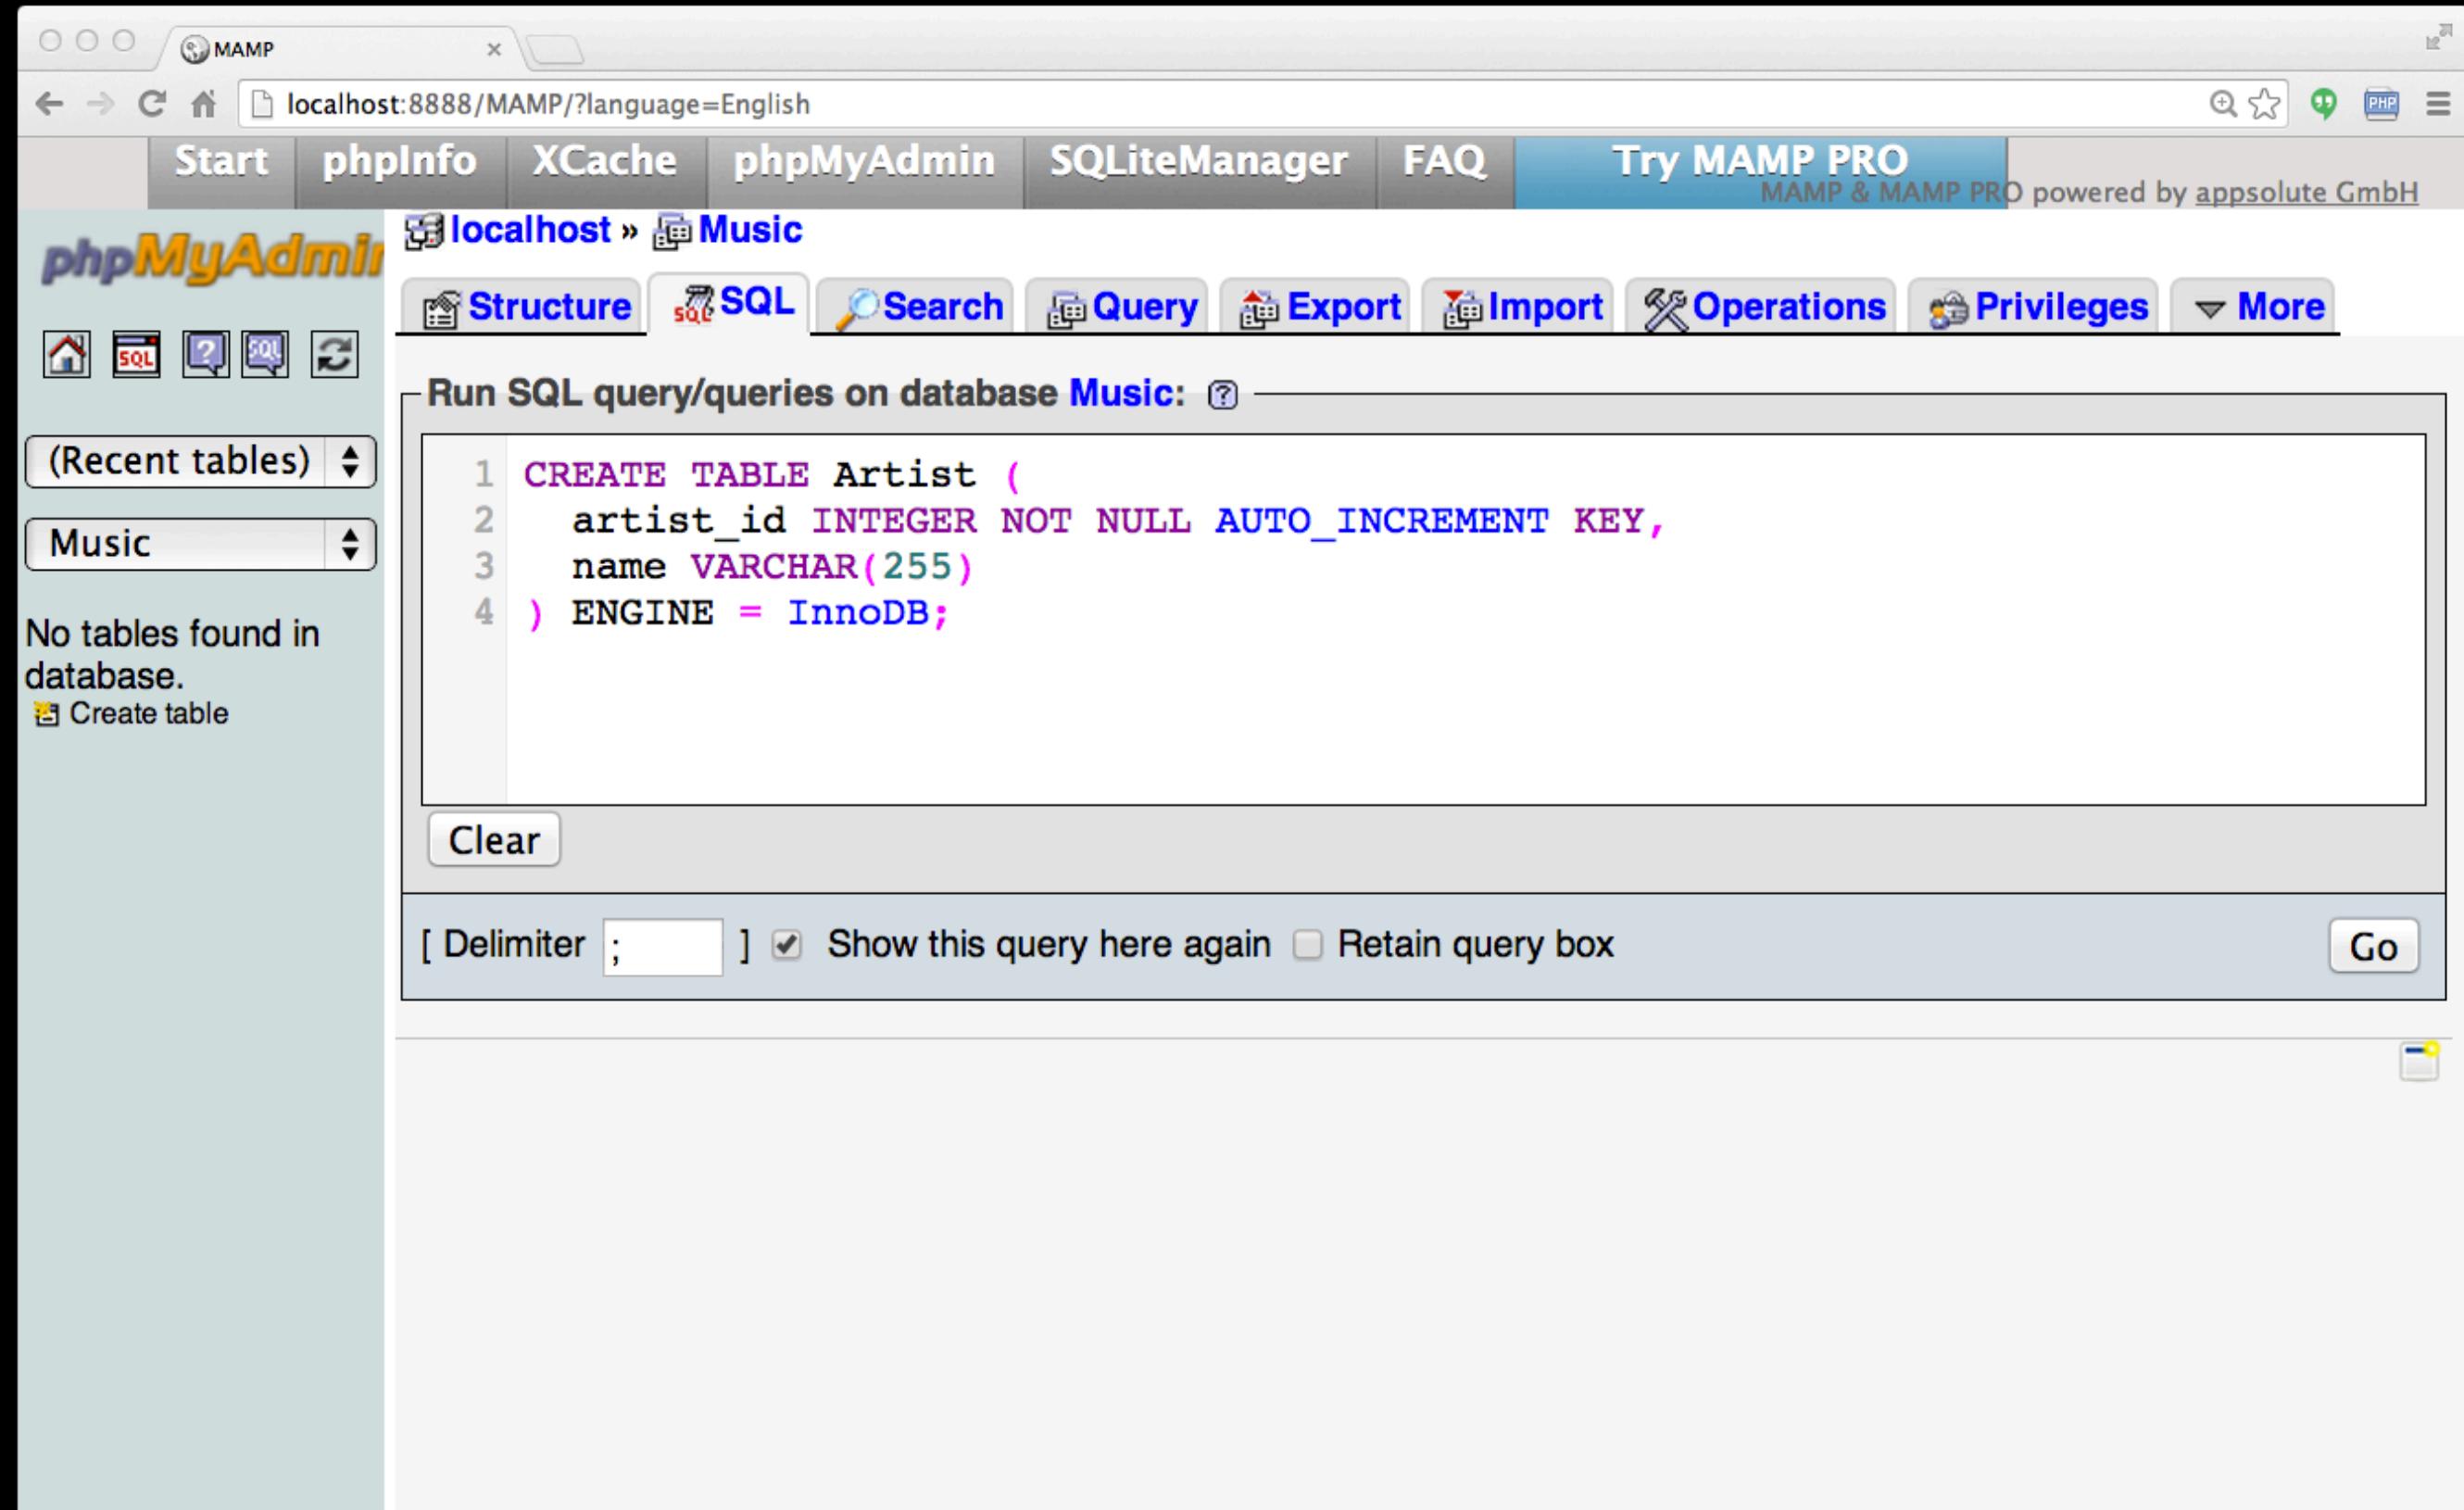
```
CREATE TABLE Album (
    album_id INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    title VARCHAR(255),
    artist_id INTEGER,
    INDEX USING BTREE (title),
    CONSTRAINT FOREIGN KEY (artist_id)
        REFERENCES Artist (artist_id)
        ON DELETE CASCADE ON UPDATE CASCADE
) ENGINE = InnoDB;
```



```
CREATE TABLE Genre (
    genre_id INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(255),
    INDEX USING BTREE (name)
) ENGINE = InnoDB;
```

```
CREATE TABLE Track (
    track_id INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    title VARCHAR(255),
    len INTEGER,
    rating INTEGER,
    count INTEGER,
    album_id INTEGER,
    genre_id INTEGER,
    INDEX USING BTREE (title),

    CONSTRAINT FOREIGN KEY (album_id) REFERENCES Album (album_id)
        ON DELETE CASCADE ON UPDATE CASCADE,
    CONSTRAINT FOREIGN KEY (genre_id) REFERENCES Genre (genre_id)
        ON DELETE CASCADE ON UPDATE CASCADE
) ENGINE = InnoDB;
```



The screenshot shows the phpMyAdmin interface running on a local MAMP server. The database selected is 'Music'. In the SQL tab, the following SQL code is entered:

```
1 CREATE TABLE Artist (
2     artist_id INTEGER NOT NULL AUTO_INCREMENT KEY,
3     name VARCHAR(255)
4 ) ENGINE = InnoDB;
```

The 'Delimiter' field contains a semicolon (;). The 'Show this query here again' checkbox is checked, while 'Retain query box' is unchecked. A 'Go' button is visible at the bottom right.

MAMP

localhost:8888/MAMP/?language=English

Start phpInfo XCache phpMyAdmin SQLiteManager FAQ Try MAMP PRO MAMP & MAMP PRO powered by apposite GmbH

localhost » Music » Artist

Browse Structure SQL Search Insert Export Import Operations Triggers

#	Name	Type	Collation	Attributes	Null	Default	Extra	Action
1	artist_id	int(11)			No	None	AUTO_INCREMENT	   
2	name	varchar(255)	utf8_general_ci		Yes	NULL		   

Check All / Uncheck All With selected:      

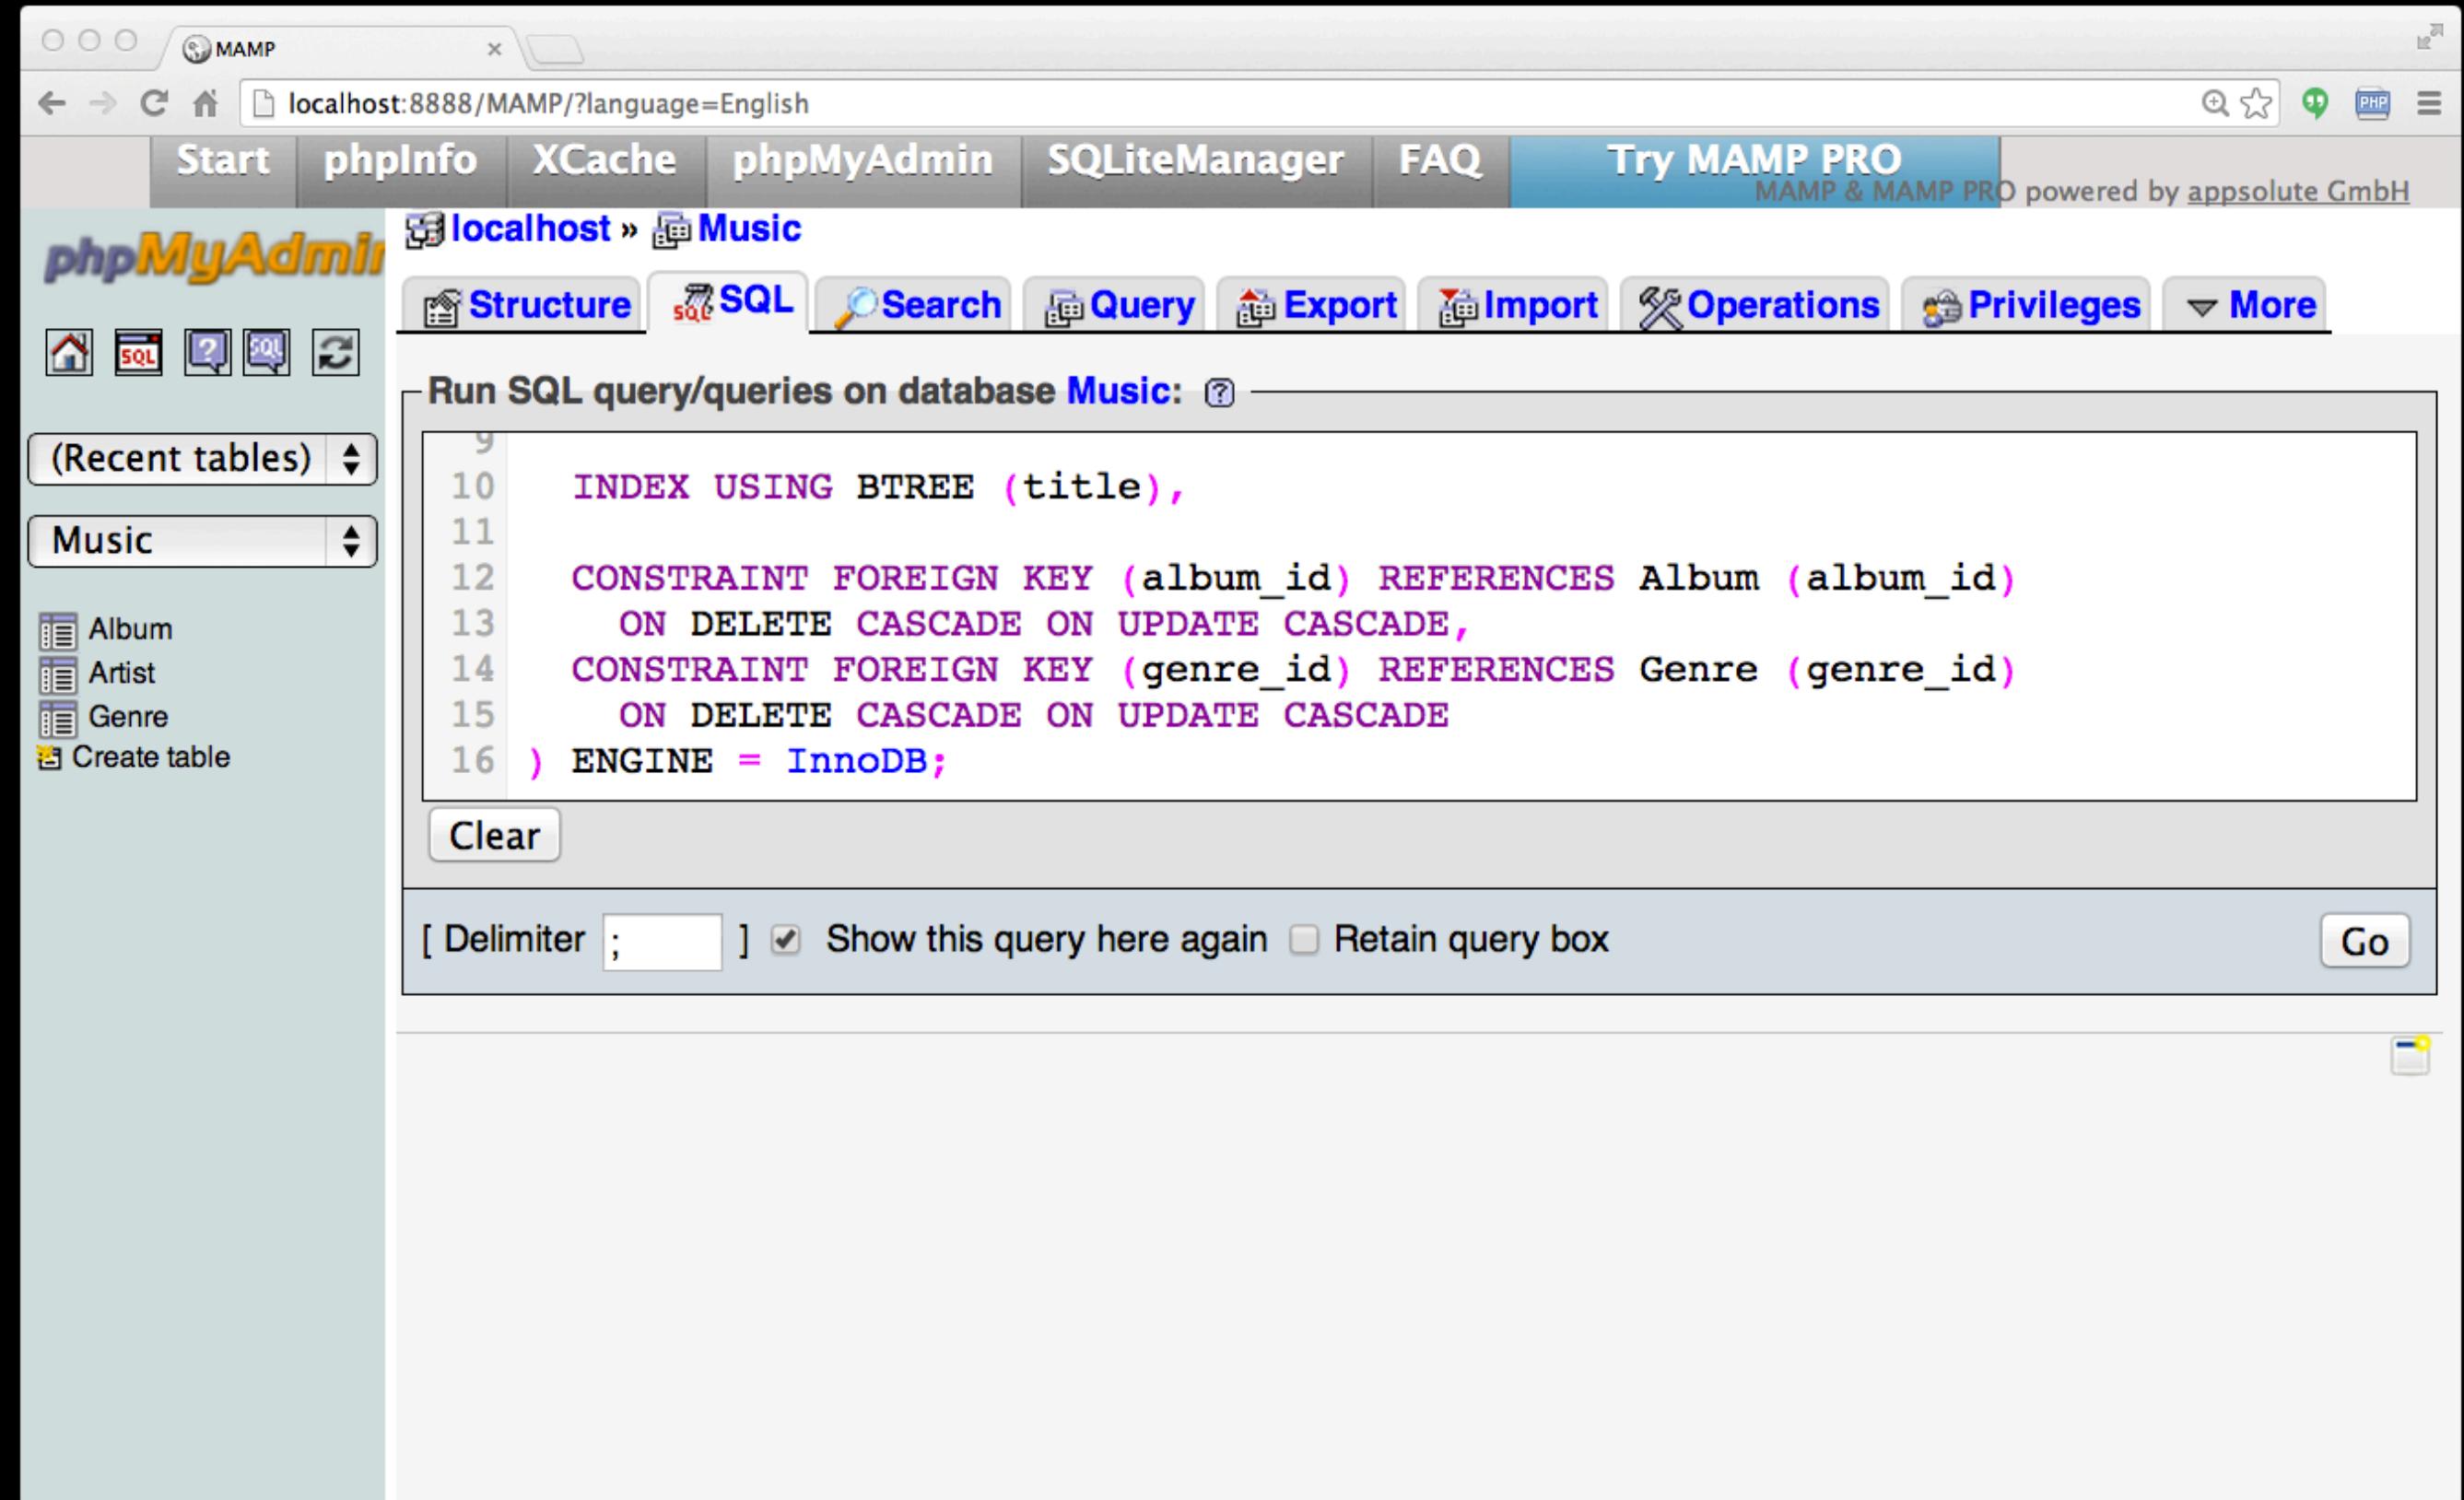
Print view Relation view Propose table structure 

Add 1 column(s) At End of Table At Beginning of Table After artist_id 

+ Indexes

Information

Space usage		Row Statistics	
Type	Usage	Statements	Value
Data	16 KiB	Format	Compact
Index	0 B	Collation	utf8_general_ci
Total	16 KiB	Next autoindex	1
		Creation	Jan 26, 2014 at 09:31 PM



The screenshot shows the phpMyAdmin interface running on a local MAMP server. The database selected is 'Music'. The SQL tab is active, displaying the following SQL code:

```
9 INDEX USING BTREE (title),
10
11
12 CONSTRAINT FOREIGN KEY (album_id) REFERENCES Album (album_id)
13     ON DELETE CASCADE ON UPDATE CASCADE,
14 CONSTRAINT FOREIGN KEY (genre_id) REFERENCES Genre (genre_id)
15     ON DELETE CASCADE ON UPDATE CASCADE
16 ) ENGINE = InnoDB;
```

The code creates an index on the 'title' column of the table, adds foreign key constraints for 'album_id' referencing 'Album.album_id' and 'genre_id' referencing 'Genre.genre_id', and specifies InnoDB as the engine.

Below the SQL editor, there are options for 'Clear', 'Delimiter ;', and checkboxes for 'Show this query here again' and 'Retain query box'. A 'Go' button is also present.

MAMP

localhost:8888/MAMP/?language=English

Start phpInfo XCache phpMyAdmin SQLiteManager FAQ Try MAMP PRO MAMP & MAMP PRO powered by apposite GmbH

localhost » Music » Track

Browse Structure SQL Search Insert Export Import Operations Triggers

#	Name	Type	Collation	Attributes	Null	Default	Extra				
1	track_id	int(11)			No	None	AUTO_INCREMENT				
2	title	varchar(255)	utf8_general_ci		Yes	NULL					
3	len	int(11)			Yes	NULL					
4	rating	int(11)			Yes	NULL					
5	count	int(11)			Yes	NULL					
6	album_id	int(11)			Yes	NULL					
7	genre_id	int(11)			Yes	NULL					

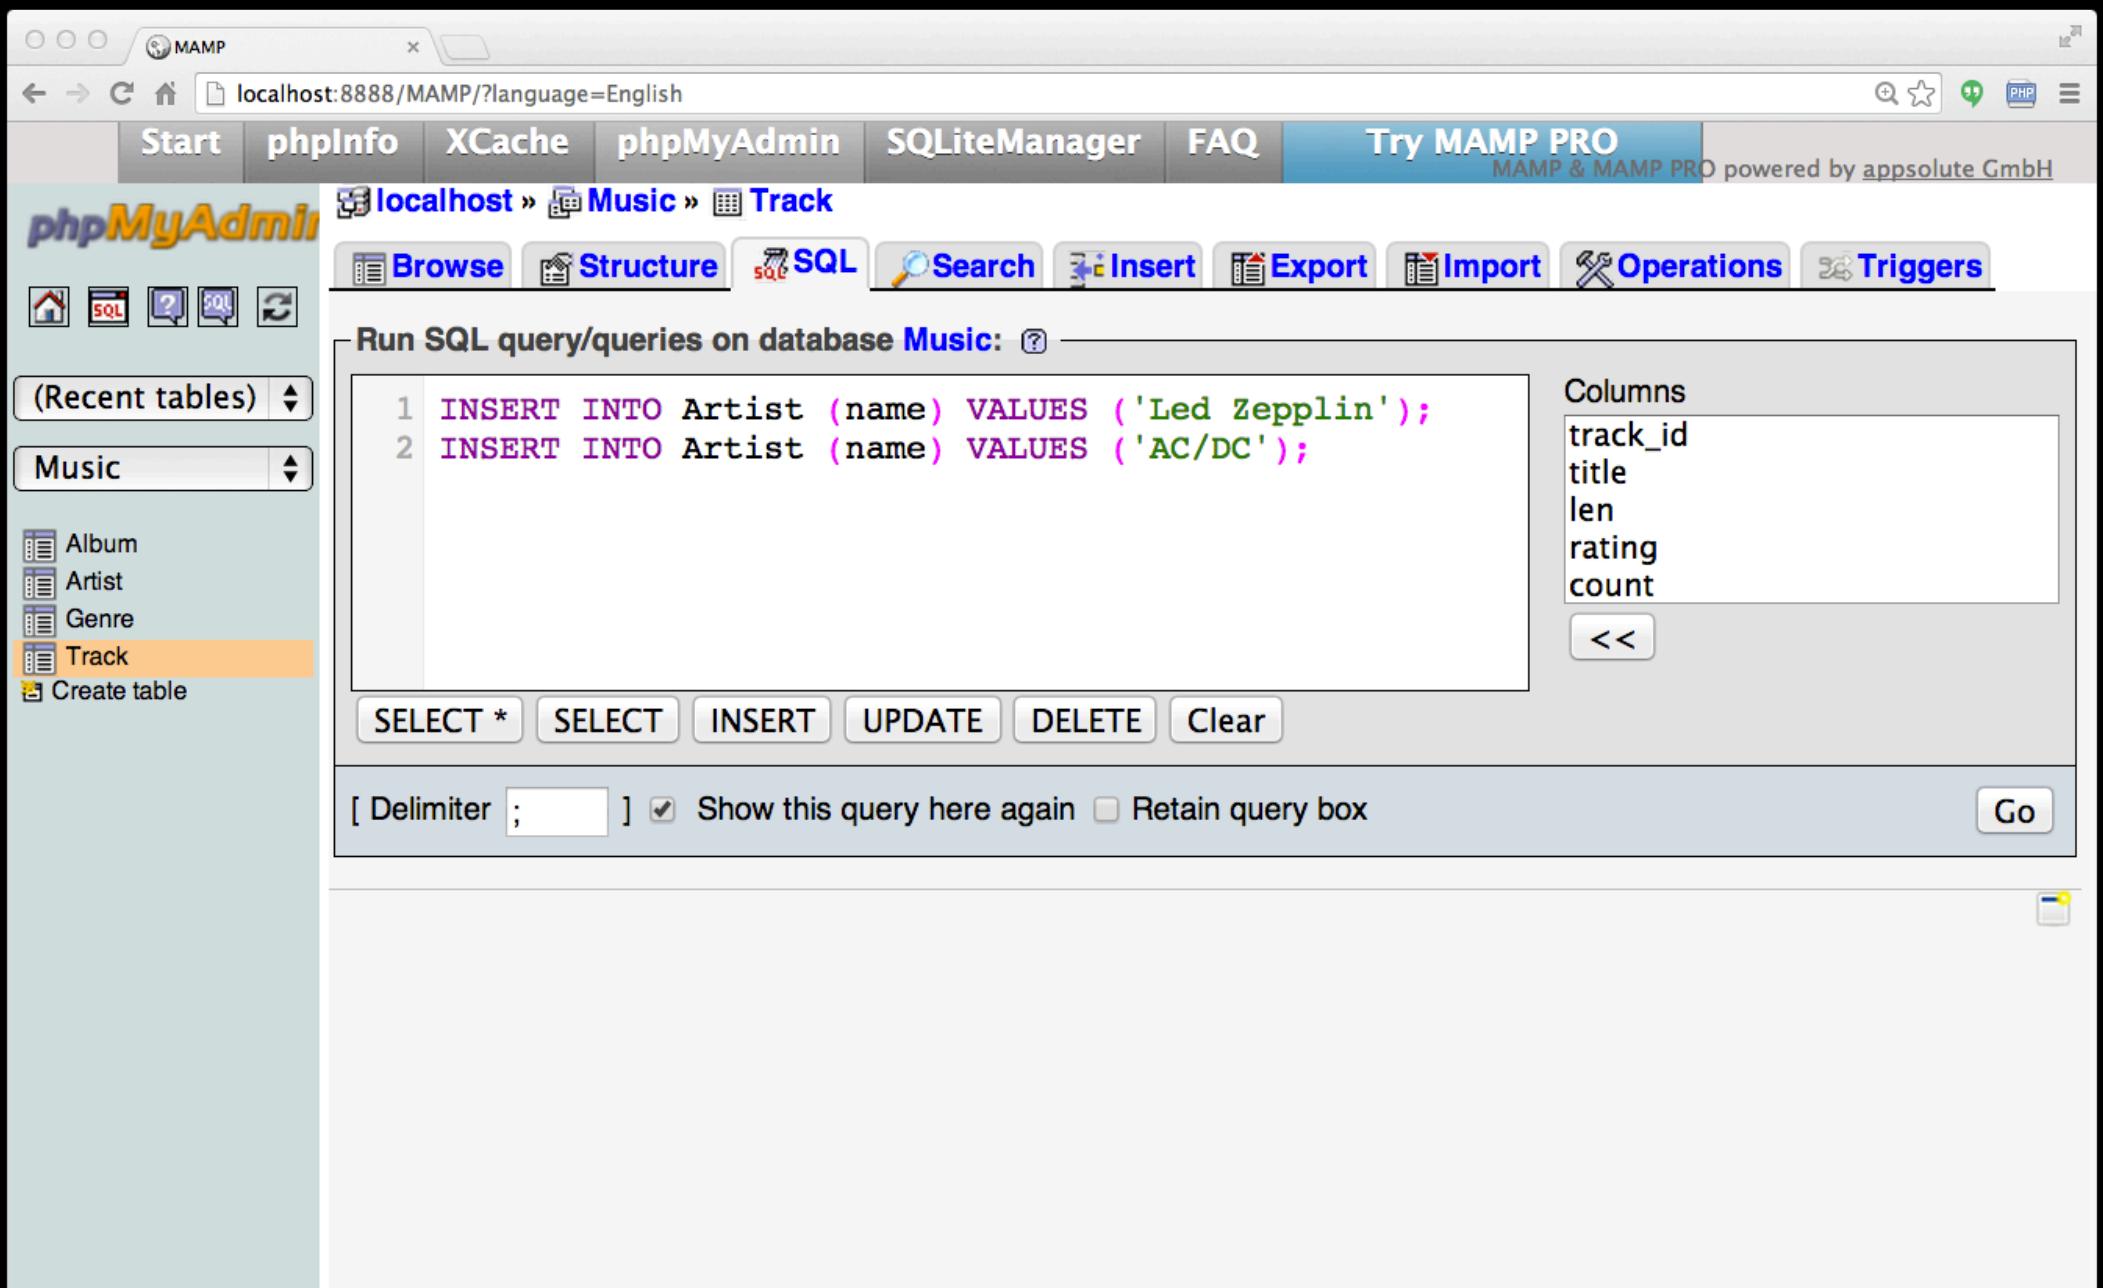
Check All / Uncheck All With selected:     

 Print view  Relation view  Propose table structure  Add 1 column(s) At End of Table At Beginning of Table After track_id 

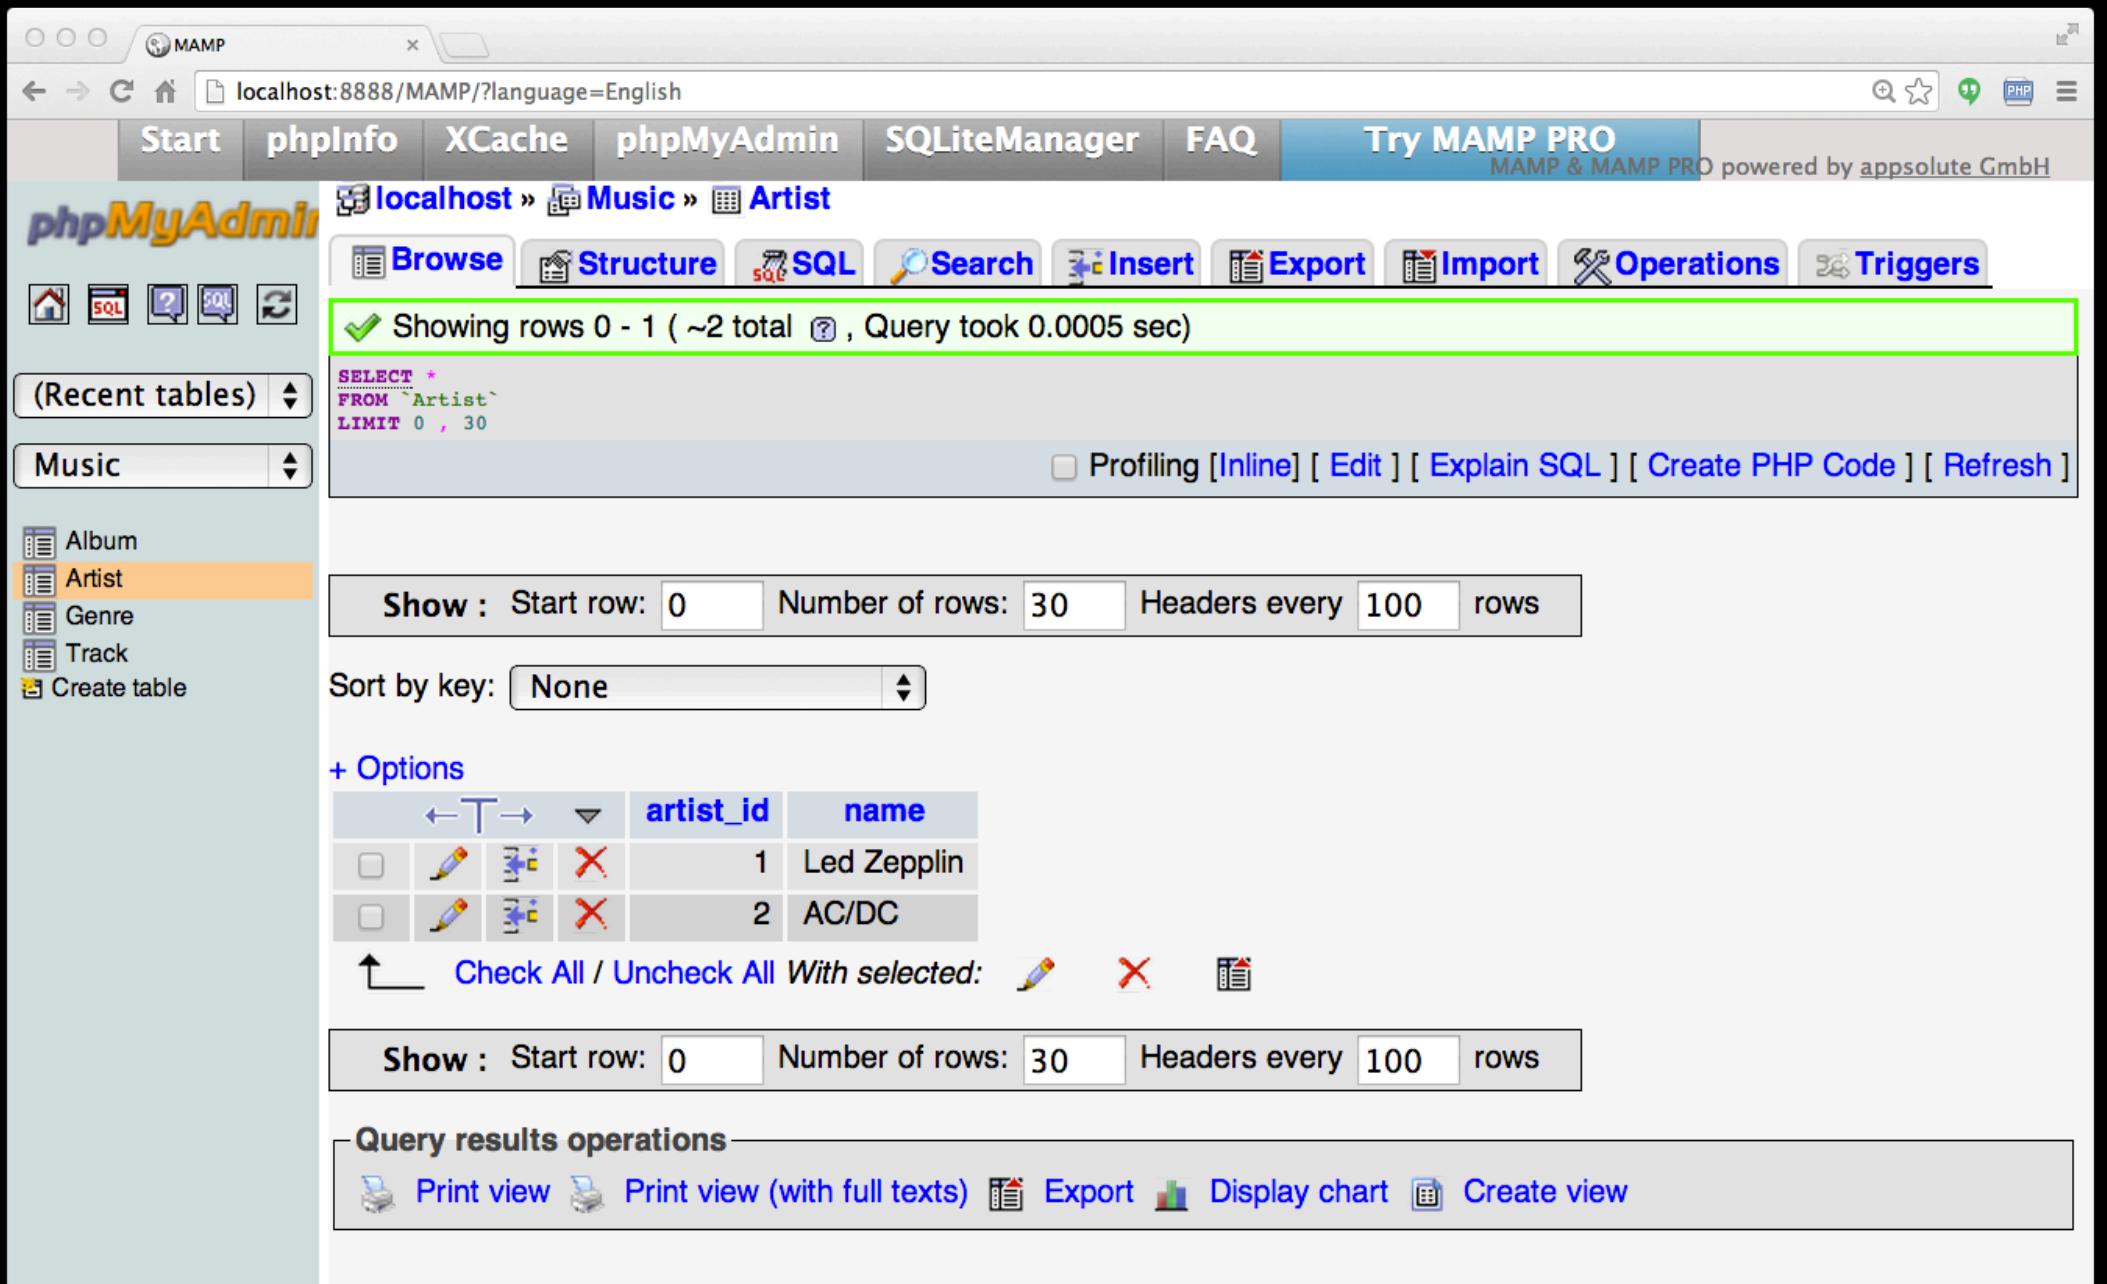
+ Indexes

Information

Space usage		Row Statistics	
Type	Usage	Statements	Value
Data	16 KiB	Format	Compact
Index	48 KiB	Collation	utf8_general_ci
Total	64 KiB	Next autoindex	1



INSERT INTO Artist (name) VALUES ('Led Zepplin');
INSERT INTO Artist (name) VALUES ('AC/DC');



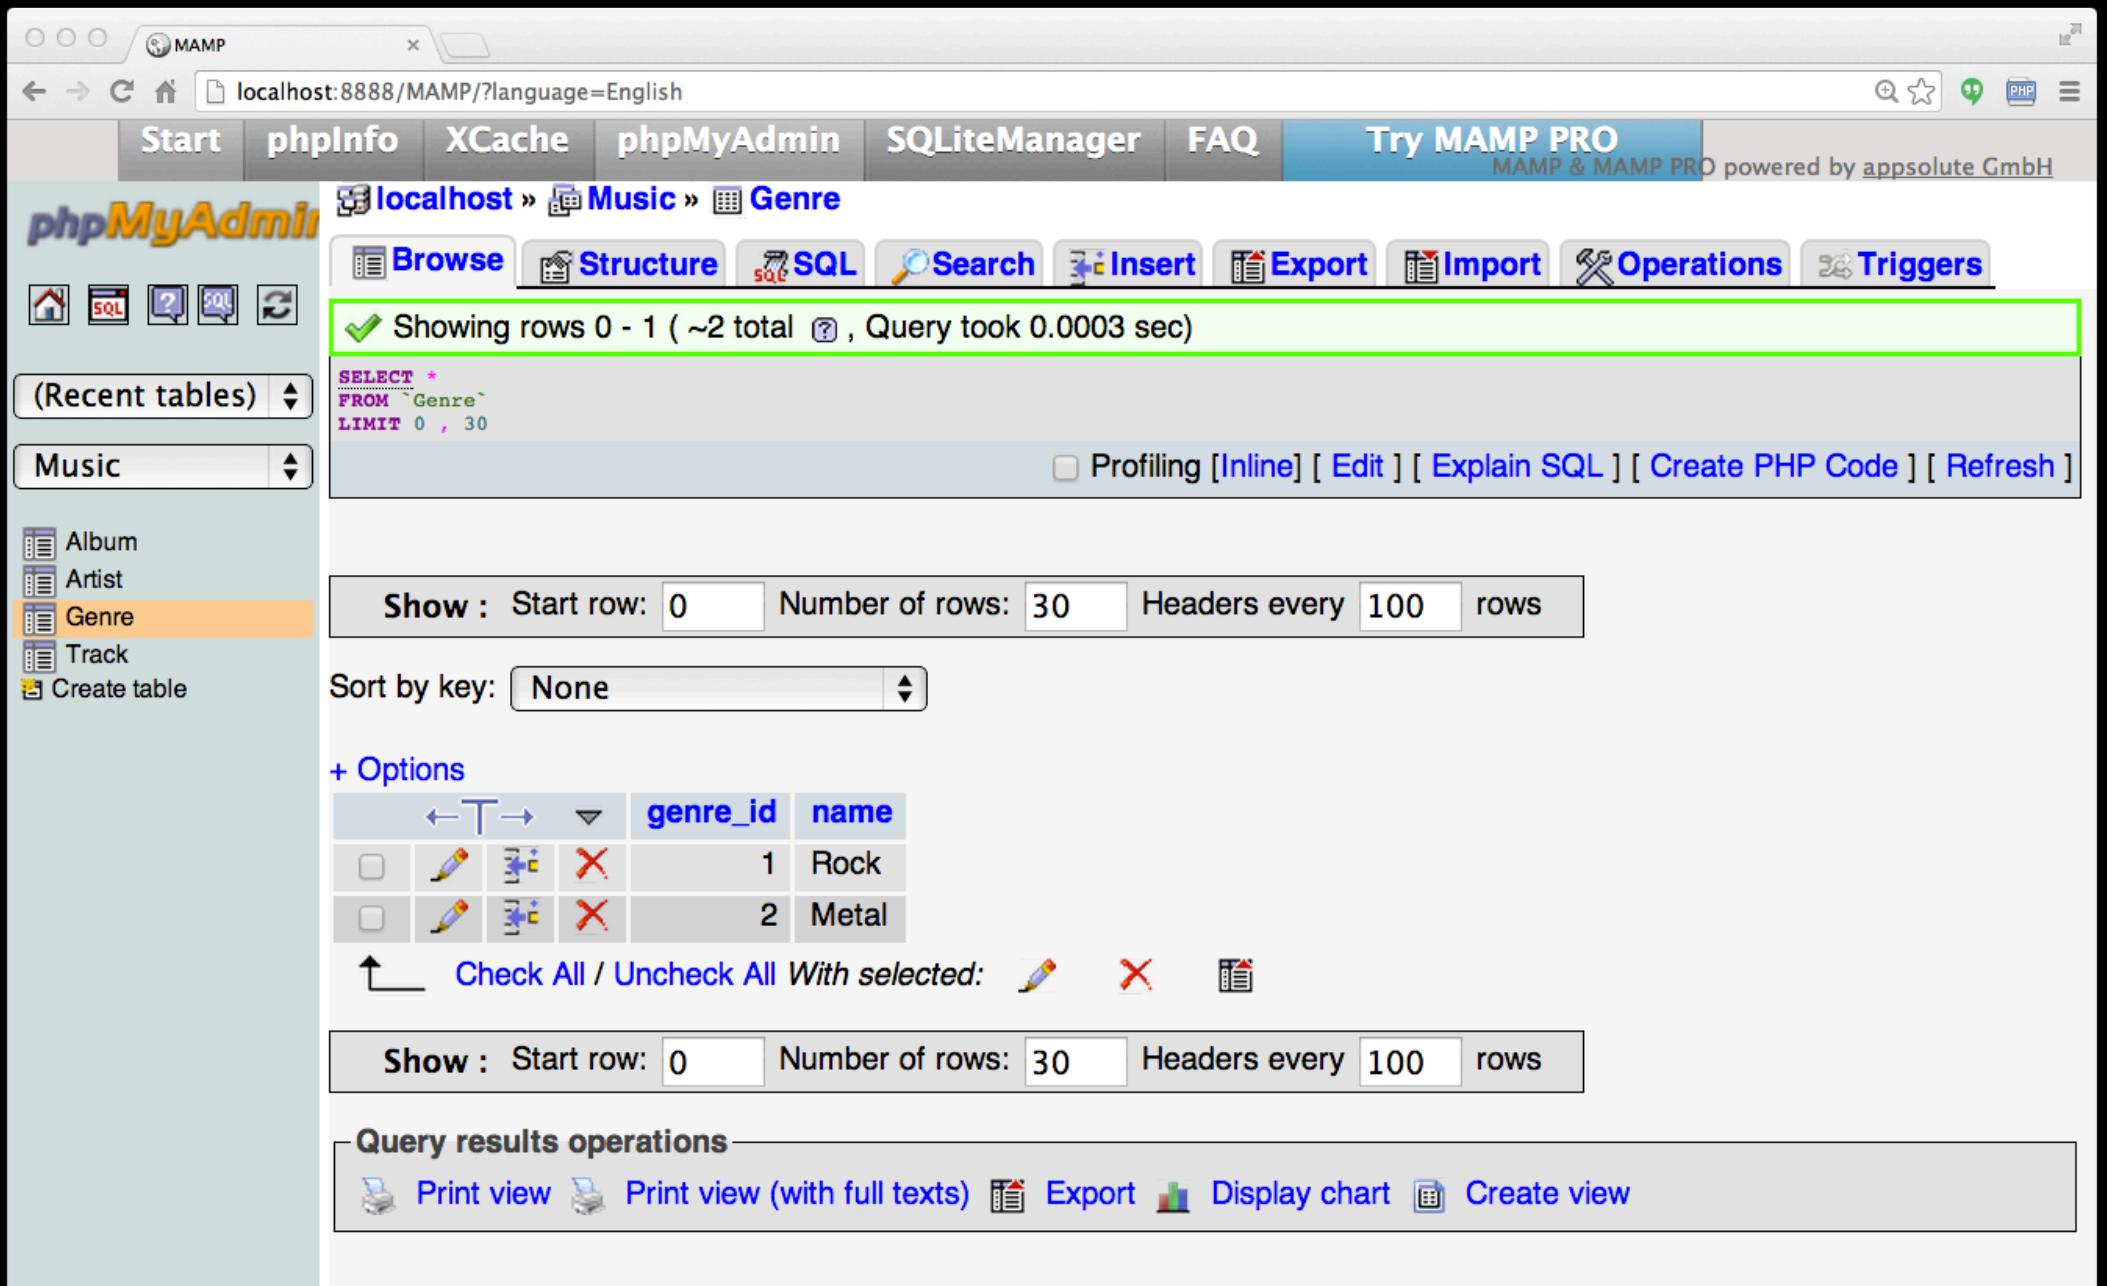
The screenshot shows the phpMyAdmin interface for a database named 'Music'. The 'Artist' table is selected. The table structure is shown with columns: artist_id and name. Two rows are present: one for 'Led Zeppelin' (artist_id 1) and one for 'AC/DC' (artist_id 2). The interface includes a SQL query history at the top, a toolbar with various operations, and a footer with navigation links.

	artist_id	name
<input type="checkbox"/>	1	Led Zeppelin
<input type="checkbox"/>	2	AC/DC

Query results operations:

- Print view
- Print view (with full texts)
- Export
- Display chart
- Create view

```
INSERT INTO Artist (name) VALUES ('Led Zepplin');
INSERT INTO Artist (name) VALUES ('AC/DC');
```



The screenshot shows the phpMyAdmin interface for a database named 'Music'. The 'Genre' table is selected in the left sidebar. The main area displays the results of a query:

```
SELECT *  
FROM `Genre`  
LIMIT 0 , 30
```

The results show two rows:

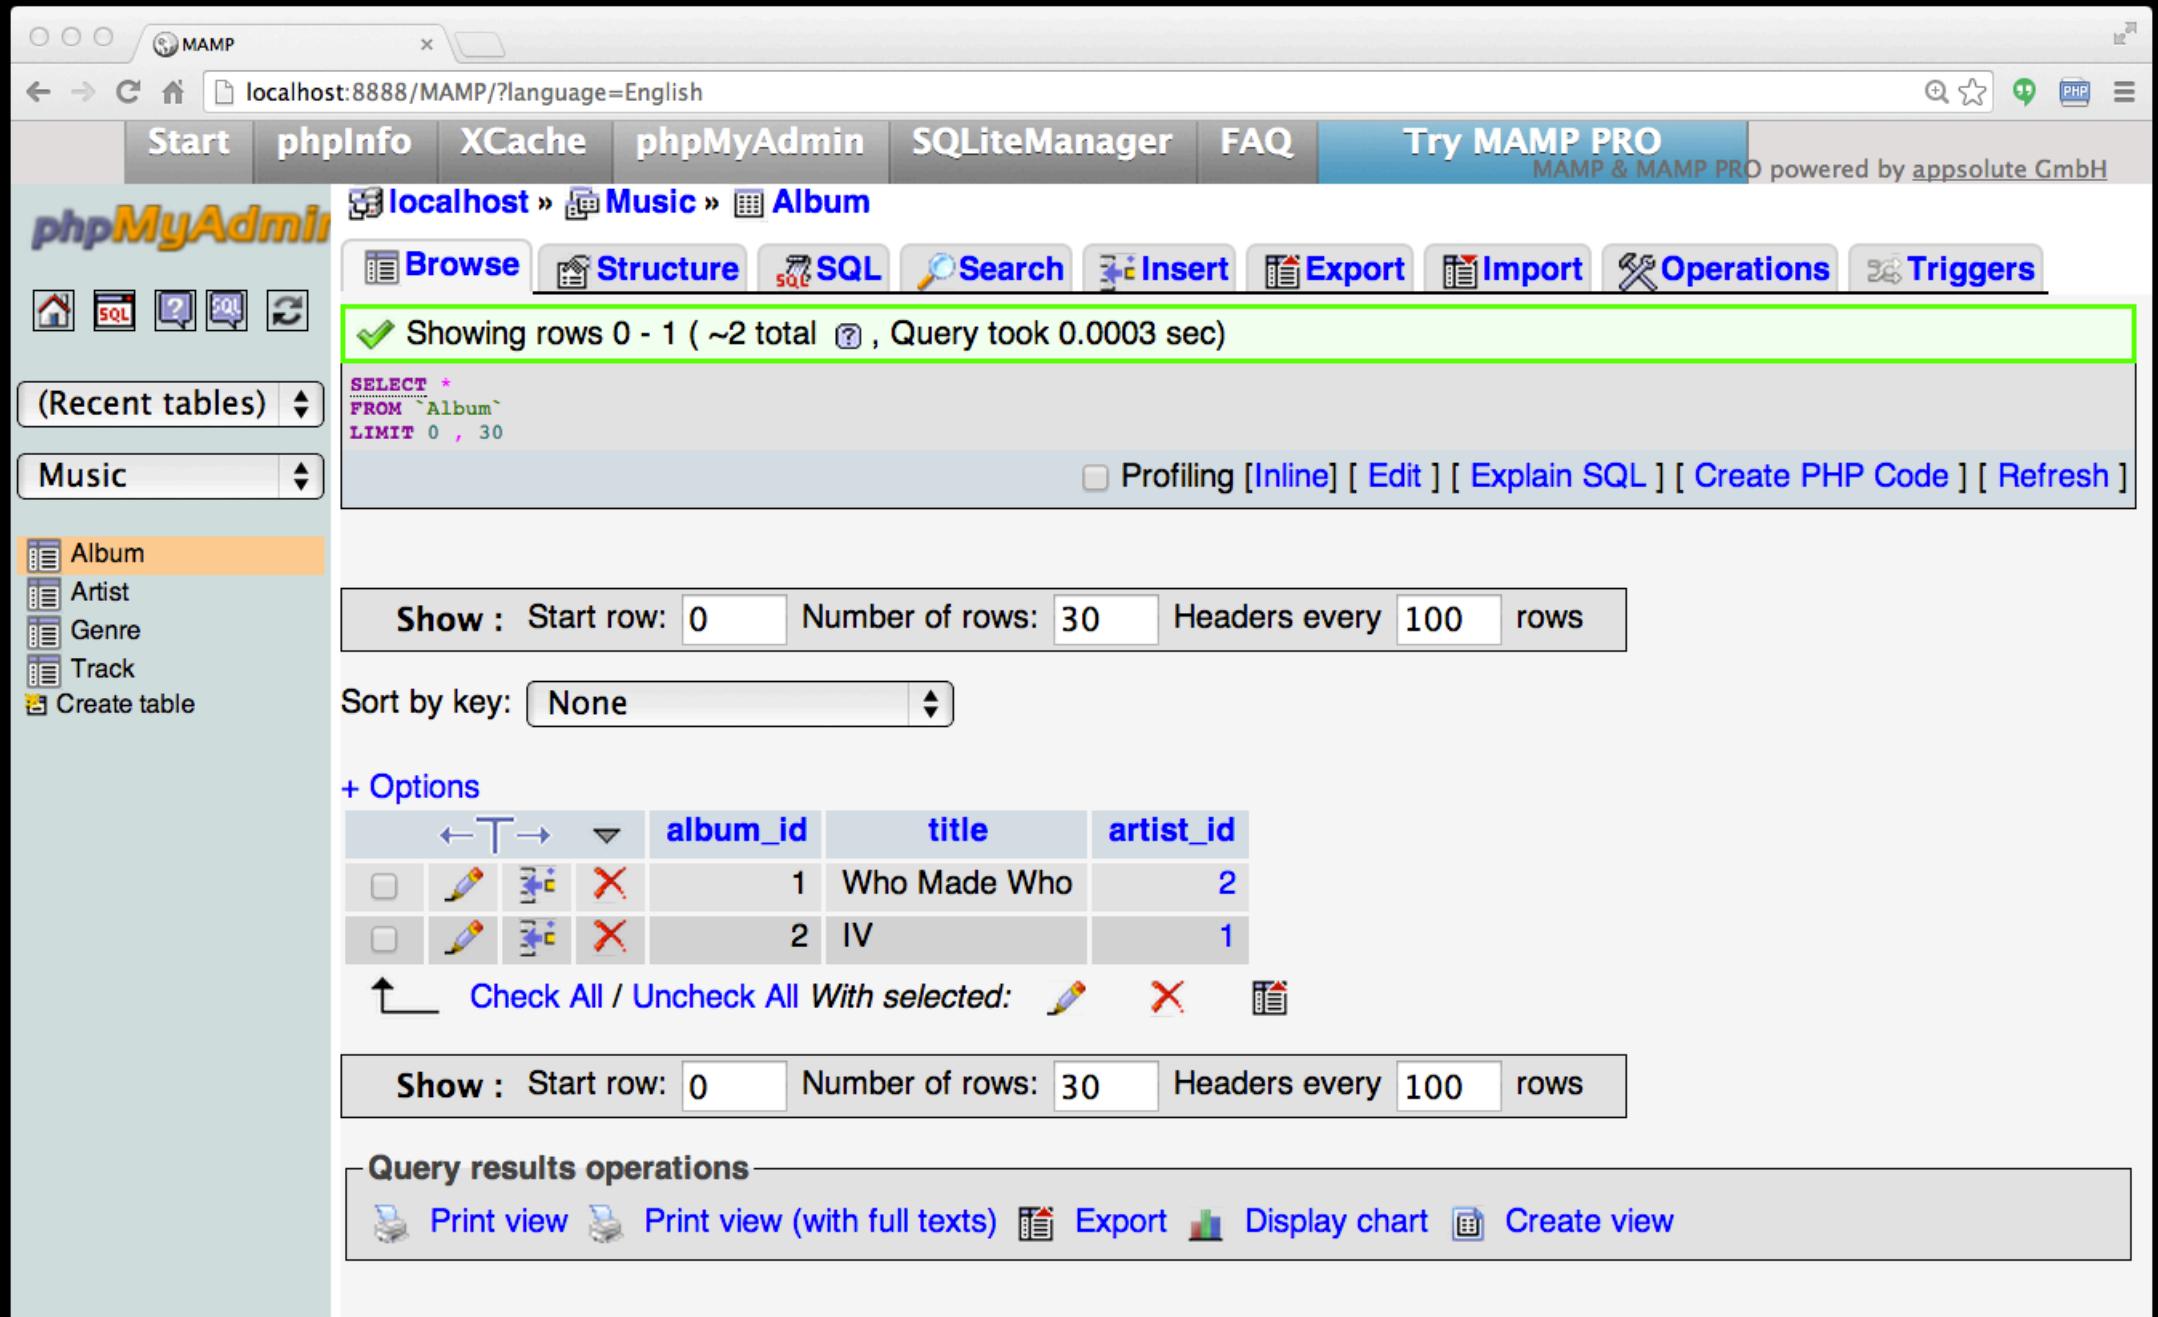
	genre_id	name
<input type="checkbox"/>	1	Rock
<input type="checkbox"/>	2	Metal

Below the table, there are buttons for 'Check All / Uncheck All' and 'With selected' (edit, delete, etc.).

At the bottom, there are additional navigation and export options:

- Print view
- Print view (with full texts)
- Export
- Display chart
- Create view

INSERT INTO Genre (name) VALUES ('Rock');
INSERT INTO Genre (name) VALUES ('Metal');



The screenshot shows the phpMyAdmin interface for a database named 'Music'. The 'Album' table is selected. The SQL query executed is:

```
SELECT *  
FROM `Album`  
LIMIT 0 , 30
```

The results show two rows:

	album_id	title	artist_id
<input type="checkbox"/>	1	Who Made Who	2
<input type="checkbox"/>	2	IV	1

Below the table, there are buttons for 'Check All / Uncheck All With selected' and additional 'Show' and 'Sort by key' options.

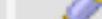
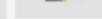
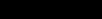
```
INSERT INTO Album (title, artist_id) VALUES ('Who Made Who', 2);  
INSERT INTO Album (title, artist_id) VALUES ('IV', 1);
```

```
INSERT INTO Track
    (title, rating, len, count, album_id, genre_id)
VALUES ('Black Dog', 5, 297, 0, 2, 1);

INSERT INTO Track
    (title, rating, len, count, album_id, genre_id)
VALUES ('Stairway', 5, 482, 0, 2, 1);

INSERT INTO Track
    (title, rating, len, count, album_id, genre_id)
VALUES ('About to Rock', 5, 313, 0, 1, 2);

INSERT INTO Track
    (title, rating, len, count, album_id, genre_id)
VALUES ('Who Made Who', 5, 207, 0, 1, 2);
```

	←	↑	→	▼	track_id	title	len	rating	count	album_id	genre_id
<input type="checkbox"/>					1	Black Dog	297	5	0	2	1
<input type="checkbox"/>					2	Stairway	482	5	0	2	1
<input type="checkbox"/>					3	About to Rock	313	5	0	1	2
<input type="checkbox"/>					4	Who Made Who	207	5	0	1	2

We Have Relationships!

Track

	← T →	▼	track_id	title	len	rating	count	album_id	genre_id
			X	1	Black Dog	297	5	0	2
			X	2	Stairway	482	5	0	2
			X	3	About to Rock	313	5	0	1
			X	4	Who Made Who	207	5	0	1

Album

	← T →	▼	album_id	title	artist_id
			X	1	Who Made Who
			X	2	IV

Artist

	← T →	▼	artist_id	name
			X	1
			X	2

Genre

	← T →	▼	genre_id	name
			X	1
			X	2



Using Join Across Tables

[http://en.wikipedia.org/wiki/Join_\(SQL\)](http://en.wikipedia.org/wiki/Join_(SQL))

Relational Power

- By removing the replicated data and replacing it with references to a single copy of each bit of data, we build a “**web**” of information that the relational database can read through very quickly - even for very large amounts of data.
- Often when you want some data it comes from a number of tables linked by these **foreign keys**.

The JOIN Operation

- The JOIN operation **links across several tables** as part of a SELECT operation.
- You must tell the JOIN **how to use the keys** that make the connection between the tables using an **ON clause**.

title	name
IV	Led Zepplin
Who Made Who	AC/DC

What we want
to see

album_id	title	artist_id
1	Who Made Who	2
2	IV	1

artist_id	name
1	Led Zepplin
2	AC/DC

The tables that
hold the data

SELECT Album.title,Artist.name FROM Album JOIN Artist ON
Album.artist_id = Artist.artist_id

How the tables
are linked

	← T →	▼	album_id	title	artist_id	
<input type="checkbox"/>			X	1	Who Made Who	2
<input type="checkbox"/>			X	2	IV	1

	← T →	▼	artist_id	name
<input type="checkbox"/>			X	1 Led Zepplin
<input type="checkbox"/>			X	2 AC/DC

Album.title Album.artist_id Artist.artist_id Artist.name

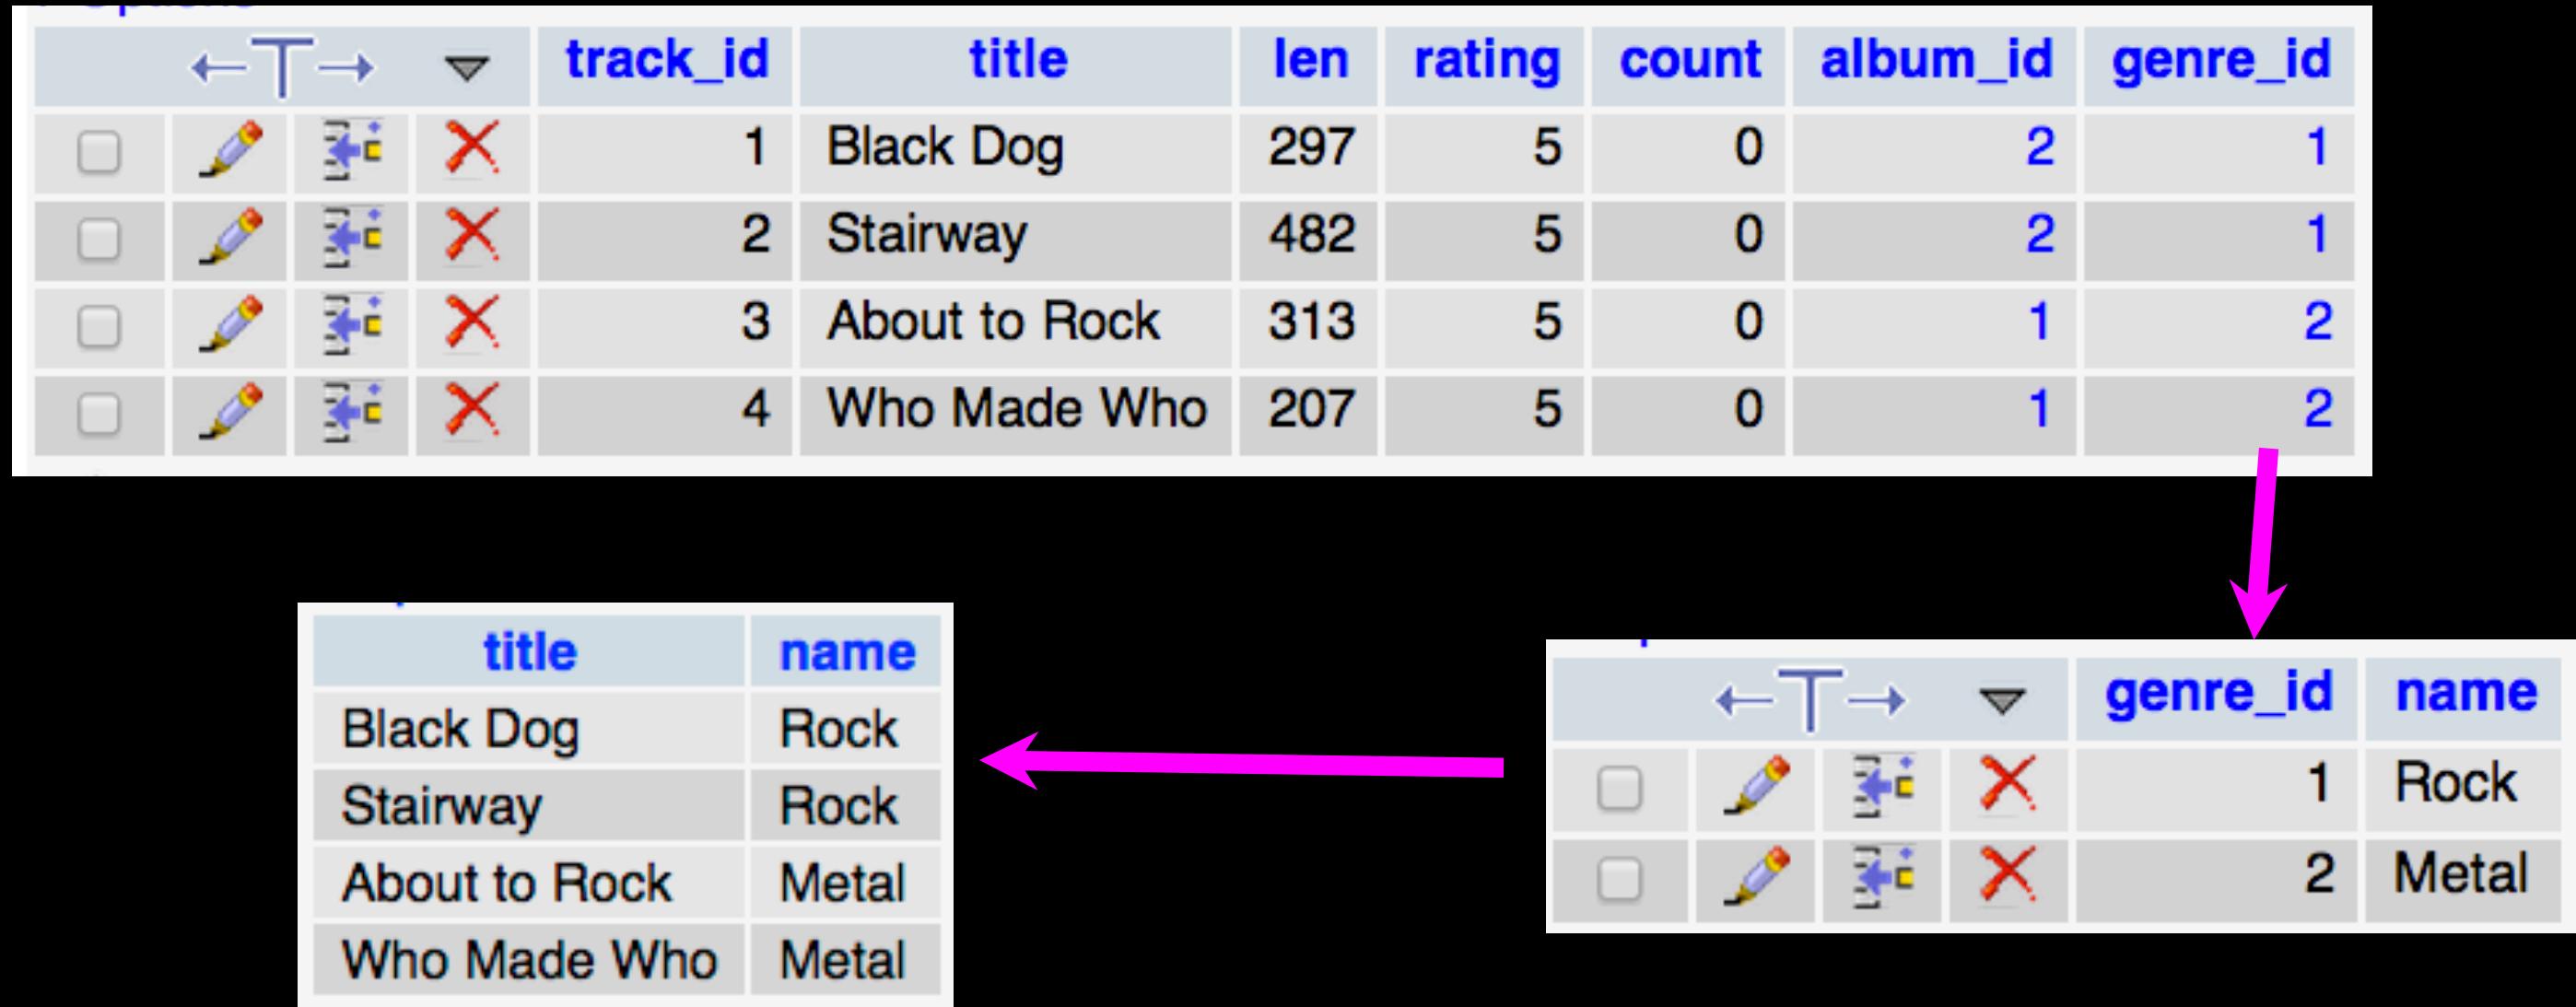
title	artist_id	artist_id	name
IV	1	1	Led Zepplin
Who Made Who	2	2	AC/DC

```
SELECT Album.title, Album.artist_id, Artist.artist_id, Artist.name  
FROM Album JOIN Artist ON Album.artist_id = Artist.artist_id
```

title	genre_id	genre_id	name
Black Dog	1	1	Rock
Black Dog	1	2	Metal
Stairway	1	1	Rock
Stairway	1	2	Metal
About to Rock	2	1	Rock
About to Rock	2	2	Metal
Who Made Who	2	1	Rock
Who Made Who	2	2	Metal

```
SELECT Track.title,  
       Track.genre_id,  
       Genre.genre_id,  
       Genre.name  
  FROM Track JOIN Genre
```

Joining two tables without an **ON** clause gives all possible combinations of rows.



```
SELECT Track.title, Genre.name FROM Track JOIN Genre  
ON Track.genre_id = Genre.genre_id
```

It Can Get Complex...

```
SELECT Track.title, Artist.name, Album.title, Genre.name  
FROM Track JOIN Genre JOIN Album JOIN Artist ON  
Track.genre_id = Genre.genre_id AND Track.album_id =  
Album.album_id AND Album.artist_id = Artist.artist_id
```

title	name	title	name
About to Rock	AC/DC	Who Made Who	Metal
Who Made Who	AC/DC	Who Made Who	Metal
Black Dog	Led Zepplin	IV	Rock
Stairway	Led Zepplin	IV	Rock

What we want
to see
The tables that
hold the data
How the tables
are linked



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input checked="" type="checkbox"/> Tin Man	3:30	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...		24
<input checked="" type="checkbox"/> Track 01	4:22	Billy Price	Danger Zone	Blues/R&B		26
<input checked="" type="checkbox"/> Track 02	2:45	Billy Price	Danger Zone	Blues/R&B		18
<input checked="" type="checkbox"/> Track 03	3:26	Billy Price	Danger Zone	Blues/R&B		22
<input checked="" type="checkbox"/> Track 04	4:17	Billy Price	Danger Zone	Blues/R&B		18
<input checked="" type="checkbox"/> Track 05	3:50	Billy Price	Danger Zone	Blues/R&B		21
<input checked="" type="checkbox"/> War Pigs/Luke's Wall	7:58	Black Sabbath	Paranoid	Metal		25

title	name	title	name
About to Rock	AC/DC	Who Made Who	Metal
Who Made Who	AC/DC	Who Made Who	Metal
Black Dog	Led Zepplin	IV	Rock
Stairway	Led Zepplin	IV	Rock
2:58	Brent	Brent's Album	1

ON DELETE CASCADE

	track_id	title	len	rating	count	album_id	genre_id
	1	Black Dog	297	5	0	2	1
	2	Stairway	482	5	0	2	1
	3	About to Rock	313	5	0	1	2
	4	Who Made Who	207	5	0	1	2

Child

	genre_id	name
	1	Rock
	2	Metal

Parent

We are telling MySQL to
"clean up" broken references

```
DELETE FROM Genre WHERE name = 'Metal'
```

ON DELETE CASCADE

	track_id	title	len	rating	count	album_id	genre_id
	1	Black Dog	297	5	0	2	1
	2	Stairway	482	5	0	2	1
	3	About to Rock	313	5	0	1	2
	4	Who Made Who	207	5	0	1	2

```
DELETE FROM Genre WHERE name = 'Metal'
```

	track_id	title	len	rating	count	album_id	genre_id
	1	Black Dog	297	5	0	2	1
	2	Stairway	482	5	0	2	1

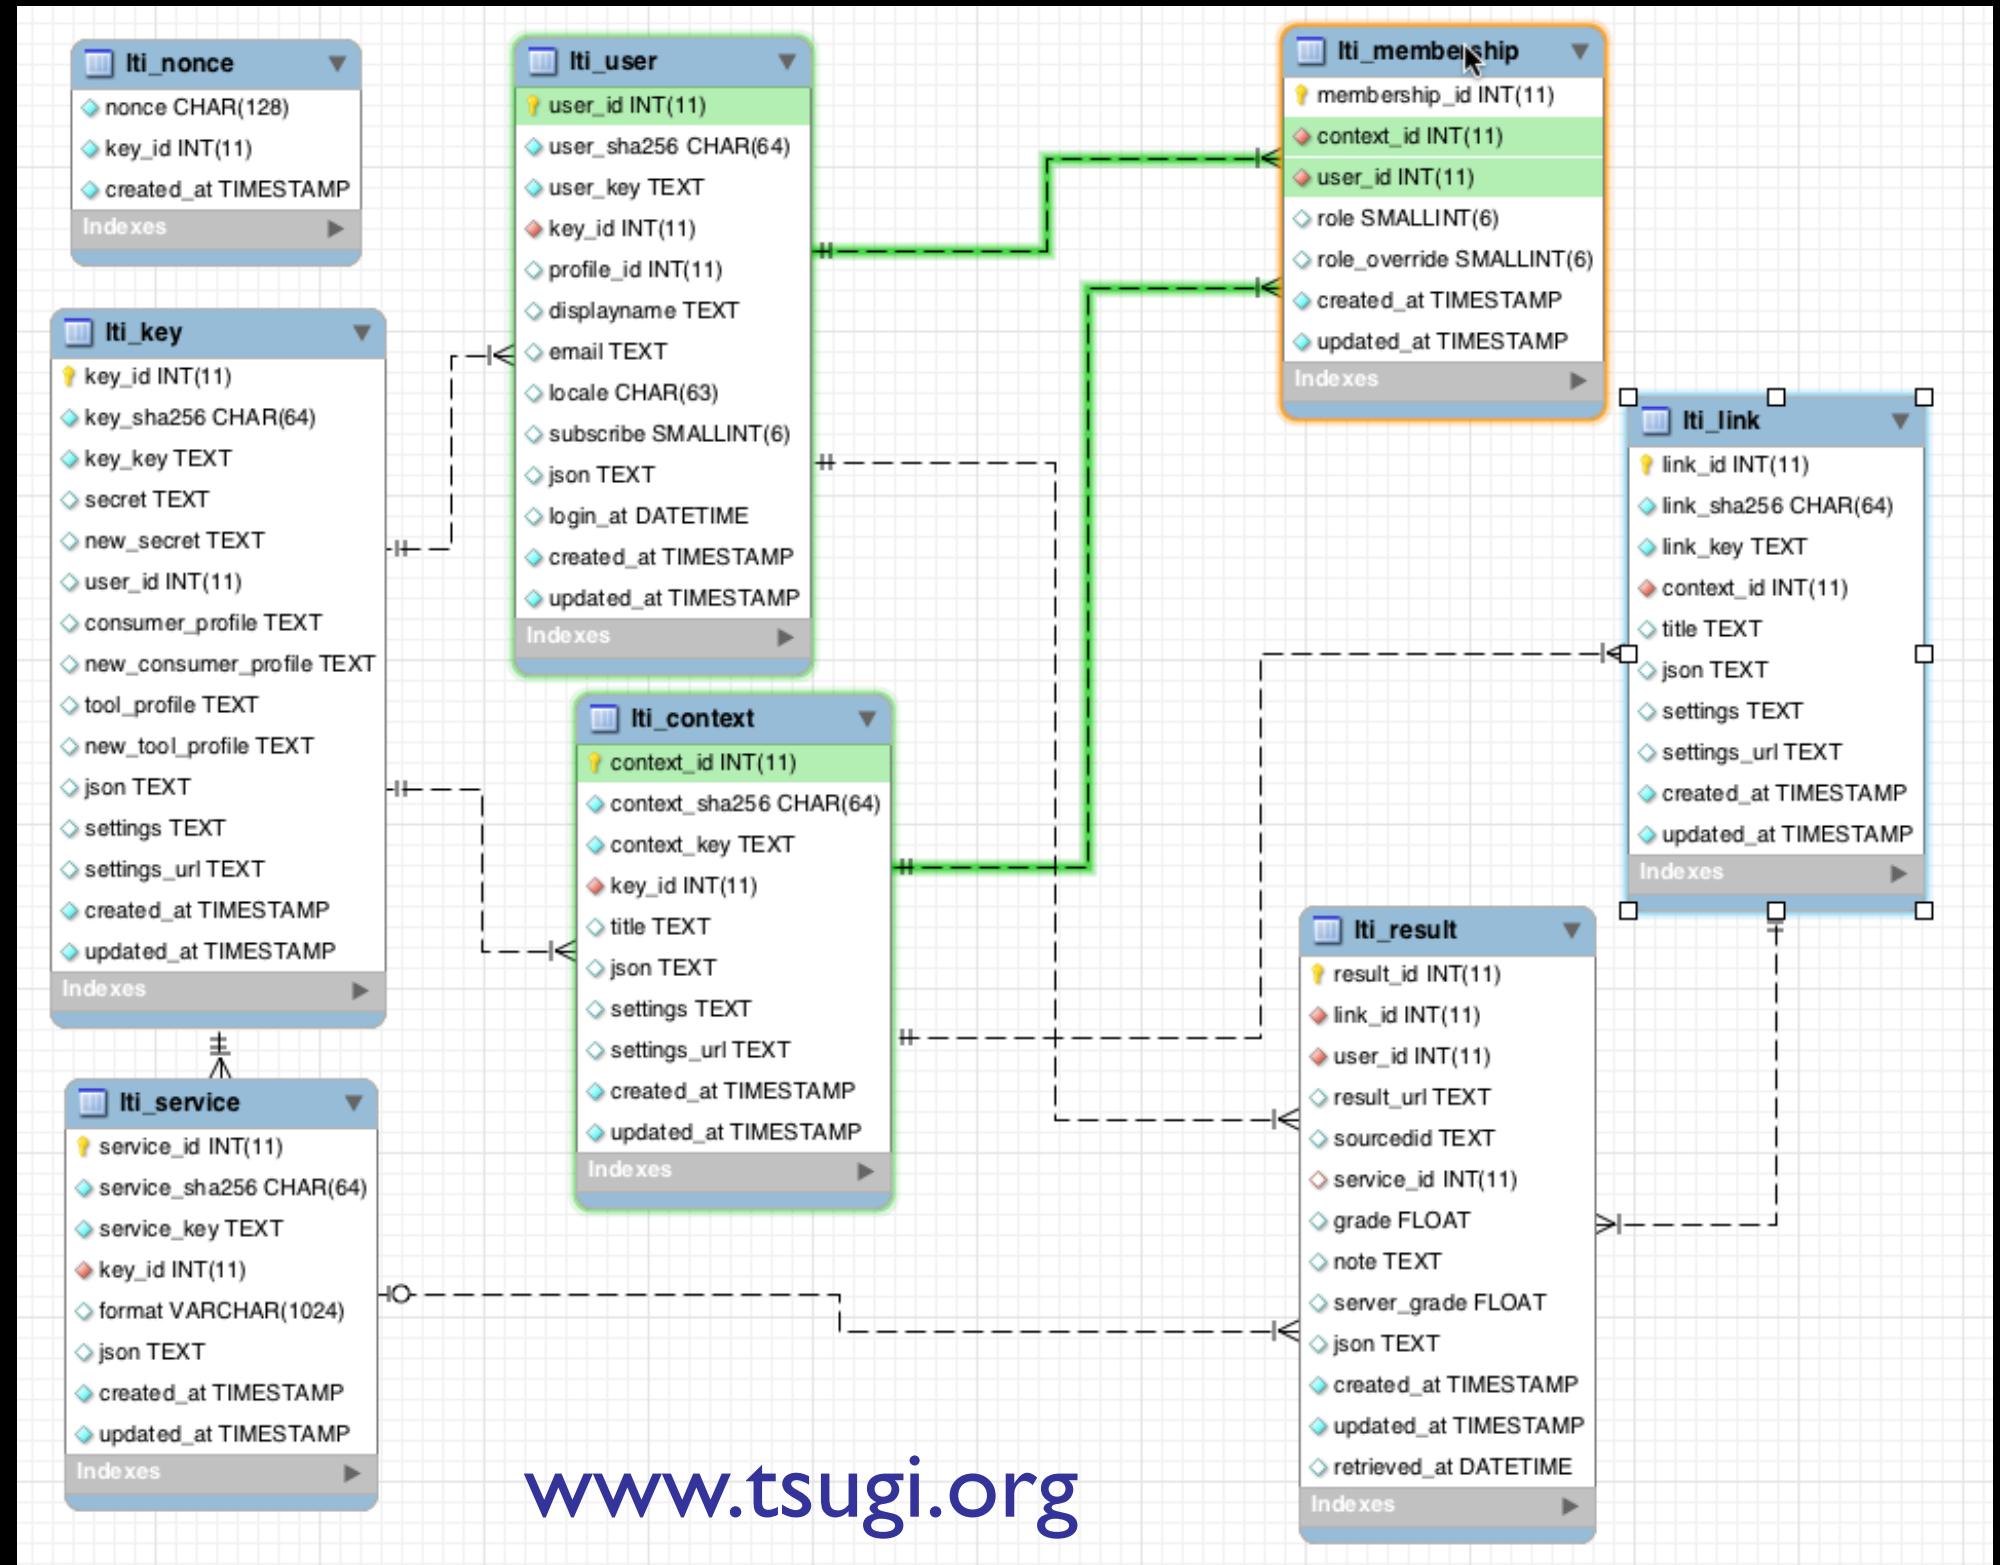
ON DELETE Choices

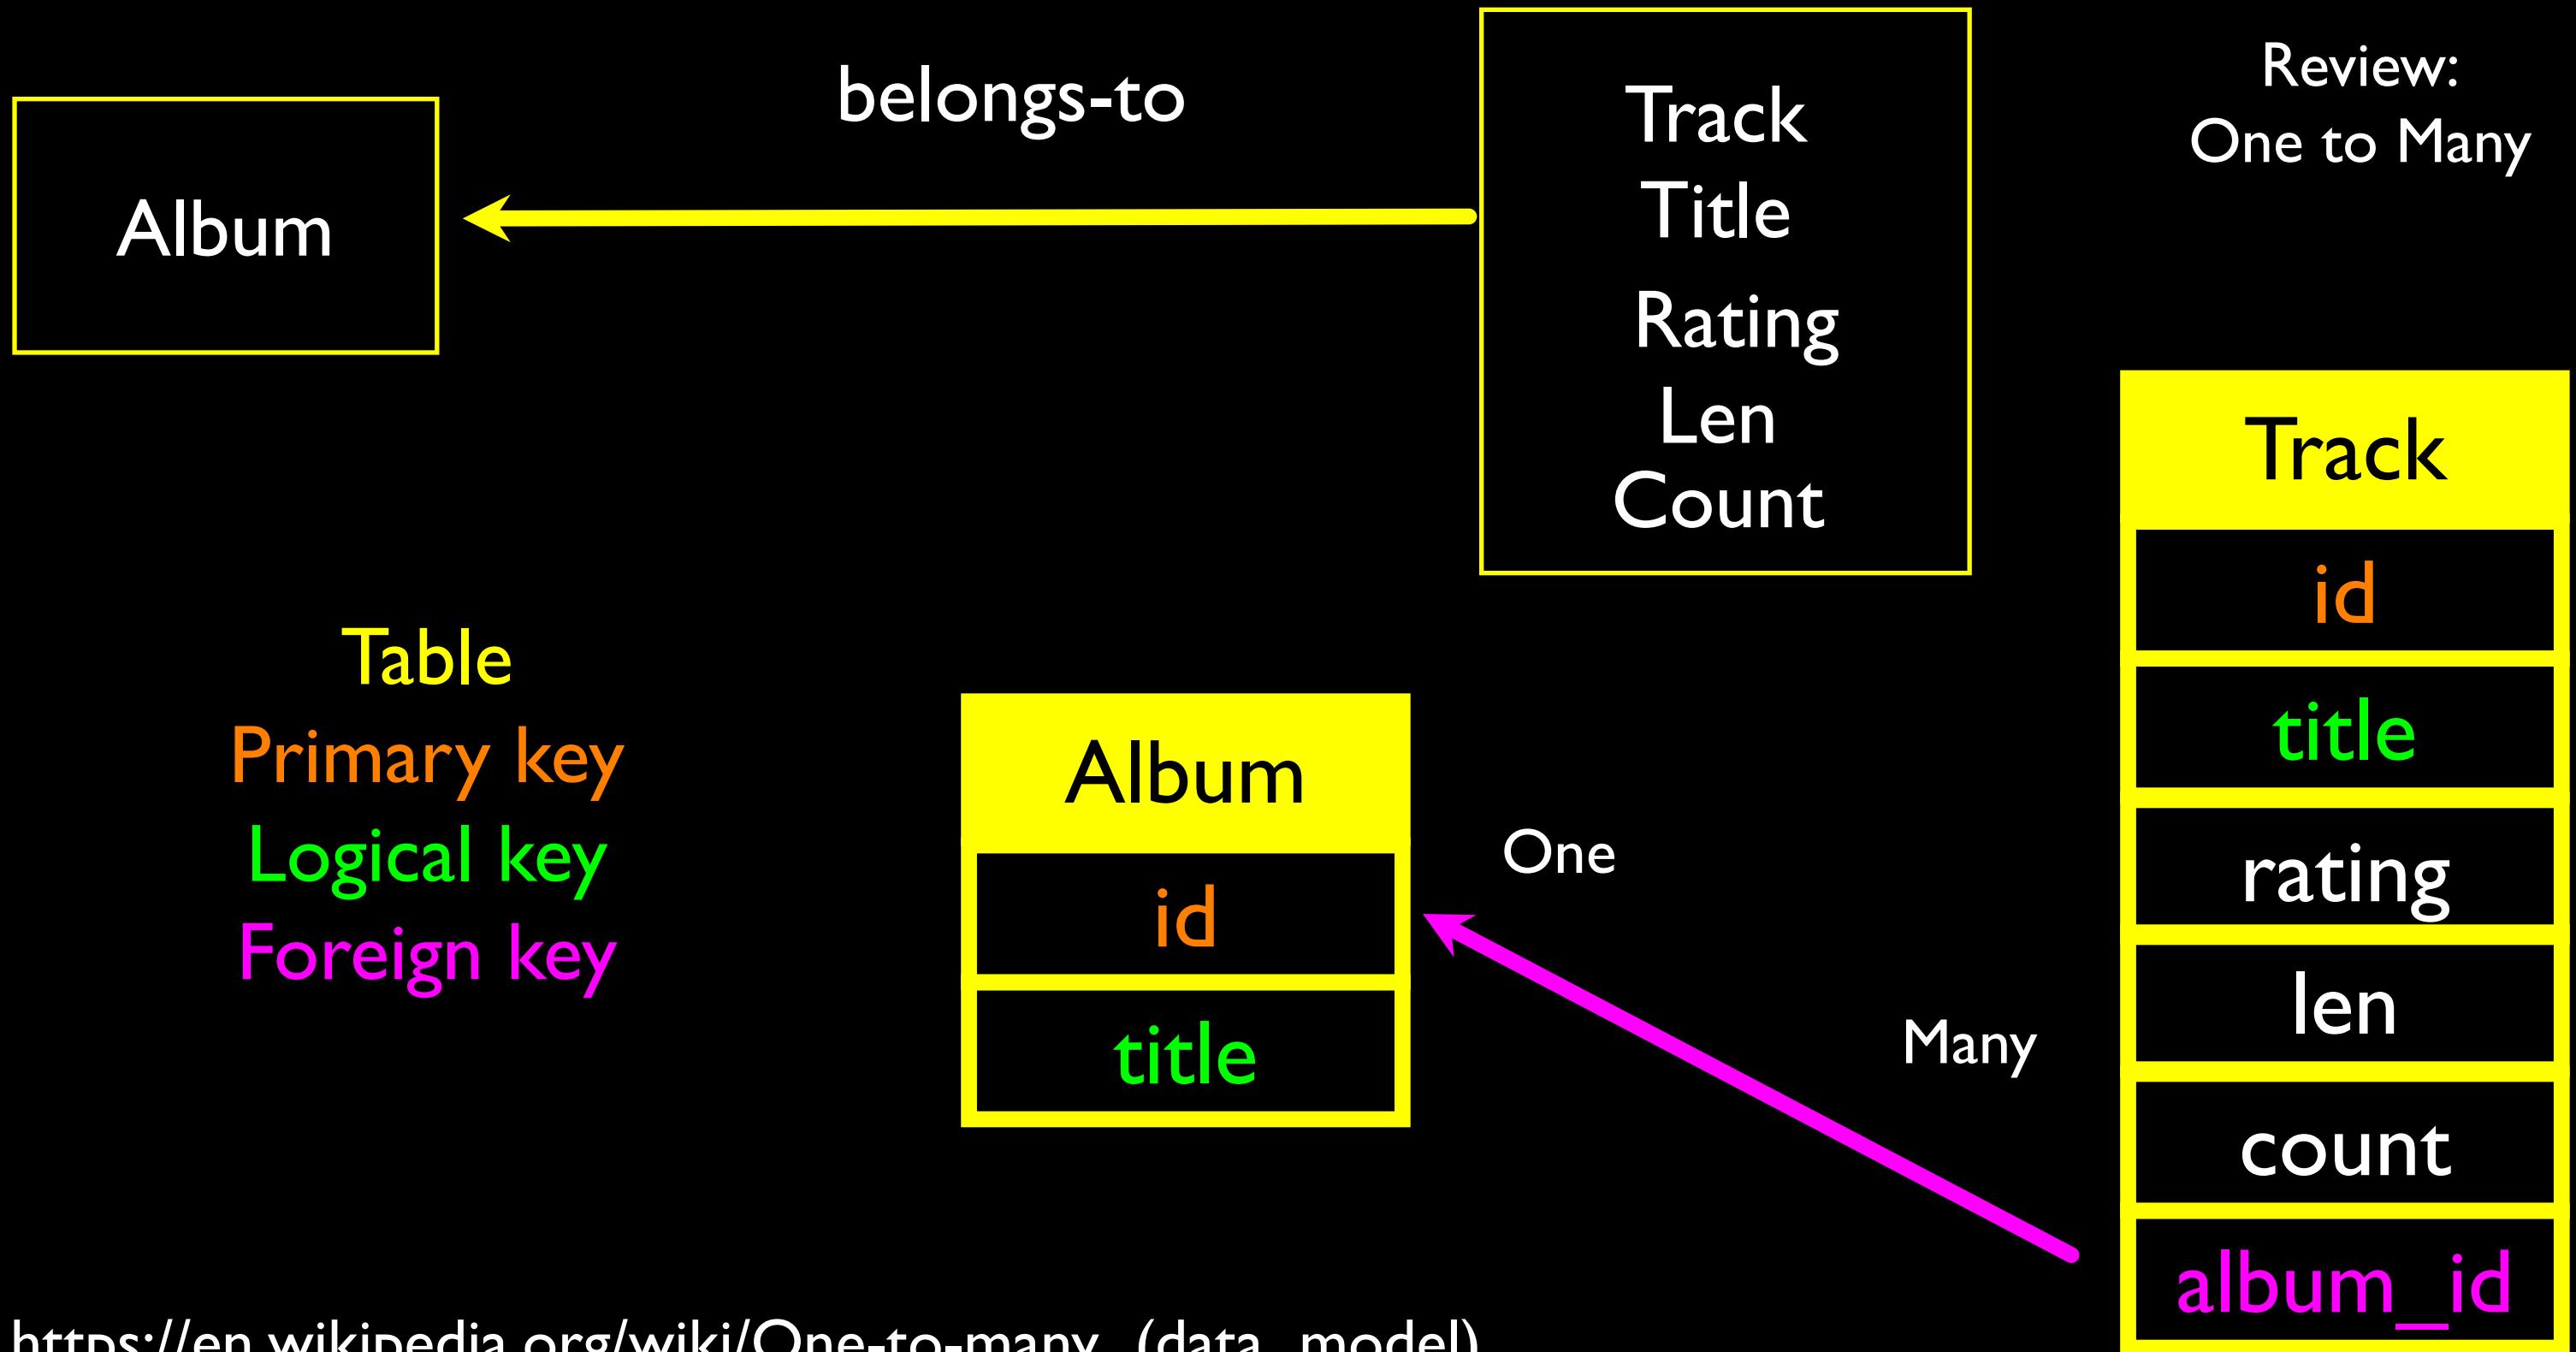
- **Default / RESTRICT** – Don't allow changes that break the constraint
- **CASCADE** – Adjust child rows by removing or updating to maintain consistency
- **SET NULL** – Set the foreign key columns in the child rows to null

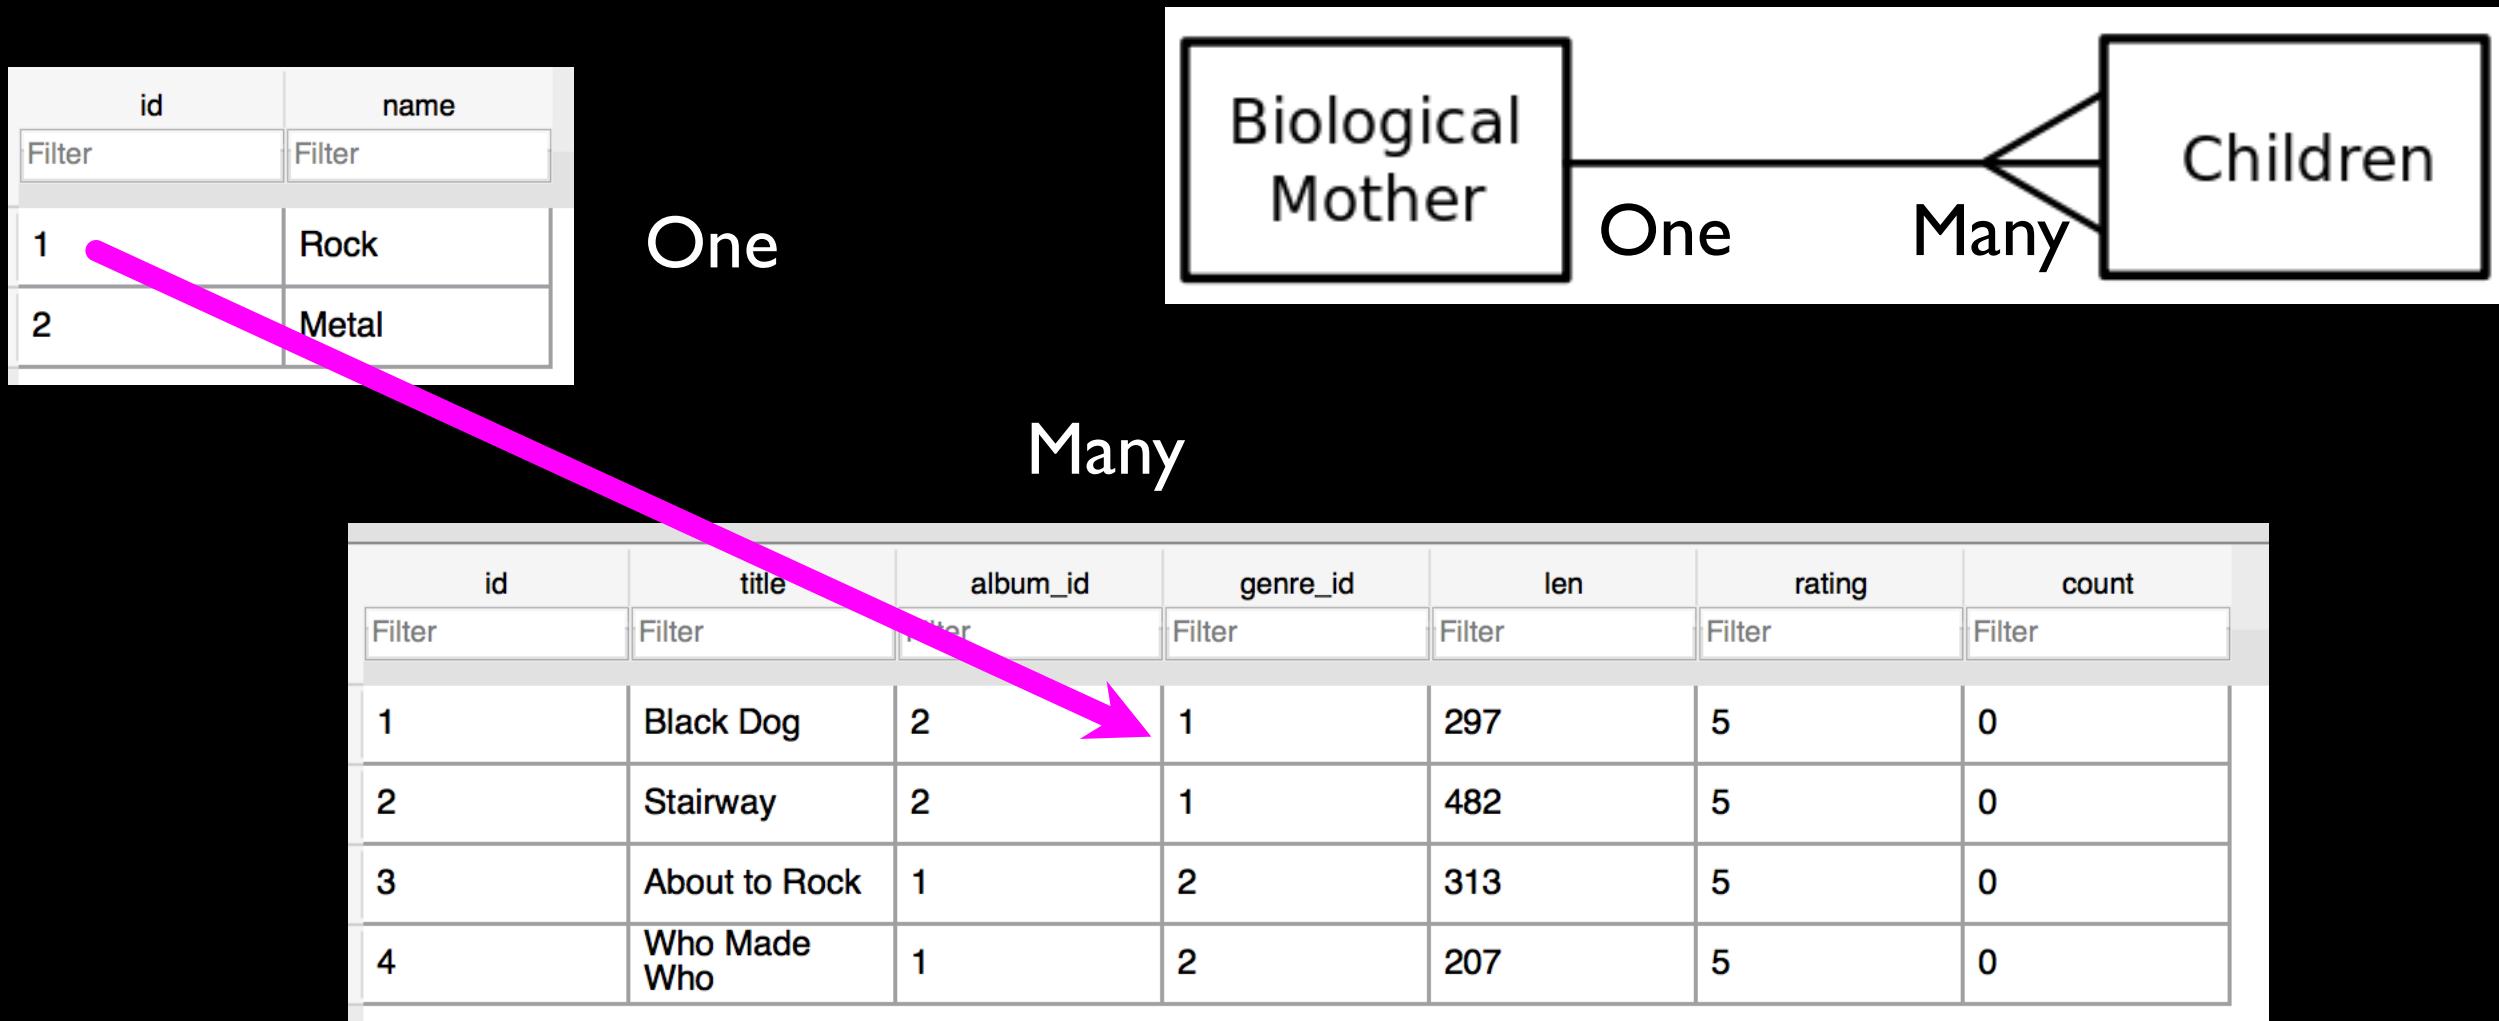
<http://stackoverflow.com/questions/1027656/what-is-mysqls-default-on-delete-behavior>



Many-to-Many Relationships



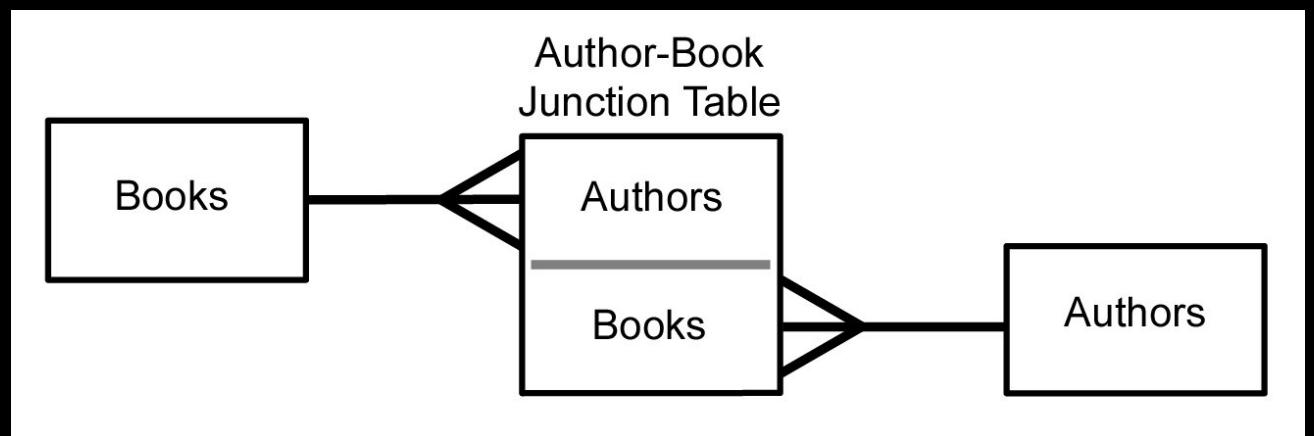
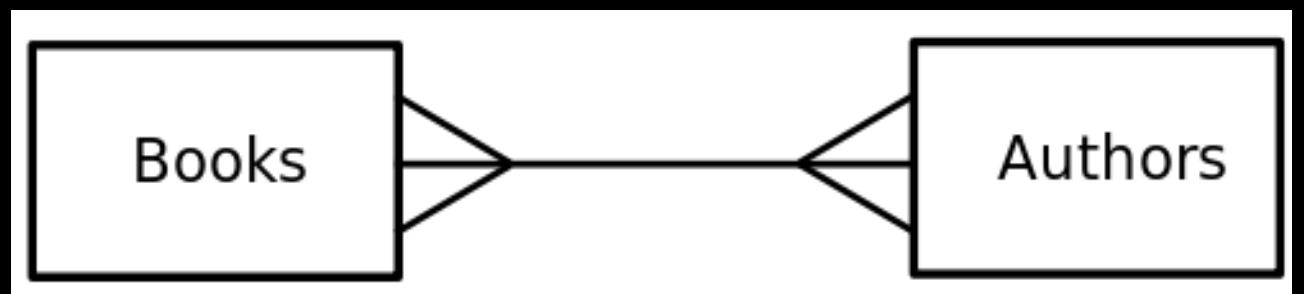


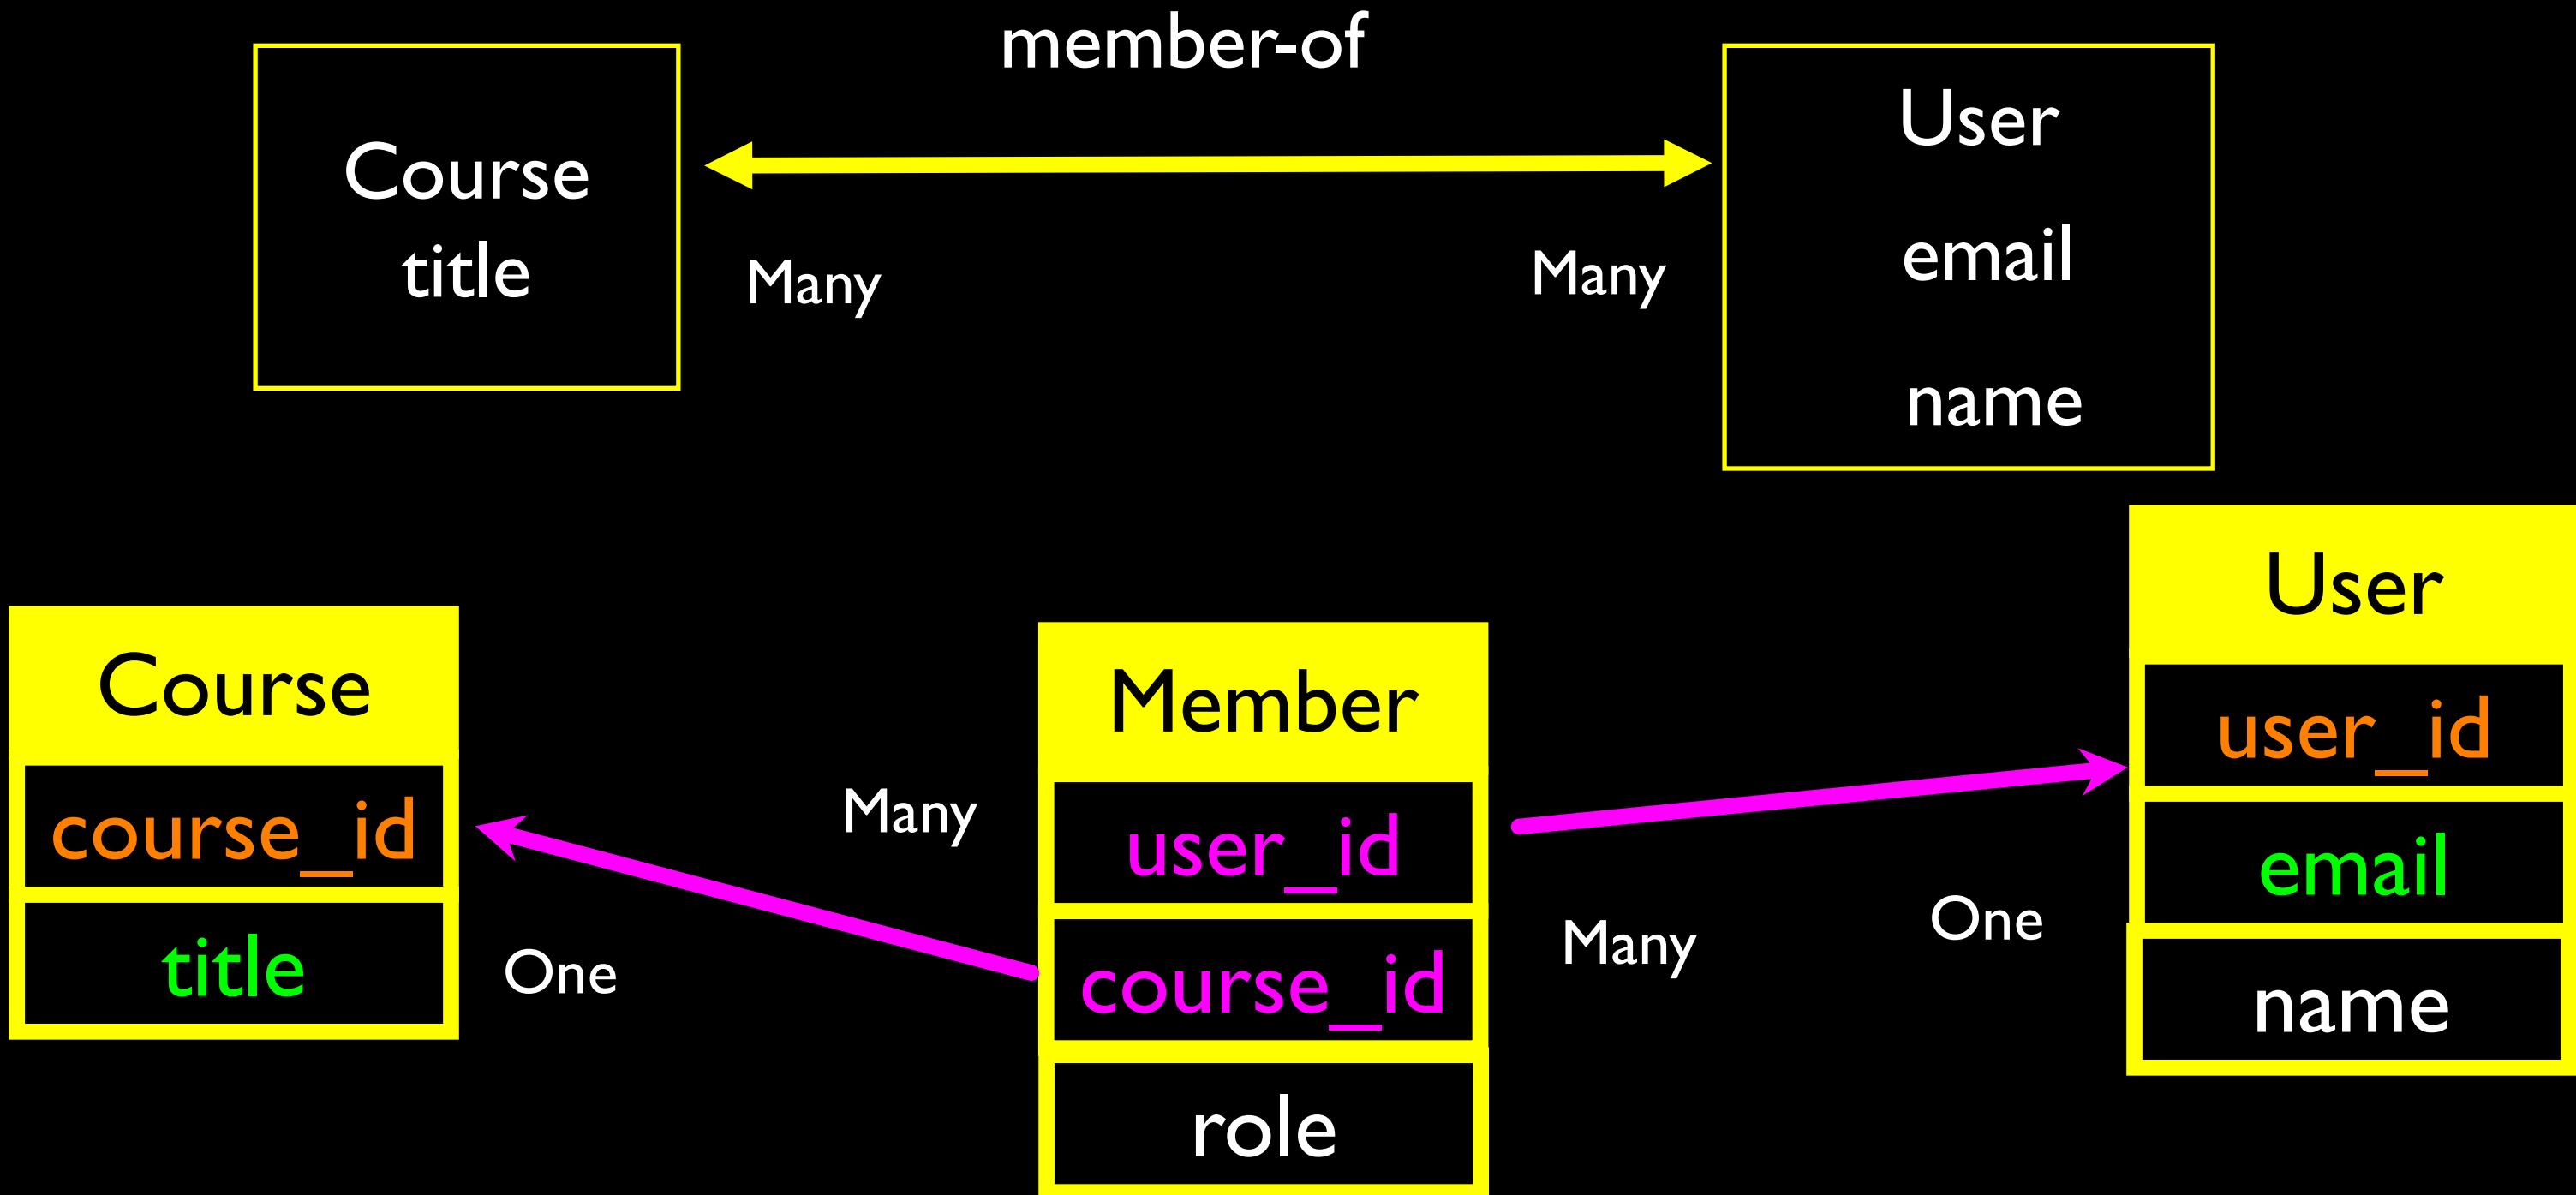


[https://en.wikipedia.org/wiki/One-to-many_\(data_model\)](https://en.wikipedia.org/wiki/One-to-many_(data_model))

Many to Many

- Sometimes we need to model a relationship that is many to many.
- We need to add a “connection” table with two foreign keys.
- There is usually no separate primary key.



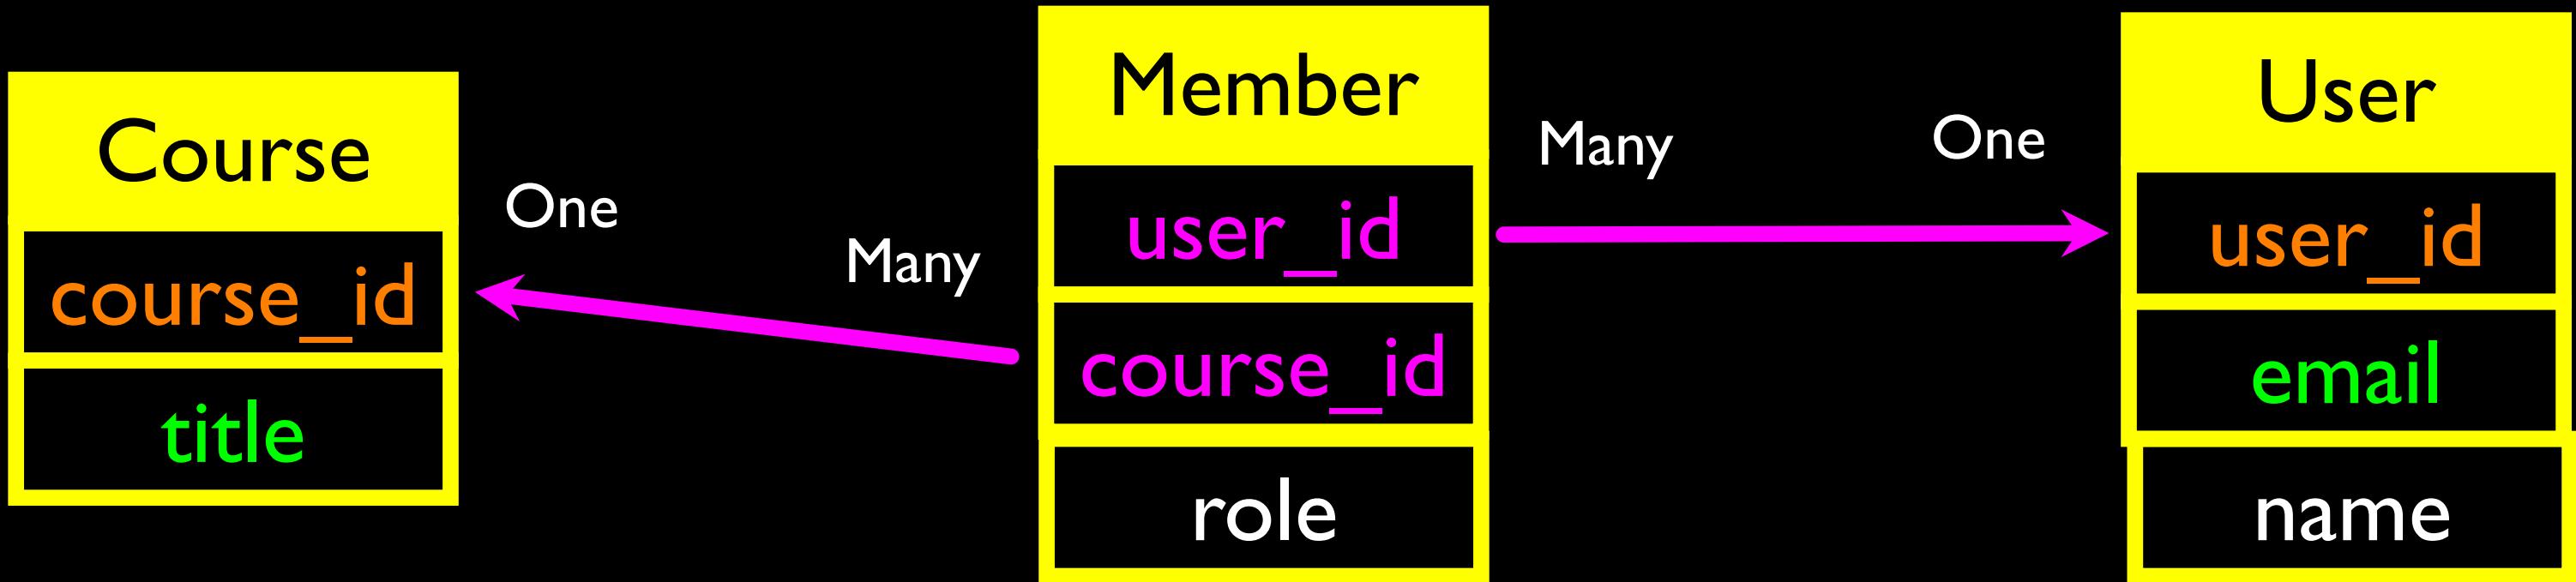




Start with a Fresh Database

```
CREATE TABLE User (
    user_id      INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    email        VARCHAR(128) UNIQUE,
    name         VARCHAR(128)
) ENGINE=InnoDB CHARACTER SET=utf8;
```

```
CREATE TABLE Course (
    course_id    INTEGER NOT NULL AUTO_INCREMENT PRIMARY KEY,
    title        VARCHAR(128) UNIQUE
) ENGINE=InnoDB CHARACTER SET=utf8;
```



```
CREATE TABLE Member (
    user_id          INTEGER,
    course_id        INTEGER,
    role             INTEGER,
    CONSTRAINT FOREIGN KEY (user_id) REFERENCES User (user_id)
        ON DELETE CASCADE ON UPDATE CASCADE,
    CONSTRAINT FOREIGN KEY (course_id) REFERENCES Course (course_id)
        ON DELETE CASCADE ON UPDATE CASCADE,
    PRIMARY KEY (user_id, course_id)
) ENGINE=InnoDB CHARACTER SET=utf8;
```

Insert Users and Courses

```
INSERT INTO User (name, email) VALUES ('Jane', 'jane@tsugi.org');  
INSERT INTO User (name, email) VALUES ('Ed', 'ed@tsugi.org');  
INSERT INTO User (name, email) VALUES ('Sue', 'sue@tsugi.org');
```

```
INSERT INTO Course (title) VALUES ('Python');  
INSERT INTO Course (title) VALUES ('SQL');  
INSERT INTO Course (title) VALUES ('PHP');
```

user_id	email	name
1	jane@tsugi.org	Jane
2	ed@tsugi.org	Ed
3	sue@tsugi.org	Sue

course_id	title
1	Python
2	SQL
3	PHP

Insert Memberships

user_id	email	name
1	jane@tsugi.org	Jane
2	ed@tsugi.org	Ed
3	sue@tsugi.org	Sue

course_id	title
1	Python
2	SQL
3	PHP

```
INSERT INTO Member (user_id, course_id, role) VALUES (1, 1, 1);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (2, 1, 0);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (3, 1, 0);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (1, 2, 0);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (2, 2, 1);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (2, 3, 1);
```

```
INSERT INTO Member (user_id, course_id, role) VALUES (3, 3, 0);
```

user_id	email	name
1	jane@tsugi.org	Jane
2	ed@tsugi.org	Ed
3	sue@tsugi.org	Sue

course_id	title
1	Python
2	SQL
3	PHP

user_id	course_id	role
1	1	1
1	2	0
2	1	0
2	2	1
2	3	1
3	1	0
3	3	0

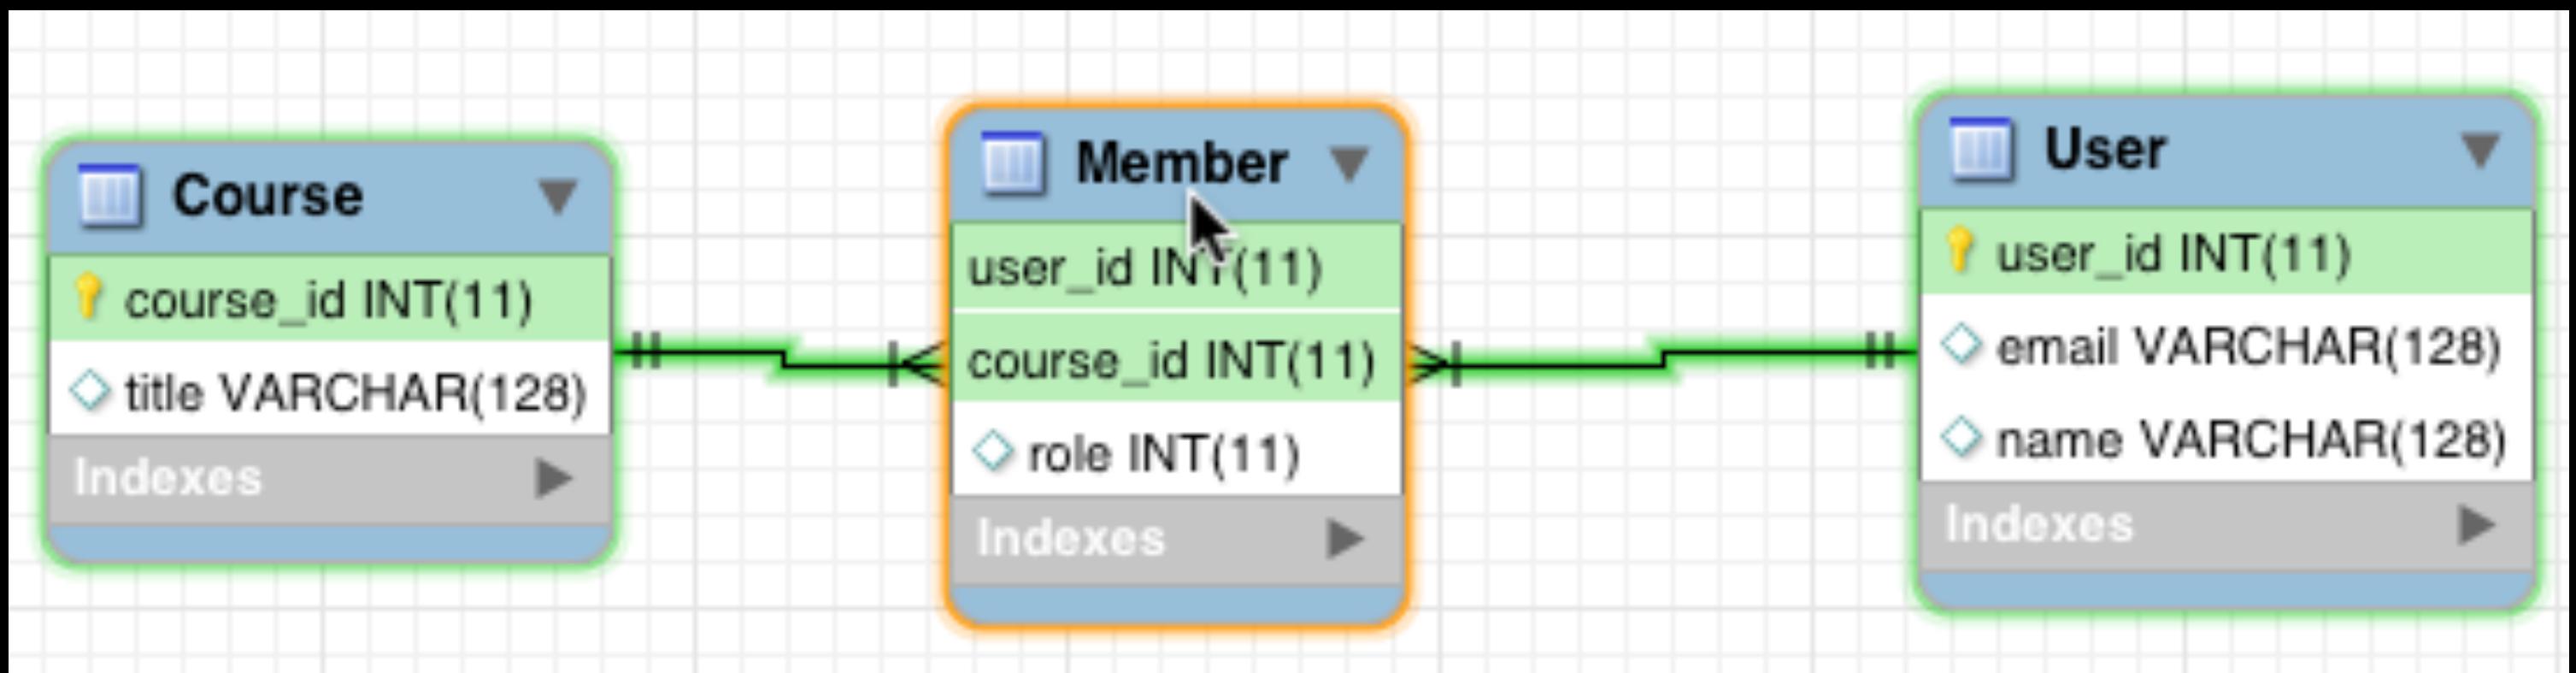
user_id	email	name
1	jane@tsugi.org	Jane
2	ed@tsugi.org	Ed
3	sue@tsugi.org	Sue

user_id	course_id	role
1	1	1
1	2	0
2	1	0
2	2	1
2	3	1
3	1	0
3	3	0

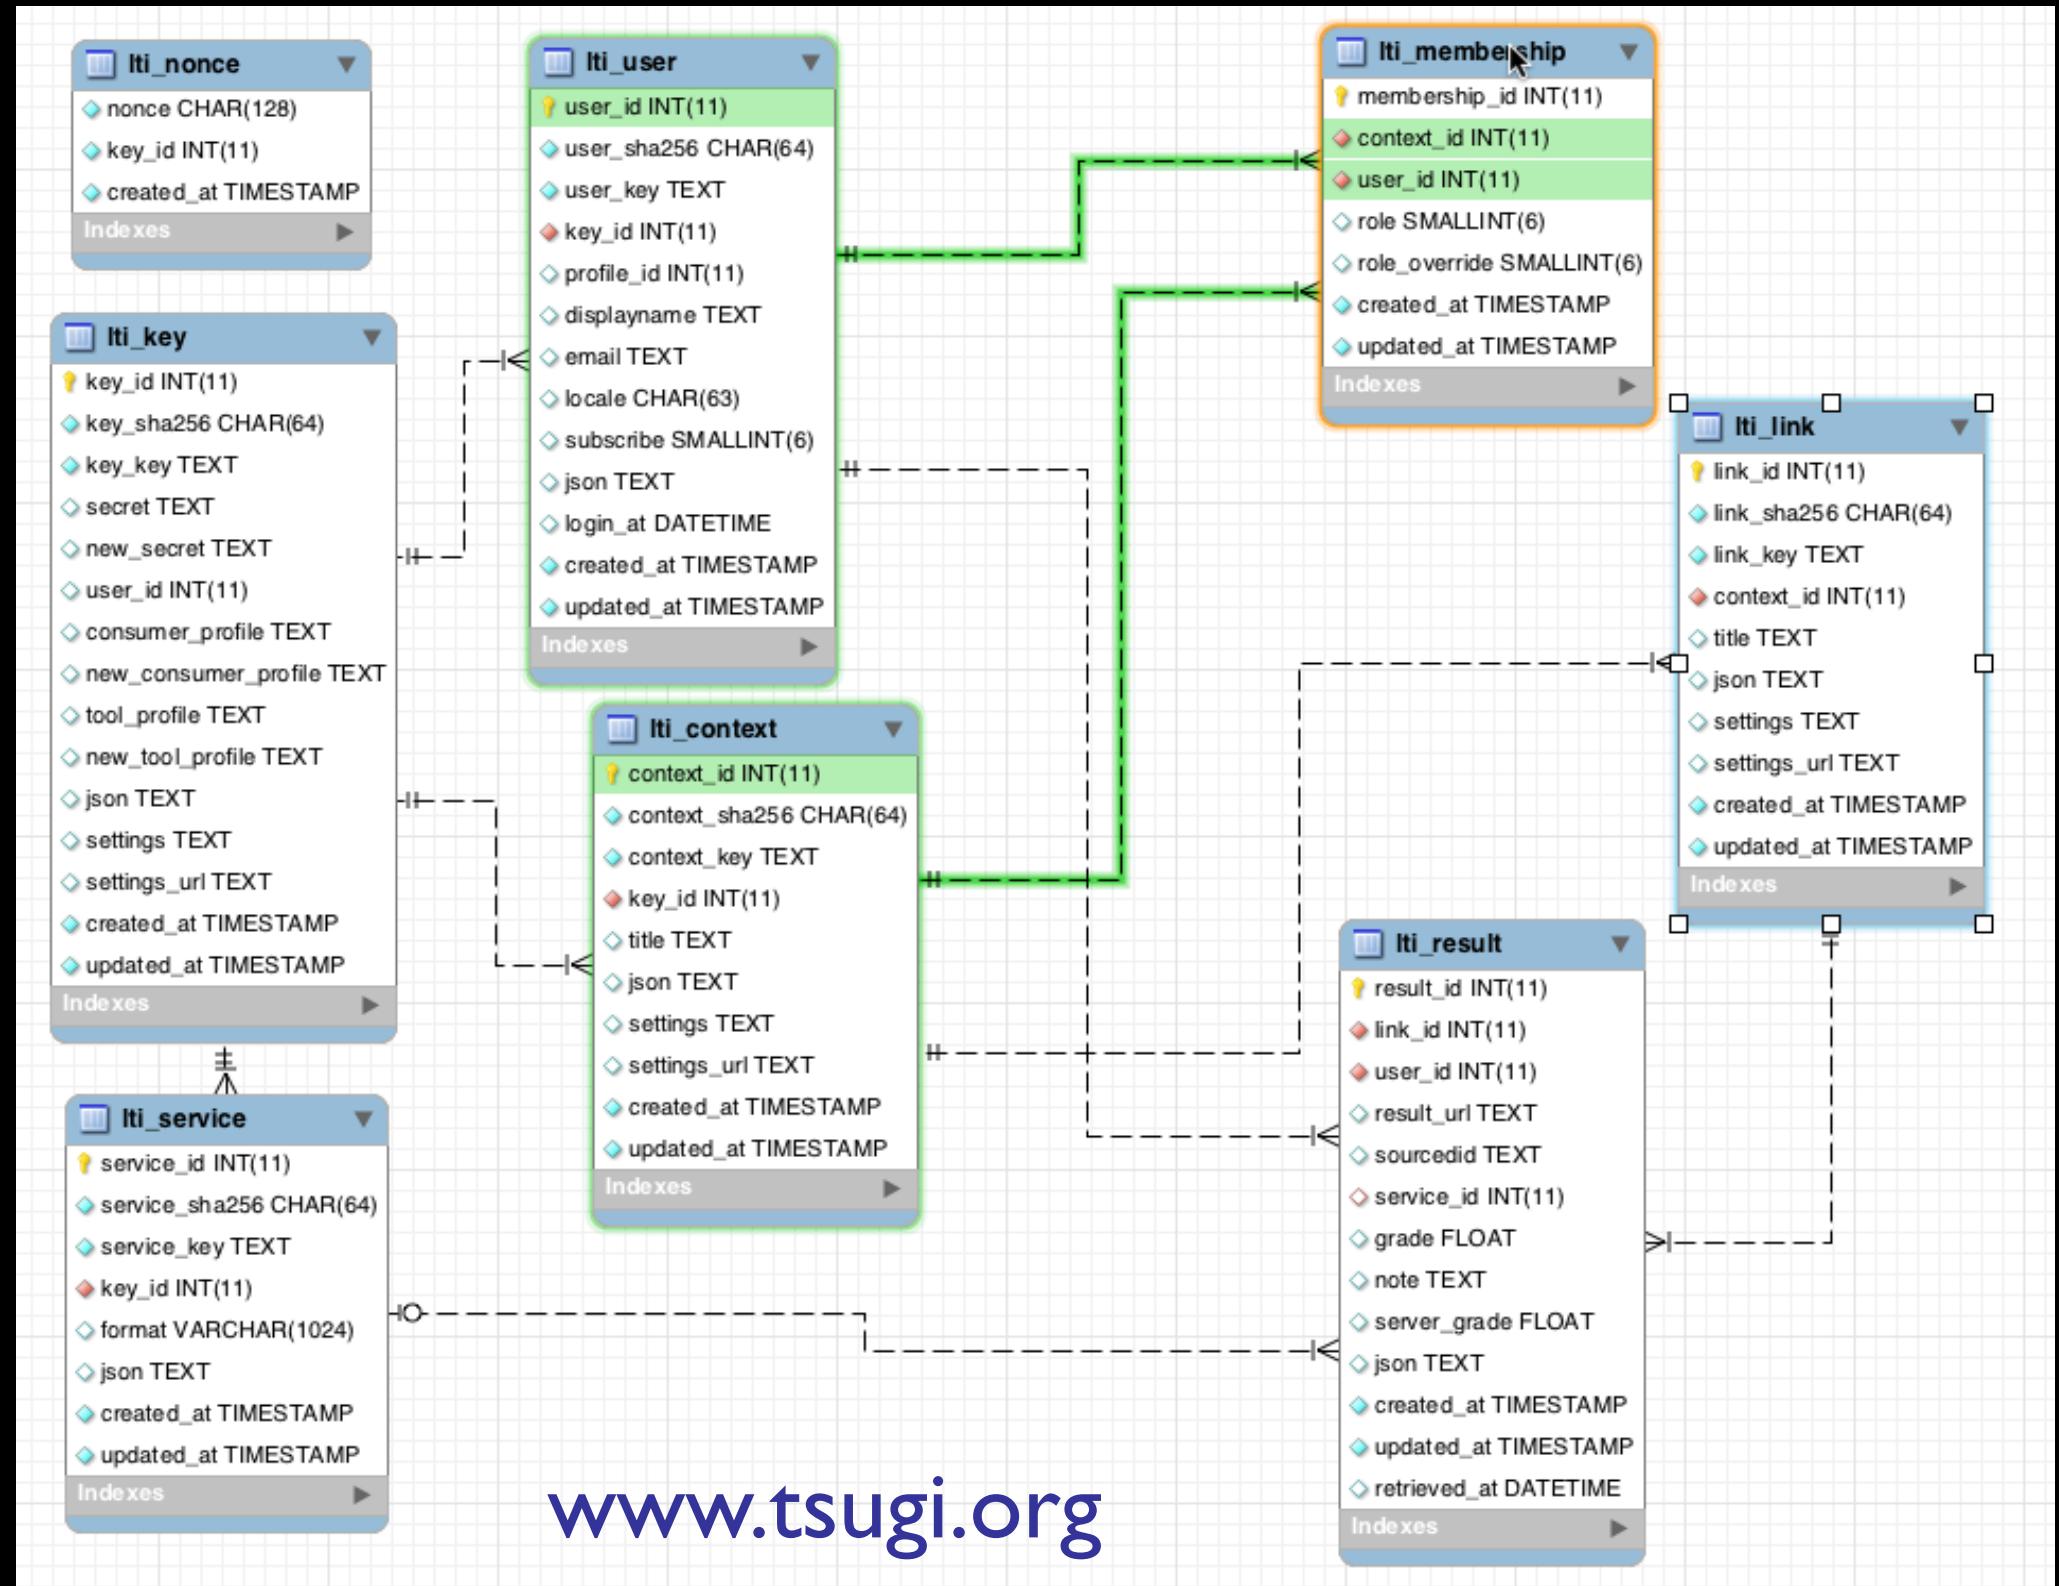
course_id	title
1	Python
2	SQL
3	PHP

name	role	title
Ed	1	PHP
Sue	0	PHP
Jane	1	Python
Ed	0	Python
Sue	0	Python
Ed	1	SQL
Jane	0	SQL

```
SELECT User.name, Member.role, Course.title
FROM User JOIN Member JOIN Course
ON Member.user_id = User.user_id AND
Member.course_id = Course.course_id
ORDER BY Course.title, Member.role DESC, User.name
```



<https://www.mysql.com/products/workbench/>



Complexity Enables Speed

- Complexity makes speed possible and allows you to get very fast results as the data size grows.
- By **normalizing the data and linking it with integer keys**, the overall **amount of data** which the relational database must **scan** is far lower than if the data were simply flattened out.
- It might seem like a **tradeoff** - spend some time designing your database so it continues to be fast when your application is a success.

Summary

- Relational databases allow us to **scale** to very large amounts of data.
- The key is to have **one copy of any data element** and use relations and joins to link the data to multiple places.
- This greatly **reduces the amount of data that must be scanned** when doing complex operations across large amounts of data.
- Database and SQL design is a bit of an **art form**.

Acknowledgements / Contributions



These slides are Copyright 2010- Charles R. Severance (www.dr-chuck.com) as part of www.wa4e.com and made available under a Creative Commons Attribution 4.0 License. Please maintain this last slide in all copies of the document to comply with the attribution requirements of the license. If you make a change, feel free to add your name and organization to the list of contributors on this page as you republish the materials.

Initial Development: Charles Severance, University of Michigan
School of Information

Insert new Contributors and Translators here including names and dates

Continue new Contributors and Translators here